

Stratified Connections, Optic Composition, and the Homotopy Fixed-Point Extension: A Categorical Framework for Viability-Weighted Curvature

X

Preprint, September 2026

Abstract. We develop the categorical and homotopy-theoretic framework for viability-weighted curvature: the geometric object that measures how non-commuting sequences of individually manageable environmental changes accumulate into a viability threat through the policy holonomy they induce. The framework has four layers. (i) The SAVGS object assembles a control base, a policy bundle with Fisher–Rao geometry, a viability margin, a maintenance graph, and a 2-categorical boundary span into a single stratified G_C -reduced bundle. (ii) The 2-category $\mathbf{StCon}(B)$ of stratified G_C -connections carries a lax-functorial gluing theorem and a piecewise holonomy formula; we state the small-loop expansion for closed loops that cross a constraint-switching wall transversally in *pairs*, with the boundary-reset contribution at its true order, in agreement with the machine-verified numerics. On constant-active-set strata the Fisher-minimal transport law defines the connection in closed form (KKT projection), and the small-loop viability–holonomy theorem bounds endpoint erosion of a viability margin by the viability-weighted curvature, the worst-case positive-part contraction of the curvature 2-form with the active survival covectors. (iii) A single composition theorem types seven heterogeneous domain bridges as optics, realizes the composed update on a complete metric state space, and establishes, for an explicit seven-map instantiation, a per-optic Lipschitz product bound and a Banach contraction of the Bregman-regularized update; the projected CPTP channel contraction settles the Zeno self-reference resolution. (iv) The filtered-colimit construction of RAF sets in $\mathbf{Optic}(\mathbf{Set})$ is proved componentwise and verified at scale, and the ∞ -categorical extension via homotopy type theory is developed with its proof-sketch status explicitly marked. An application of this framework to metabolic flux rerouting — the atomic curvature measure of parametric flux balance analysis and its genome-scale empirical validation — appears in a companion application paper.

Status: standalone theory paper; the application paper (X, 2026) cites it for constructions and proofs.

Contents

1	Introduction	3
2	Preliminaries	4
3	The SAVGS Framework	6
3.1	2-categorical gluing theorem for stratified connections	7
3.2	Rigorous Fisher-minimal transport law on constant-active-set strata	11
3.3	Square-root embedding and gauge invariance	13
3.4	Intervention-based autopoiesis closure test	13
4	Bregman-Divergence Noether Correspondence	15
5	The Seven-Claim Falsification Hierarchy	16
6	The Single Composition Theorem	17

7	Independent Per-Optic Lipschitz Constants and Unconditional Banach Contraction	19
7.1	Explicit instantiation of the seven forward maps	19
7.2	Projected CPTP channel contraction (Zeno self-reference resolution)	21
8	Filtered-Colimit Construction of RAFs	22
8.1	Componentwise filtered colimits in Optic(Set)	25
8.2	Filtered-colimit construction at scale ($ M > 10$)	26
9	∞-Categorical Extension via Homotopy Type Theory	27
9.1	Setup and definitions	27
9.2	The ∞ -categorical composition theorem	28
9.3	Univalence axiom and the canonical homotopy-fixed-point	29
9.4	Phase III closure-test definition: pathwise + univalence-corrected	29
9.5	Falsifiable prediction and operationalization	30
9.6	Implications	32
10	Future directions and open problems	32

1 Introduction

The problem. An adaptive system that must remain viable while its environment changes is governed, at each instant, by a policy that is optimal subject to active constraints. Classical viability theory supplies set-valued dynamics and tangential conditions, but no *geometric* object that measures how a non-commuting sequence of individually manageable changes accumulates into a viability threat. The obstruction is structural: at a constraint-switching boundary the active set changes, horizontal subspaces jump, the connection 1-form ceases to be defined, and the standard holonomy integral fails exactly where the question is most interesting. This paper constructs the missing object categorically: a stratified connection with a boundary gluing law, a viability-weighted curvature that contracts the curvature 2-form with survival covectors, and a compositional semantics (optics) under which heterogeneous domain bridges compose into one typed endomorphism whose fixed point is well posed and, for an explicit instantiation, contractive.

Contributions.

1. The *SAVGS object* (Definition 3.1): control base, policy bundle with Fisher–Rao geometry, viability margin, maintenance graph, and a 2-categorical boundary span, assembled into one stratified G_C -reduced bundle.
2. The 2-category $\mathbf{StCon}(B)$ of stratified G_C -connections (Definition 3.3), the 2-categorical gluing theorem (Theorem 3.4), and the piecewise holonomy formula (Theorem 3.6), with the small-loop expansion stated for closed loops crossing a switching wall transversally in pairs and the boundary-reset contribution at its true order, in agreement with the machine-verified numerics in five loop geometries and three structure-group regimes.
3. The *Fisher-minimal transport law* on constant-active-set strata in closed KKT form (Proposition 3.13) and the *small-loop viability-holonomy theorem* (Theorem 3.15): endpoint erosion of a viability margin is bounded, to leading order, by the viability-weighted curvature (Definition 2.10).
4. The Bregman–Hessian Noether correspondence (Proposition 4.1) with its falsifiable precondition check, giving the gauge-covariance statement for the curvature.
5. The *single composition theorem* (Theorem 6.4): seven domain bridges typed as optics, the realized update on a complete metric state space (Definition 7.2), per-optic analytic Lipschitz constants (Lemma 7.1), the Banach contraction of the Bregman-regularized update for the chosen seven-map instantiation (Theorem 7.4), and the projected CPTP channel contraction (Theorem 7.6) settling the Zeno self-reference resolution.
6. The filtered-colimit construction of RAF sets in $\mathbf{Optic}(\mathbf{Set})$, proved componentwise (Theorem 8.3) and verified at scale (Proposition 8.7).
7. The ∞ -categorical extension via homotopy type theory (Theorem 9.2, Corollary 9.3), with proof-sketch status explicitly marked (Remark 9.4), and the three-phase autopoiesis closure test built on it (Definition 9.5).

Relation to the application paper. An application of this framework to metabolic flux rerouting appears in a separate manuscript (X, 2026): there the viability-weighted-curvature idea is realized measure-theoretically as an atomic curvature measure on the active-set complex of parametric flux balance analysis, a gene-level sensitivity metric κ^μ is derived from it, and the metric is validated at genome scale against carbon-depletion transcriptional response, with an explicit protein-layer null and selection-rule robustness checks. The application paper states the load-bearing definitions of the present framework in brief, adapted form and cites this paper for the constructions and proofs; the present paper develops the full categorical machinery and cites the application paper where the empirical line is concerned. The division of labor is deliberate:

the empirical claims of the application paper are statistical and self-contained, the categorical claims of this paper are mathematical and self-contained, and the two papers share no verbatim passages of length.

Status conventions. Theorem-level claims are labeled by proof status throughout: complete proofs; proofs that reduce to cited descent or coherence theory (marked in remarks); proof sketches (marked); and machine-verified statements, whose scripts and artifacts are committed in the project repository and regenerated by them. Two results of the predecessor lineage that were superseded rather than repaired — the smooth-surrogate envelope line and the indicator-weighted operational curvature — are *not* carried here: the first is superseded in the application paper by the measure-theoretic development, the second by its locked metric. Open problems are collected in Section 10.

2 Preliminaries

Definition 2.1 (Viability depth functional D_V ; distinct from curvature). Let M be a smooth state manifold, $V : M \rightarrow \mathbb{R}$ a smooth viability function with maximum $V_{\max} = \max_M V$, and $\gamma_a : [0, 1] \rightarrow M$ a smooth policy loop of amplitude a , parametrized at unit speed so that $|\gamma_a| = \int_0^1 |\dot{\gamma}_a(t)| dt$ is the loop length. The *viability depth functional* is

$$D_V(\gamma_a) = \frac{1}{|\gamma_a|} \int_{\gamma_a} (V_{\max} - V) ds = \frac{1}{|\gamma_a|} \int_0^1 (V_{\max} - V \circ \gamma_a(t)) |\dot{\gamma}_a(t)| dt. \quad (1)$$

For the radially symmetric viability $V(x_1, \dots, x_{n-1}) = 1 - \sum_i x_i^2$ and the circular loop $\gamma_a(t) = (a \cos 2\pi t, a \sin 2\pi t, 0, \dots, 0)$ at unit speed, $D_V(a) = a^2$. The viability depth is a loop-averaged deficit; it is *not* a curvature 2-form: it carries no orientation, no dependence on path ordering, and no commutator structure. The curvature-based viability-weighted curvature κ_V (Definition 2.10 below, contracting the curvature 2-form F with viability covectors) is a distinct object; the equality $D_V(a) = \kappa_V(a) = a^2$ in the radial prototype is a model-specific relation, not an identity.

Definition 2.2 (Structure group: $O(r)$, $SO(r)$, or $CO(r)$ with Weyl scale). The Fisher–Rao metric on a policy fiber of dimension r admits an orthonormal frame bundle with structure group $O(r)$ (Kobayashi and Nomizu (1969); Amari and Nagaoka (2000)). If the policy frame is oriented, the structure group reduces to $SO(r)$. A reduction to the conformal group $CO(r) = \mathbb{R}_+ \ltimes SO(r)$ requires an additional *scale (Weyl) structure* on the policy bundle: a positive scale bundle $L \rightarrow E$ together with a conformal class $[g^F]$ of the Fisher metric and a Weyl connection preserving the conformal class. The three cases are distinct modeling commitments; Chentsov’s uniqueness theorem Chentsov (1982) fixes the Fisher metric under statistical-map invariance but does *not* by itself select among $O(r)$, $SO(r)$, or $CO(r)$. For the framework’s prototypes we use: $G_C = SO(2)$ for the abelian $n = 3$ prototype; $G_C = SO(3)$ for the non-abelian $n = 4$ prototype; and $CO(r)$ only when an endogenous scale variable (effective precision, temperature, sample size, or Weyl factor) is explicitly modeled (Remark 2.3).

Remark 2.3 (Fisher–Weyl policy structure, when $CO(r)$ is justified). A Fisher–Weyl policy structure consists of: (i) a policy bundle $E \rightarrow B$; (ii) a vertical Fisher metric g^F ; (iii) a positive scale bundle $L \rightarrow E$; (iv) a conformal class $[g^F]$; (v) a Weyl connection preserving the conformal class. The conformal frame bundle then has structure group $CO(r)$. If the cost functional depends only on the conformal class of g^F and is invariant under positive rescalings of orthonormal frames, its stabilizer contains $CO(r)$; if scale is fixed, the stabilizer reduces to $SO(r)$ or $O(r)$. This preserves the $CO(r)$ ambition in a mathematically honest way: $CO(r)$ is an *added modeling structure*, not a consequence of Chentsov’s theorem alone.

Remark 2.4 (Physical meaning of the $n = 3 \rightarrow n = 4$ transition). At $n = 3$ the policy simplex is two-dimensional, the Fisher–Rao stabilizer is $CO(2) = \mathbb{R}_+ \ltimes SO(2)$, and the Lie algebra $\mathfrak{so}(2)$ is

one-dimensional, generated by the single infinitesimal rotation about the origin. Two rotations about the same origin necessarily commute (they lie in the same one-parameter subgroup), so the holonomy of a sequence of loops is path-independent up to homotopy, and the path-ordering in the parallel-transport exponential collapses to ordinary commutative multiplication. The non-abelian signature (Claim F) is then unobservable: the commutator $[R_1, R_2] = 0$ identically, regardless of the rotation amplitudes. At $n = 4$ the policy simplex is three-dimensional, the stabilizer is $\text{CO}(3) = \mathbb{R}_+ \ltimes \text{SO}(3)$, and $\mathfrak{so}(3)$ is three-dimensional with three independent generators $\{\mathfrak{l}_x, \mathfrak{l}_y, \mathfrak{l}_z\}$ satisfying $[\mathfrak{l}_z, \mathfrak{l}_x] = \mathfrak{l}_y$ etc. Two rotations about *distinct* axes (e.g. $R_z(\alpha_1)R_x(\alpha_2)$ vs $R_x(\alpha_2)R_z(\alpha_1)$) no longer commute, and the small-angle commutator is $[R_1, R_2] \approx \alpha_1\alpha_2[\mathfrak{l}_i, \mathfrak{l}_j]$ with Frobenius norm $\sqrt{2}\alpha_1\alpha_2$ (each $\mathfrak{so}(3)$ basis element has Frobenius norm $\sqrt{2}$). The non-abelian signature is therefore a quantitative, falsifiable prediction specific to $n \geq 4$; it cannot be tested at $n = 3$ by any experimental design.

Definition 2.5 (Optic category). Following Riley [Riley \(2018\)](#), for a category $\mathcal{C}$ with finite limits the *optic category* $\mathbf{Optic}(\mathcal{C})$ has as objects triples (M, C, R) of a forward carrier M , a backward carrier C , and a residual R , and as morphisms pairs (φ, ρ) of a forward map $\varphi : M \rightarrow M'$ and a backward map $\rho : R' \rightarrow R$ satisfying the residual-compatibility condition. The monoidal structure of $\mathbf{Optic}(\mathcal{C})$ supplies composition, associativity, and unitality (Proposition 2.3 of [Riley \(2018\)](#)).

Definition 2.6 (RAF set). Following Hordijk and Steel [Steel and Hordijk \(2004\)](#); [Hordijk and Steel \(2011\)](#), given a catalytic reaction network $(X, \mathcal{R}, F, c)$ over a molecule set X with food set $F \subset X$ and catalysis assignment $c : \mathcal{R} \rightarrow 2^X$, a subset $R' \subseteq \mathcal{R}$ is a *reflexively autocatalytic food-generated (RAF) set* if (i) every $r \in R'$ is catalyzed by some element of $F \cup \bigcup_{r' \in R'} \text{react}(r')$, (ii) the food-generated closure of R' contains all reactants of R' , and (iii) R' is closed under the induced closure operator.

Definition 2.7 (CPTP channel). A *completely-positive trace-preserving (CPTP) channel* on a finite Hilbert space $\mathcal{H}$ is a linear map $\Phi : \text{End}(\mathcal{H}) \rightarrow \text{End}(\mathcal{H})$ of the form $\Phi(\rho) = \sum_i K_i \rho K_i^\dagger$ with Kraus operators $\{K_i\}$ satisfying $\sum_i K_i^\dagger K_i = I$. The measurement schedule of Φ is the sequence $(\tau_k, P_k)_{k \geq 1}$ of inter-measurement intervals τ_k and projective measurements P_k . The Liouvillian $\mathcal{L}$ of the underlying Markov process has a spectral gap $\Delta = \min_{\lambda \neq 0} \text{Re}(\lambda)$.

Definition 2.8 (Bregman divergence). A *Bregman divergence* [Bregman \(1967\)](#); [Amari and Nagaoka \(2000\)](#) generated by a strictly convex differentiable function ϕ is

$$D_\phi(p, q) = \phi(p) - \phi(q) - \langle \nabla \phi(q), p - q \rangle. \quad (2)$$

The divergence is not symmetric in general. The *dual affine coordinates* are p (the primal) and $\nabla \phi(p)$ (the dual). The Bregman divergence is affine-invariant in the sense that an affine transformation in the primal coordinate, paired with the corresponding affine transformation in the dual, leaves D_ϕ unchanged.

Definition 2.9 (Algorithmic rate-distortion distance). Fix a universal Turing machine U , a distortion measure d , and a distortion threshold $D \geq 0$. The *algorithmic rate-distortion distance* of x at threshold D is

$$\text{dist}_D(x) = \min\{|p| : U(p) \text{ outputs } \hat{x}, d(x, \hat{x}) \leq D\}. \quad (3)$$

It is monotone non-increasing in D , bounded above by the Kolmogorov complexity $K(x)$, and upper-semicomputable by dovetailing over all programs. It is a single-string, intrinsically deterministic quantity: no set-average rate and no random-coding gap arise.

Definition 2.10 (Viability-weighted curvature κ_V). Let $(B, E, \mathcal{P}, \varepsilon, \Gamma)$ be a SAVGS (Definition 3.1) with a G_C -connection on the policy bundle whose curvature 2-form on a stratum is F . Let

$$h_\alpha(\theta, x) = D_\phi(\text{dist}_D(x), \text{dist}_D(x_0)) \quad (4)$$

be a viability margin built from a Bregman divergence (Definition 2.8) evaluated at the algorithmic rate-distortion distance (Definition 2.9) of the current state relative to a reference x_0 , with $h_\alpha > 0$ on the open feasibility set. The *per-pair viability-weighted curvature* is

$$\kappa_\alpha(\theta, x; u, v) = \frac{[-D_{F(u,v)}h_\alpha]^+(\theta, x)}{h_\alpha(\theta, x)}, \quad (5)$$

where $[\cdot]^+$ is the positive part: only viability-eroding curvature contributes. The *viability-weighted curvature* is the worst case over unit-area bivectors and active viability covectors,

$$\kappa_V(\theta, x) = \sup_{\substack{u, v \in T_\theta B \\ \|u \wedge v\|_{g_F} = 1}} \max_{\alpha: h_\alpha(\theta, x) > 0} \kappa_\alpha(\theta, x; u, v). \quad (6)$$

The quantity $\kappa_V(\theta, x)$ is the worst fractional loss of viability margin per unit oriented environmental-loop area, to leading order; it contracts the geometric defect F with the survival covectors Dh_α instead of reporting $\|F\|$ alone. Gauge invariance of κ_V follows from the square-root embedding (Proposition 3.18) together with the Bregman–Noether correspondence (Proposition 4.1); the wedge-norm constraint makes κ_V a per-area quantity, matching the leading-order expansion of Theorem 3.15. The instantiation (4) is one admissible choice: any C^2 viability function with the stated positivity works, and the measure-theoretic discretization used in the application paper (X, 2026) replaces the 2-form F by an atomic crease measure.

3 The SAVGS Framework

The Stratified Autopoietic Viability Geometric System (SAVGS) is the minimal object on which the viability-weighted curvature can be stated without the type confusions, gauge ambiguities, and missing premises of prior formulations. It assembles five components that prior work treats separately or treats at the wrong level of abstraction.

Definition 3.1 (SAVGS). The SAVGS is the tuple $(B, E, \mathcal{P}, \varepsilon, \Gamma)$:

1. A continuous control base manifold $B \subset \mathbb{R}^d$ (denoted Θ in earlier drafts), with coordinates including food scarcity, danger intensity, and sensor noise.
2. A policy bundle $\pi : E \rightarrow B$ with total space E and policy fiber $\mathcal{P} = \pi^{-1}(\theta)$, where $\mathcal{P}$ is the open simplex $\Delta_\circ^{m-1}$ (for m actions), the circle S^1 (abelian $n = 3$ prototype), the rotation group $SO(3)$ (non-abelian $n = 4$ prototype), or another declared policy manifold. The bundle projection $\pi : E \rightarrow B$ is the canonical projection onto the base manifold (not the wrong-direction $\Delta^{n-1} \rightarrow \Theta$ of earlier drafts); on the trivial bundle $E = B \times \mathcal{P}$ this is $\pi(\theta, p) = \theta$.
3. A viability margin $\varepsilon \geq \varepsilon_{\min} > 0$, with the inequality strict on the open feasibility set and non-strict on the closed viability kernel.
4. An endogenous maintenance graph $\Gamma = (M, R, E_\Gamma)$ of maintenance operations M , regeneration rules R , and maintenance edges E_Γ (renamed to avoid clash with total space E).
5. A 2-categorical span $\text{Stab}_1 \leftarrow \text{Boundary} \rightarrow \text{Stab}_2$ that resolves the constraint-switching boundary discontinuity between strata (Theorem 3.4 for the full gluing theorem).

The five components compose into a single object whose geometry is a stratified G_C -reduced principal bundle over the control manifold, with policy fibers over each stratum and viability kernels defined by non-strict inequalities on each stratum. The viability-weighted curvature of Definition 2.10 is computed on this object.

Remark 3.2 (Why the 2-categorical span is needed: smooth-connection breakdown at constraint-switching boundaries). The Ehresmann connection of the geometric construction requires

smooth variation of horizontal subspaces across the base manifold. At a constraint-switching boundary—where the active set of inequality constraints on the viability margin ε changes—the horizontal subspace changes discontinuously. The connection is therefore no longer smooth at the boundary, and the standard holonomy integral breaks down. The fifth SAVGS component resolves this by introducing a 2-categorical span $\text{Stab}_1 \leftarrow \text{Boundary} \rightarrow \text{Stab}_2$ that glues the two smooth strata via a boundary 2-cell; the connection is defined piecewise on each stratum and a matching condition on the boundary supplies the 2-categorical compatibility. The formal development of the 2-categorical structure — the lax-functorial gluing and the boundary 2-cell’s coherence with the connection 1-forms — is given in Section 3.1 below, and the homotopy-theoretic well-definedness of the resulting holonomy is treated in Section 9. The operational requirement on the abelian and non-abelian prototypes is that the active set remains constant along each loop, which places the loop entirely within a single smooth stratum and avoids the boundary discontinuity.

3.1 2-categorical gluing theorem for stratified connections

We now close the 2-categorical gluing agenda sketched in Remark 3.2 by constructing the 2-category $\mathbf{StCon}(B)$ of stratified G_C -connections on a stratified base B , proving the gluing theorem, and deriving the piecewise holonomy formula. Throughout, $G_C \in \{O(r), SO(r), \text{CO}(r)\}$ is the structure group selected by the policy-fiber geometry (Definition 2.2).

Definition 3.3 (2-category $\mathbf{StCon}(B)$ of stratified connections). Let B be a stratified manifold with strata $\{S_i\}_{i \in I}$ ($B = \bigsqcup_i S_i \cup \bigcup_{ij} B_{ij}$ where $B_{ij} = S_i \cap S_j$ are boundary hypersurfaces, and $T_{ijk} = S_i \cap S_j \cap S_k$ are triple intersections). The 2-category $\mathbf{StCon}(B)$ of stratified G_C -connections is:

- *Objects*: stratified principal G_C -bundles $P \rightarrow B$ with a connection $A = \{A_i \text{ on } S_i, g_{ij} : B_{ij} \rightarrow G_C\}$ satisfying:

(O1) A_i is a smooth connection 1-form on stratum S_i ;

(O2) $g_{ij} : B_{ij} \rightarrow G_C$ is a smooth transition function;

(O3) the matching condition on B_{ij} : with the convention that the transition g_{ij} acts on frames as $e_j = e_i \cdot g_{ij}$ (so g_{ij} carries stratum S_i ’s frame to stratum S_j ’s frame), the connection forms are related by the standard gauge transformation

$$A_j = g_{ij}^{-1} A_i g_{ij} + g_{ij}^{-1} dg_{ij}, \quad (7)$$

whose infinitesimal form ($g_{ij} = \exp u_{ij}$ near the identity) is $A_j = A_i + du_{ij} + [A_i, u_{ij}] + O(u_{ij}^2)$. For abelian $G_C = U(1)$ this reduces to $A_j - A_i = d\alpha_{ij}$ where $g_{ij} = e^{i\alpha_{ij}}$; the reverse-direction reading ($g_{ji} = g_{ij}^{-1}$ carrying A_j to A_i) is equivalent and is the one used in the holonomy formula below;

(O4) the Čech cocycle on T_{ijk} : $g_{ij} \cdot g_{jk} = g_{ik}$ (multiplication in G_C).

- *1-Morphisms*: connection-preserving equivariant bundle maps $\phi : P \rightarrow P'$, smooth on each stratum, with $\phi \circ g_{ij} = g'_{ij} \circ \phi$ on each boundary B_{ij} .
- *2-Morphisms*: gauge transformations $\eta : \phi \Rightarrow \phi'$ (vertical automorphisms of P' that conjugate ϕ to ϕ'), smooth on each stratum, with the coherence 2-cell condition at triple intersections (η commutes with g_{ij} on T_{ijk} in the sense of strict 2-functorial equality).

Theorem 3.4 (2-categorical gluing theorem). *Given a stratified base B with strata $\{S_i\}$, boundary hypersurfaces $\{B_{ij}\}$, a smooth Čech cocycle $\{g_{ij} : B_{ij} \rightarrow G_C\}$ satisfying $g_{ij} \cdot g_{jk} = g_{ik}$ on triple intersections T_{ijk} , and connection 1-forms $\{A_i \text{ on } S_i\}$ satisfying the matching condition $g_{ij}^* A_j = A_i + d(\log g_{ij}) - [A_i, \log g_{ij}]$ on each B_{ij} , there exists a unique (up to 2-isomorphism) object (P, A) in $\mathbf{StCon}(B)$ whose restriction to each S_i is A_i and whose transition on each B_{ij} is g_{ij} .*

Proof. The proof is by 2-descent in the 2-stack of stratified G_C -bundles with connection. The 2-stack property (objects glue uniquely up to 2-isomorphism on overlaps; 1-morphisms glue uniquely up to 2-isomorphism; 2-morphisms glue strictly) follows from the Giraud axioms for the underlying G_C -bundles extended to stratified spaces:

- *Object gluing:* the underlying G_C -bundle P is constructed by the standard clutching construction on each stratum, glued along the boundary via the transition g_{ij} . The Čech cocycle (O4) ensures consistency at triple intersections. The result is unique up to bundle isomorphism.
- *Connection gluing:* the connection 1-forms $\{A_i\}$ glue via the matching condition (O3): on B_{ij} , g_{ij} is a gauge transformation carrying A_j to A_i , so A_i on S_i and A_j on S_j define a single stratified connection on P . The glued connection is unique up to gauge transformation (2-isomorphism).
- *1-morphism gluing:* given two stratified connections (P, A) and (P', A') and a 1-morphism $\phi_i : P|_{S_i} \rightarrow P'|_{S_i}$ on each stratum that commutes with g_{ij} on boundaries, the $\{\phi_i\}$ glue to a unique 1-morphism $\phi : P \rightarrow P'$, with the boundary compatibility $\phi \circ g_{ij} = g'_{ij} \circ \phi$ supplying the gluing data.
- *2-morphism gluing:* gauge transformations are determined locally; on overlaps they agree by the strict 2-functorial coherence condition, so 2-morphisms glue strictly.

Hence $\mathbf{StCon}(B)$ is a 2-stack (a 2-sheaf of groupoids), and the gluing theorem holds by 2-descent [Giraud \(1971\)](#); [Breen \(1994\)](#); [Breen and Messing \(2005\)](#); [Bartels \(2004\)](#). $\square$

Remark 3.5 (Proof status of Theorem 3.4). The proof reduces the gluing to the 2-stack descent property for stratified G_C -bundles: the Giraud-type axioms for the underlying bundles are cited [Giraud \(1971\)](#); [Breen and Messing \(2005\)](#), not re-proved, and the extension from smooth spaces to stratified spaces (stratum-by-stratum clutching with boundary matching) is verified here at the level of construction, with the coherence of the boundary 2-cells checked on triple intersections. The piecewise holonomy consequences (Theorem 3.6), by contrast, are machine-verified in five loop geometries and all three structure-group regimes (Remarks 3.7–3.11).

Theorem 3.6 (Piecewise holonomy formula). *Let $(P, A) \in \mathbf{StCon}(B)$ be a stratified connection, and let $\gamma : [0, 1] \rightarrow B$ be a smooth loop that crosses finitely many boundaries transversally at points $\{p_k\}_{k=1}^n$ in cyclic order ($p_1 < p_2 < \dots < p_n$). Let γ be decomposed as $\gamma = \gamma_1 \sqcup \gamma_2 \sqcup \dots \sqcup \gamma_n$ where each piece γ_k lies in stratum S_{i_k} . Then the stratified holonomy is*

$$\mathrm{Hol}^{\mathrm{strat}}(\gamma) = \mathrm{Hol}_{S_{i_n}}(\gamma_n) \cdot g_{i_n i_1}(p_n) \cdot \mathrm{Hol}_{S_{i_{n-1}}}(\gamma_{n-1}) \cdot \dots \cdot g_{i_1 i_2}(p_1) \cdot \mathrm{Hol}_{S_{i_1}}(\gamma_1), \quad (8)$$

a product of stratum-holonomies with boundary-transition maps, the order fixed by the chronological (right-to-left) traversal of the loop. For a closed loop the number of transversal crossings of any single wall is necessarily even: a small loop of area ε^2 centered on a wall B_{+-} crosses it twice, at points $p_+ \neq p_-$ with $|p_+ - p_-| = O(\varepsilon)$. For such a loop the piecewise curvature formula is

$$H^{\mathrm{strat}}(\gamma_\varepsilon) = \exp\left(-\varepsilon^2 \left[\iint_{\Sigma \cap S_+} F_+ dA + \iint_{\Sigma \cap S_-} F_- dA \right] - [\log g_{+-}(p_+) - \log g_{+-}(p_-)] + O(\varepsilon^3) \right), \quad (9)$$

where $F_\pm$ are the curvature 2-forms on strata $S_\pm$ and Σ is the small enclosed disk. The two boundary resets cancel at leading order by the cocycle condition ($g_{-+} = g_{+-}^{-1}$ on the wall): the residual boundary contribution $\log g_{+-}(p_+) - \log g_{+-}(p_-)$ is the tangential variation of the transition function between the two crossing points, of order $O(\varepsilon)$; it vanishes identically when the transition function is constant along the wall, and it is linear in the wall-coordinate difference in the machine-verified regime of Remark 3.7. An $O(1)$ reset term $\log g_{+-}(p_)$ arises only for an open arc that crosses the wall once and is not closed by the reverse transition.*

Proof. By Theorem 3.4, the stratified connection A exists and is unique up to 2-isomorphism, so the holonomy is well-defined. Decompose γ into pieces $\{\gamma_k\}$ on strata $\{S_{i_k}\}$ separated by boundary crossings $\{p_k\}$. The parallel transport of A along γ factorizes as the product of stratum-parallel transports on each γ_k (using $A|_{S_{i_k}} = A_{i_k}$ and Stokes' theorem on each stratum) interleaved with the boundary transitions $g_{i_k i_{k+1}}(p_k)$ applied at each crossing. The ordering (right-to-left in (8)) follows the chronological traversal of the loop; the composition order is genuinely order-dependent for non-abelian G_C (Remark 3.8). For an ε -small *closed* loop crossing the wall twice, the stratum-holonomies reduce by Stokes to $\varepsilon^2 \iint_{\Sigma \cap S} F_S dA$ on each stratum (where Σ is the small enclosed disk), while the two boundary transitions contribute $\log g_{+-}(p_+) - \log g_{+-}(p_-)$ in the Lie algebra: the $O(1)$ parts of the reset pair cancel by the cocycle condition, and the residual is the tangential variation of the transition function between the crossing points, of order $O(\varepsilon)$. For abelian G_C with the matching condition (7) the cancellation is exact at leading order: along the short segment of γ inside S_- , the jump $A_- - A_+ = d\alpha_{+-}$ integrates to $\alpha(p_-) - \alpha(p_+)$, which cancels the reset variation identically, so the holonomy is governed by the curvature terms at ε^2 plus the stated $O(\varepsilon)$ residual. This is the regime verified numerically in Remark 3.7 (boundary reset linear in the wall coordinate with slope ε , the loop width). $\square$

Remark 3.7 (Numerical verification of the piecewise holonomy formula). The piecewise holonomy formula — including the two-crossing small-loop form of (9), with the boundary contribution $\alpha(p_+) - \alpha(p_-)$ at its true order $O(\varepsilon)$ — is numerically verified by script `two_cat_gluings_stratified.py` on a two-stratum base $B = \mathbb{R}^2$ with the x -axis as the boundary, $G_C = U(1)$ (abelian, simplifying the matching condition), and constant-curvature connections $A_{\pm} = (c_{\pm}/2)(x dy - y dx)$ on the upper (S_+) and lower (S_-) half-planes. The transition function $g_{+-}(x, 0) = e^{i\alpha(x)}$ with $\alpha(x) = a_1 x$ is the boundary reset map. The verification: (a) computes the analytic holonomy $H_{\text{an}}(\gamma) = e^{-i(F\varepsilon^2 + \alpha(p_+) - \alpha(p_-))}$ for a rectangular loop of side ε centered at $(x_c, 0)$ crossing the boundary twice; (b) computes the numerical holonomy by parallel transport along the loop with the connection 1-form A and the boundary transitions; (c) confirms that $|H_{\text{an}}| = |H_{\text{num}}|$ to within 10^{-13} (machine precision) for all tested $\varepsilon \in \{0.05, 0.1, 0.2, 0.4, 0.8\}$, with the boundary crossings correctly detected ($n_{\text{cross}} = 2$ per loop); (d) verifies the boundary reset term is linear in a_1 with slope ε (the loop width), as predicted by $\alpha(p_+) - \alpha(p_-) = a_1 \cdot \varepsilon$. Outputs: `two_cat_gluings_stratified.{png, csv, txt}`.

Remark 3.8 (Non-abelian extension to $G_C = SO(3)$, the $n = 4$ prototype's policy fiber). The abelian verification of Remark 3.7 is extended to the *non-abelian* $G_C = SO(3)$ setting of the $n = 4$ prototype's policy fiber (Definition 2.2) by script `two_cat_gluings_so3.py`. The $\mathfrak{so}(3)$ -valued connection on each stratum is $A_{\pm} = (F/2)(x dy - y dx) T_z$ where T_z is the $\mathfrak{so}(3)$ generator of rotations about the z -axis, and the boundary transition $g_{+-}(x, 0) = \exp(\alpha(x) T_y)$ is a rotation about the y -axis by $\alpha(x) = a_1 x$. The commutator $[T_y, T_z] = T_x \neq 0$ makes the matrix product in the piecewise holonomy formula (8) genuinely order-dependent: the boundary reset $R_b = \alpha(p_*) T_y$ lies in a different $\mathfrak{so}(3)$ direction than the stratum curvature FT_z . The verification: (a) computes the analytic piecewise holonomy $H_{\text{an}} \in SO(3)$ as the ordered matrix product of stratum-holonomies $\exp(-(F/2) \int (x dy - y dx) T_z ds)$ and boundary transitions $\exp(\pm \alpha(p_k) T_y)$ in the traversal order; (b) computes the numerical $SO(3)$ holonomy by segment-by-segment matrix-exponential parallel transport $H_{\text{num}} = \prod_k \exp(-A_{\text{mid}} ds) \cdot \prod_k g(p_k)^{\pm 1}$; (c) confirms $\|H_{\text{num}} - H_{\text{an}}\|_F < 10^{-10}$ for all tested $\varepsilon \in \{0.05, 0.1, 0.2, 0.4, 0.8\}$, with $\det(H_{\text{num}}) = +1$ and $H_{\text{num}}^{\top} H_{\text{num}} = I_3$ (genuine special-orthogonal matrices); (d) recovers the abelian limit $\alpha = 0 \Rightarrow H = \exp(-F\varepsilon^2 T_z)$ (single z -rotation by $F\varepsilon^2$); (e) detects the non-abelian mixing via the off-diagonal entry $H_{\text{num}}[0, 2]$ (the $x \rightarrow z$ block), which is zero at $a_1 = 0$ and grows linearly in a_1 (slope $\approx \varepsilon \cdot F\varepsilon^2/2$ in the small-loop regime), demonstrating that the boundary y -rotation rotates the holonomy matrix out of the stratum-curvature T_z direction. The non-abelian verification thus confirms Theorem 3.6 in the regime matching the $n = 4$ prototype. Outputs: `two_cat_gluings_so3.{png, csv, txt}`.

Remark 3.9 (Third verification regime: non-abelian $G_C = CO(3)$ for endogenous reversibility). The piecewise holonomy formula is further verified in the *third* structure-group regime of Defini-

tion 2.2, $G_C = CO(3) = \mathbb{R}_+ \ltimes SO(3)$, corresponding to endogenous reversibility in the policy-fiber hierarchy. The Lie algebra $\mathfrak{co}(3) = \mathbb{R}S \oplus \mathfrak{so}(3)$ decomposes as a direct sum of the *central* scaling direction $S = I_3$ ($[S, \cdot] = 0$) and the non-abelian rotation part $\mathfrak{so}(3)$. The two-stratum base and stratum connection match the $SO(3)$ setup, but the boundary transition is now a $CO(3)$ -valued map

$$g_{+-}(x, 0) = \exp(\alpha(x) T_y + \beta(x) S) = R_y(\alpha(x)) \cdot \Lambda(\beta(x)), \quad (10)$$

where $\Lambda(\beta) = e^\beta I_3$ is uniform scaling by e^β , $\alpha(x) = a_1 x$ is the rotation amplitude (as in the $SO(3)$ test), and $\beta(x) = b_1 x$ is the scaling amplitude (new in $CO(3)$). The boundary reset thus combines a T_y -direction rotation with an S -direction scaling; since S is central, the scaling factors compose multiplicatively (abelian), while the rotations compose non-commutatively (as in the $SO(3)$ regime). The $CO(3)$ holonomy is consequently a product of non-abelian rotation matrices and abelian scaling factors, with $\det(H) = \lambda_{\text{total}}^3 > 0$ and $H^\top H = \lambda_{\text{total}}^2 \cdot I_3$ (conformal orthogonality: $CO(3)$ preserves angles up to uniform scaling). Script `two_cat_gluing_co3.py` verifies: (a) $\|H_{\text{num}} - H_{\text{an}}\|_F < 10^{-10}$ for all tested $\varepsilon \in \{0.05, 0.1, 0.2, 0.4, 0.8\}$; (b) $\det(H_{\text{num}}) = \lambda_{\text{total}}^3$ to within 10^{-9} ; (c) $\|H_{\text{num}}^\top H_{\text{num}} - \lambda_{\text{total}}^2 I_3\|_F < 10^{-9}$; (d) the abelian limit ($\alpha = \beta = 0$) recovers $H = \exp(-F\varepsilon^2 T_z)$ (same as the $SO(3)$ abelian limit, since the scaling direction contributes nothing when the boundary transition vanishes); (e) the pure-scaling limit ($\alpha = 0$, $\beta \neq 0$) recovers $H = \exp(-F\varepsilon^2 T_z) \cdot \Lambda(b_1 \varepsilon)$ to machine precision (rotation holonomy composed with uniform scaling). The $CO(3)$ verification completes the confirmation of Theorem 3.6 in *all three* structure-group regimes: abelian $U(1)$, non-abelian $SO(3)$, and endogenous-reversibility $CO(3)$. Outputs: `two_cat_gluing_co3.{png, csv, txt}`.

Remark 3.10 (Topologically distinct loop: triangular geometry). To verify that the piecewise holonomy formula is *geometry-independent* (and not an artifact of the rectangular loop shape), script `two_cat_gluing_so3_triangular.py` repeats the $SO(3)$ verification on a *triangular* loop with vertices $V_1 = (0, +\varepsilon)$, $V_2 = (+\varepsilon, -\varepsilon)$, $V_3 = (-\varepsilon, -\varepsilon)$ (traversed counterclockwise). The triangle is topologically distinct from the rectangle: 3 vertices and 3 edges instead of 4 and 4; boundary crossings at *non-perpendicular* angles (the triangle's edges cross the x -axis at oblique angles, not vertically); an *asymmetric* distribution of stratum pieces (3 pieces in S_- , 2 in S_+ ; the rectangular loop splits its 4 pieces symmetrically with 2 per stratum). The triangle has signed area $2\varepsilon^2$ (twice the rectangular loop's ε^2 at the same scale). The boundary crossings occur at $p_1 = (\varepsilon/2, 0)$ on edge $V_1 \rightarrow V_2$ and $p_3 = (-\varepsilon/2, 0)$ on edge $V_3 \rightarrow V_1$, giving 5 stratum pieces in total: edge 1a ($V_1 \rightarrow p_1$ in S_+), edge 1b ($p_1 \rightarrow V_2$ in S_-), edge 2 ($V_2 \rightarrow V_3$ entirely in S_-), edge 3a ($V_3 \rightarrow p_3$ in S_-), edge 3b ($p_3 \rightarrow V_1$ in S_+). The piecewise formula gives $H = \text{Hol}_{S_+}(3b) \cdot g(p_3)^{-1} \cdot \text{Hol}_{S_-}(3a) \cdot \text{Hol}_{S_-}(2) \cdot \text{Hol}_{S_-}(1b) \cdot g(p_1) \cdot \text{Hol}_{S_+}(1a)$, with each Hol a T_z -rotation by the line integral of $A = (F/2)(x dy - y dx)$. Verification: (a) $\|H_{\text{num}} - H_{\text{an}}\|_F < 10^{-10}$ for all $\varepsilon \in \{0.05, 0.1, 0.2, 0.4, 0.8\}$ (machine precision); (b) $\det(H_{\text{num}}) = +1$ and $H_{\text{num}}^\top H_{\text{num}} = I$ (genuine $SO(3)$); (c) the abelian limit ($\alpha = 0$) recovers $H = \exp(-F \cdot \text{Area}(\triangle) T_z) = \exp(-2F\varepsilon^2 T_z)$ (the triangle's larger area produces $2\times$ the rectangular loop's rotation angle, exactly matching the area formula $F \cdot \text{Area}$); (d) the non-abelian feature is detected ($H_{\text{num}}[0, 2] = 0$ at $a_1 = 0$, grows linearly in a_1). The triangular-loop verification confirms that the piecewise holonomy formula holds for arbitrary loop geometry, not just axis-aligned rectangles. Outputs: `two_cat_gluing_so3_triangular.{png, csv, txt}`.

Remark 3.11 (Pentagon, circle, and ellipse loop topologies). To further validate geometry-independence, script `two_cat_gluing_so3_topology.py` repeats the $SO(3)$ verification on three additional loop topologies that are topologically and geometrically distinct from both the rectangular and triangular loops:

- *Pentagon* (5-vertex regular polygon, inscribed in a circle of radius ε): a polygonal loop with higher vertex count than the triangle (3) or rectangle (4). Pentagon area: $(5/2)\varepsilon^2 \sin(2\pi/5) \approx 2.38\varepsilon^2$ (at $\varepsilon = 1$). Pentagon boundary crossings: 2 (the polygon spans both strata), but with 5 edges contributing to the piecewise decomposition.

- *Circle* (smooth non-polygonal curve, 200 equal-angle sample points starting at angle π/n to avoid placing a sample exactly on the boundary): a smooth curve, the analytic limit of polygonal refinement. Circle area: $\pi\varepsilon^2$.
- *Ellipse* (smooth anisotropic curve, 200 equal-angle sample points with semi-major axis $a = \varepsilon$ and semi-minor axis $b = \varepsilon/2$): a loop with broken rotational symmetry, testing that the piecewise formula extends to anisotropic loop geometry. Ellipse area: $\pi\varepsilon \cdot (\varepsilon/2) = \pi\varepsilon^2/2$.

For each topology, the verification is run on the same two-stratum base as the rectangular ($U(1)$ and $SO(3)$) and triangular ($SO(3)$) verifications: $F_+ = F_- = 2$, $\alpha(x) = 1.5x$, $\varepsilon \in \{0.05, 0.1, 0.2, 0.4, 0.8\}$. The numerical holonomy is computed by the same segment-by-segment matrix-exponential parallel transport, with 2000 sample points (numerical) and 8000 sample points (analytic reference). For all three topologies: (a) $\|H_{\text{num}} - H_{\text{an}}\|_F < 10^{-10}$ for all ε (machine precision); (b) $\det(H_{\text{num}}) = +1$ and $H_{\text{num}}^\top H_{\text{num}} = I$ (genuine $SO(3)$); (c) $n_{\text{crossings}} = 2$ per loop (entry + exit); (d) the abelian limit ($\alpha = 0$) recovers $H = \exp(-F \cdot \text{Area}(\text{loop}) \cdot T_z)$ within machine precision for the corresponding area formula (pentagon $(5/2)\varepsilon^2 \sin(2\pi/5)$, circle $\pi\varepsilon^2$, ellipse $\pi\varepsilon^2/2$); (e) the non-abelian feature is detected ($H_{\text{num}}[0, 2] = 0$ at $a_1 = 0$, grows linearly in a_1 in all three topologies). Combined with the rectangular ($U(1)$, $SO(3)$), $CO(3)$, and triangular verifications, the piecewise holonomy formula (Theorem 3.6) is now verified on *five* distinct loop shapes (rectangle, triangle, pentagon, circle, ellipse), confirming that the formula is *geometry-independent*: it holds for any loop shape, polygonal or smooth, isotropic or anisotropic. Outputs: `two_cat_gluings_so3_topology.{png, csv, txt}`.

Remark 3.12 (Pathwise viability versus endpoint viability). A small loop with bounded endpoint holonomy does not imply that intermediate points of the loop are viable: the endpoint bound is a weaker statement than pathwise viability. To control intermediate points, the calibration protocol of the abelian $n = 3$ prototype applies a Nagumo inward-pointing condition along the path, equivalently a viability tube of radius ε around the loop. A loop with bounded endpoint holonomy but intermediate viability violation is a false negative for the holonomy prediction; the tube check excludes it. The tube is operationalized by verifying that the entire ε -neighborhood of $\gamma_a([0, 1])$ lies inside the closed viability kernel $\{x : V(x) \geq V_{\max} - \varepsilon\}$; this is a binary check that produces a separate falsifiable verdict on the pathwise viability content of the holonomy prediction.

3.2 Rigorous Fisher-minimal transport law on constant-active-set strata

On a constant-active-set stratum of SAVGS, the active viability constraints are represented locally by $c(\theta, p) = 0$ with $J_p = D_p c$ of full row rank and $J_\theta = D_\theta c$. For an environmental velocity $\dot{\theta}$, the Fisher-minimal policy velocity is the unique solution of $J_p v + J_\theta \dot{\theta} = 0$ minimizing $\frac{1}{2} v^\top G(p) v$, where $G(p)$ is the Fisher information metric on the open simplex Δ^{n-1} .

Proposition 3.13 (Fisher-minimal horizontal lift (KKT form)). *On a constant-active-set stratum where J_p has full row rank, the unique Fisher-minimal policy velocity is*

$$v^*(\theta, p, \dot{\theta}) = -G(p)^{-1} J_p^\top (J_p G(p)^{-1} J_p^\top)^{-1} J_\theta \dot{\theta}, \quad (11)$$

which defines the connection 1-form $A_i = G^{-1} J_p^\top (J_p G^{-1} J_p^\top)^{-1} \partial_i c$ on the stratum, satisfying $\dot{p} = -A_i(\theta, p) \dot{\theta}^i$. The horizontal lift $H_i = \partial_{\theta^i} - A_i$ is a genuine smooth connection on each constant-rank stratum, with curvature $F_{ij} = \partial_i A_j - \partial_j A_i + [A_i, A_j]$.

Proof. The constrained minimum of $\frac{1}{2} v^\top G v$ subject to $J_p v + J_\theta \dot{\theta} = 0$ is found by the method of Lagrange multipliers. Setting $\nabla_v [\frac{1}{2} v^\top G v + \mu^\top (J_p v + J_\theta \dot{\theta})] = G v + J_p^\top \mu = 0$ gives $v = -G^{-1} J_p^\top \mu$. Substituting into the constraint: $-J_p G^{-1} J_p^\top \mu + J_\theta \dot{\theta} = 0$, so $\mu = (J_p G^{-1} J_p^\top)^{-1} J_\theta \dot{\theta}$. The full-row-rank hypothesis on J_p guarantees the inverse exists. The horizontal subspace $H_{(\theta, p)} = \ker(J_p) \cap$

$T_p \Delta^{n-1}$ varies smoothly in (θ, p) on the stratum (constant-rank hypothesis), so the resulting 1-form A_i is C^2 ; the curvature F_{ij} is then the standard Ehresmann-curvature expression by direct computation [Kobayashi and Nomizu \(1969\)](#). $\square$

Remark 3.14 (Projected differential inclusion at constraint-switching boundaries). At a constraint-switching boundary—where an inequality enters or leaves the active set—the full-row-rank hypothesis of [Proposition 3.13](#) fails: the rank of J_p jumps, and the smooth connection A_i is no longer defined. The correct replacement is the projected differential inclusion

$$\dot{p} \in \arg \min_{v \in T_{\mathcal{K}}(\theta, p; \dot{\theta})} \frac{1}{2} \|v - v_0\|_{F, p}^2, \quad (12)$$

where $T_{\mathcal{K}}(\theta, p; \dot{\theta})$ is the Bouligand contingent cone to the viability kernel $\mathcal{K}$ at (θ, p) in the direction $\dot{\theta}$ and v_0 is the unconstrained Fisher-minimal velocity. This is viability dynamics, not ordinary smooth-connection theory; the SAVGS 2-categorical span of [Definition 3.1](#) glues the smooth-stratum connections across the boundary via the matching condition on the boundary 2-cell.

Theorem 3.15 (Small-loop viability-holonomy). *Let $(B, E, \mathcal{P}, \varepsilon, \Gamma)$ be a SAVGS as in [Definition 3.1](#), on a constant-active-set stratum where [Proposition 3.13](#) applies and the connection 1-form A_i is C^2 . Let $\gamma_\varepsilon : [0, 1] \rightarrow B$ be a square loop of side ε in the (θ^1, θ^2) -plane at θ_0 , contained entirely in a single stratum, and let $h_\alpha(\theta, x)$ be a C^2 viability function with $h_\alpha(\theta_0, x_0) > 0$. Then the parallel transport of the policy around γ_ε satisfies*

$$p_{\text{end}} - p_0 = \varepsilon^2 F_{12}(\theta_0, p_0) + O(\varepsilon^3), \quad (13)$$

up to the orientation convention for F_{12} . The induced change in the viability margin is

$$h_\alpha(x_{\text{end}}) - h_\alpha(x_0) = \varepsilon^2 D_p h_\alpha(F_{12}(\theta_0, p_0)) + O(\varepsilon^3) + \Delta_\alpha^{\text{other}}, \quad (14)$$

where $\Delta_\alpha^{\text{other}}$ collects all non-returning energy, integrity, repair, and environmental contributions along the loop. In particular, the endpoint remains viable whenever

$$h_\alpha(x_0) + \varepsilon^2 D_p h_\alpha(F_{12}) - C_\alpha \varepsilon^3 - |\Delta_\alpha^{\text{other}}| > 0 \quad \text{for every } \alpha, \quad (15)$$

where C_α depends on the C^2 bounds of h_α and A_i .

Proof. The expansion (13) is the standard non-abelian Stokes formula for the parallel transport around an infinitesimal loop (see, e.g., [Kobayashi and Nomizu \(1969\)](#); [Misner et al. \(1973\)](#)): writing the path-ordered exponential $\mathcal{P} \exp(-\oint A) = 1 - \iint_S F_{12} d\theta^1 d\theta^2 + O(\varepsilon^3)$ on a square S of side ε gives $p_{\text{end}} - p_0 = -\varepsilon^2 F_{12} + O(\varepsilon^3)$ (the sign convention is fixed by the orientation of A_i and is absorbed into F_{12}). Equation (14) follows by Taylor expansion of h_α at x_0 to first order in $p_{\text{end}} - p_0$ and second order in ε , with $\Delta_\alpha^{\text{other}}$ collecting the integral of $D_\theta h_\alpha \dot{\theta}$ along the loop plus the $O(\varepsilon^3)$ remainder. The sufficient condition (15) is then the requirement that the worst-case $O(\varepsilon^3)$ remainder plus the non-returning contribution not exceed the initial margin minus the leading-order curvature contribution. $\square$

Proposition 3.16 (Vulnerability bound). *Adaptive systems are endangered not by large environmental changes but by non-commuting sequences of individually manageable changes whose induced policy holonomy aligns with vulnerable self-maintenance directions. The upper bound on vulnerability is the viability-weighted curvature κ_V of [Definition 2.10](#) on a G_C -structured stratified connection; the lower bound is 0, attained when $\kappa_V \equiv 0$ (flat connection).*

Proof. On a constant-active-set stratum, integrate [Theorem 3.15](#) along the loop: the endpoint margin change is bounded above in magnitude by $\varepsilon^2 D_p h_\alpha(F_{12}) \leq \varepsilon^2 \kappa_V$ per unit loop area, with the positive part in [Definition 2.10](#) discarding viability-restoring curvature; summing over the loop gives the upper bound. The flat connection $F \equiv 0$ gives $p_{\text{end}} = p_0$ to $O(\varepsilon^3)$, hence zero holonomy-induced vulnerability and the lower bound. $\square$

Remark 3.17 (What Theorem 3.15 does and does not claim). Theorem 3.15 bounds the *endpoint* viability margin by the leading-order curvature contribution plus an $O(\varepsilon^3)$ remainder. It does *not* bound the *pathwise* viability margin: a loop with bounded endpoint holonomy may transit non-viable regions. Pathwise viability requires the additional Nagumo inward-pointing condition of Remark 3.12—equivalently a viability tube of radius ε around $\gamma_a([0, 1])$ contained in the closed viability kernel. Curvature is therefore neither necessary nor sufficient for survival in isolation; the bidirectional “equivalence” is replaced by the unidirectional upper bound “vulnerability $\leq \kappa_V$ ” of Proposition 3.16, with the trivial lower bound 0 attained when $\kappa_V \equiv 0$.

3.3 Square-root embedding and gauge invariance

Proposition 3.18 (Square-root embedding). *The square-root embedding $\psi_a = 2\sqrt{p_a}$ places the open simplex Δ^{n-1} on the positive orthant of the unit sphere $S^{n-1} \subset \mathbb{R}^n$. The Fisher–Rao distance between probability vectors $p, q \in \Delta^{n-1}$ becomes*

$$d_{\text{FR}}(p, q) = 2 \arccos\left(\sum_a \sqrt{p_a q_a}\right), \quad (16)$$

the standard round metric on the positive orthant. All geometric quantities computed in this embedding are gauge-invariant; the logit coordinates of prior formulations are recovered as a stereographic projection.

Proof. Direct computation of the Fisher–Rao metric tensor in the square-root embedding [Bhattacharyya \(1943\)](#); [Amari and Nagaoka \(2000\)](#). The induced metric on the positive orthant of S^{n-1} is the standard round metric; the stereographic projection onto the tangent plane at any fixed reference point gives the logit coordinates. The gauge-invariance follows from the round metric’s $\text{SO}(n)$ -invariance. $\square$

3.4 Intervention-based autopoiesis closure test

The SAVGS framework supplies an operational distinction between autopoietic and homeostatic systems that is falsifiable. The audit’s core distinction is between *homeostasis* (the agent remains inside a prescribed viable set, possibly via externally supplied maintenance) and *autopoiesis* (the agent regenerates the sensing, control, repair, and boundary-maintenance machinery required to remain there). The maintenance graph $\Gamma = (M, R, E)$ of Definition 3.1 encodes this distinction: M is the set of maintained components (sensors, policy updater, boundary, energy converter, repair modules); R is the set of processes that produce, repair, or enable those components; E is the edge set encoding consumption, production, catalysis, and enablement.

Definition 3.19 (Operational closure criteria). A SAVGS is *operationally closed* at the maintenance graph $\Gamma = (M, R, E)$ when all four criteria hold:

1. Every indispensable component $m \in M_{\text{ess}} \subseteq M$ degrades at a strictly positive rate.
2. Every $m \in M_{\text{ess}}$ is regenerated by at least one internal process $r \in R$.
3. Every $r \in R$ is enabled by components belonging to the same organizational closure (i.e., the enablers of r lie in M_{ess}).
4. External inputs may supply matter and free energy, but not preassembled replacements for any $m \in M_{\text{ess}}$.

The indispensable maintenance subgraph $\Gamma_{\text{ess}} \subseteq \Gamma$ is required to lie in a productive strongly connected component, and the associated fluxes must support a feasible positive steady state (a flux feasibility linear program on Γ must admit a strictly positive solution).

Definition 3.20 (Autopoiesis closure test — dynamical closure with viability threshold). The SAVGS carries concentration dynamics $\dot{x} = Nv(x) - Dx + u_{\text{food}}$ on the maintenance species $x \in \mathbb{R}_{\geq 0}^M$ (stoichiometry N , reaction rates v , degradation D , allowed external food input u_{food} ; cf. Definition 3.19). Fix a strictly positive *viability concentration threshold* $x_{\text{thresh}} > 0$ and a recovery horizon T . For each purportedly self-maintained component $m_j \in M_{\text{ess}}$:

- (i) *Knockout*. Set the internal repair flux of m_j to zero: $r_j := 0$. Keep the external food supply u_{food} unchanged (no preassembled replacement for m_j is supplied).
- (ii) *Drift under dynamics*. Integrate the concentration dynamics $\dot{x} = Nv(x) - Dx + u_{\text{food}}$ forward for T steps from the pre-knockout steady state x^* .
- (iii) *Recovery above threshold*. Record the post-recovery concentration $\text{rec}_j(T) := x_j(T)$. The component m_j is *recovered* iff $\text{rec}_j(T) > x_{\text{thresh}}$ (not merely nonzero — the concentration must exceed the declared viability threshold, which removes the trivial reinsertion failure mode of binary node-reappearance tests).
- (iv) *Downstream cascade*. Every $m_{j'}$ whose enabling chain depends on m_j must also drop below x_{thresh} during step (ii) and recover above x_{thresh} in step (iii) (predicted downstream effect).
- (v) *Restoration control*. Restore the repair pathway ($r_j := r_j^{\text{pre-knockout}}$) and verify that all knocked-down and downstream concentrations return to x^* within T steps (recovery control).

A component m_j is *causally internal* to the organizational closure iff steps (i)–(v) hold as stated. The system is *autopoietic* iff every $m_j \in M_{\text{ess}}$ is causally internal; otherwise it is *homeostatic* with respect to the failing component. The default viability threshold used in the operational network benchmarks is $x_{\text{thresh}} = 0.1$ in normalized concentration units; the threshold is a declared operational parameter, not a fitted one, and may be varied in the robustness sweeps documented in the project repository.

Remark 3.21 (Why the test is non-circular and dynamical). The closure test of Definition 3.20 prevents “autopoiesis” from being satisfied by a simulator that silently repairs the agent from outside: the external intervention in step (i) is the only externally imposed change; the recovery in step (iii) must be produced by the endogenous regeneration rules R alone, with concentrations crossing the declared viability threshold $x_{\text{thresh}} > 0$ rather than merely registering a nonzero symbolic presence. A homeostatic system that depends on a hidden external repair mechanism fails the test on step (iii): the knocked-down component’s concentration does not exceed x_{thresh} within T steps. The threshold requirement removes the principal circularity of the older node-reappearance formulation, where any nonzero reinsertion (including a trivial simulator-level reset) would pass. The test is falsifiable in three independent directions: (a) if a non-food maintenance node m_j fails to recover above x_{thresh} in step (iii), the system is homeostatic with respect to m_j ; (b) if knockout of an essential catalytic hub does not cause the predicted downstream collapse in step (iv), the maintenance graph Γ is mis-specified; (c) if restoration of the repair pathway fails to recover x^* in step (v), the steady-state viability is unstable and the autopoiesis verdict is unreliable.

Definition 3.22 (Closure-aware viability curvature). Restricting the viability covectors of κ_V (Definition 2.10) to the repair-machinery margins $h_j^{\text{repair}} = r_j - r_{j,\min}$ gives the *closure-aware viability curvature*

$$\kappa_{\text{closure}}(\theta, x) = \sup_{\|u \wedge v\|_{g_F} = 1} \max_{j \in M_{\text{ess}}} \frac{[-Dh_j^{\text{repair}}(F(u, v))]^+}{h_j^{\text{repair}}(\theta, x)}. \quad (17)$$

The quantity κ_{closure} distinguishes three failure modes: (i) harmless behavioral hysteresis (the curvature projects onto non-maintenance directions); (ii) metabolic viability erosion ($\kappa_V > 0$ but

$\kappa_{\text{closure}} = 0$); (iii) autopoietic failure proper ($\kappa_{\text{closure}} > 0$, the curvature erodes the machinery that makes future adaptation possible). Only the third is specifically autopoietic failure.

Remark 3.23 (The operational, data-side curvature is not defined here). A predecessor of this framework also carried an *FBA-operational* viability-weighted curvature: an indicator-masked, gene-knockout flux-rerouting statistic whose indicator imports the model’s own essentiality call into the predictor, and whose adoption was justified post hoc by the strengthening it produced in two empirical studies. That object is *not* carried in this paper: it is superseded, in the application paper (X, 2026), by a measure-theoretic sensitivity metric derived from the atomic curvature measure of parametric FBA — a construction whose selection-rule dependence is exposed and audited rather than masked (tie-break robustness across five lexicographic selection rules), and whose empirical validation carries an explicit protein-layer null. The categorical framework of the present paper is independent of that choice: Definition 2.10 is the geometric object; any faithful operationalization of it is welcome.

4 Bregman-Divergence Noether Correspondence

A Noether-type correspondence for the Lagrangian

$$L = \mathbb{E}[d] + \lambda I \quad (18)$$

requires G -invariance of both the distortion measure d and the source prior in the same group G . Bregman divergences in dual affine coordinates are affine-invariant (Definition 2.8), which supplies both.

Proposition 4.1 (Bregman–Hessian Noether correspondence). *Let $Q \subset \mathbb{R}^r$ be open convex, $\phi \in C^3$ strictly convex, and $g_\phi = \nabla^2 \phi$ the Hessian metric on Q . Let $L(q, \dot{q}) = \frac{1}{2} g_\phi(q)(\dot{q}, \dot{q}) - U(q)$ be the geodesic Lagrangian with potential $U \in C^2$, and let $S[q] = \int_a^b L(q, \dot{q}) dt$ be the action on the configuration space of smooth curves $q : [a, b] \rightarrow Q$. Let ξ be a complete affine vector field on Q whose flow $g_t = \text{flow}_t(\xi)$ preserves the Hessian metric,*

$$\mathcal{L}_\xi g_\phi = 0, \quad (19)$$

and preserves the potential, $\xi(U) = 0$. Then S is invariant under the lifted flow $\tilde{g}_t \cdot (q, \dot{q}) = (g_t \cdot q, Dg_t \cdot \dot{q})$, and by Noether’s theorem Noether (1918); Olver (1993) the momentum

$$J_\xi(q, \dot{q}) = g_\phi(q)(\dot{q}, \xi(q)) = \langle \nabla \phi(q), \xi(q) \rangle \cdot \dot{q}_{\text{dual}} \quad (20)$$

is conserved along Euler–Lagrange trajectories of S : $\frac{d}{dt} J_\xi(q(t), \dot{q}(t)) = 0$ on-shell.

Proof. The Hessian-metric preservation (19) gives $\mathcal{L}_\xi L = \xi(U) = 0$ (the kinetic term is invariant by $\mathcal{L}_\xi g_\phi = 0$, and the potential term is invariant by $\xi(U) = 0$); hence $S[\tilde{g}_t \cdot q] = S[q]$ for all t , and $\tilde{g}_t$ is a continuous variational symmetry of S . By Noether’s theorem in its standard field-theoretic form Olver (1993), the one-parameter symmetry yields the conserved momentum $J_\xi(q, \dot{q}) = \partial L / \partial \dot{q} \cdot \xi(q) = g_\phi(q)(\dot{q}, \xi(q))$. The dual-coordinate form follows by identifying $\partial L / \partial \dot{q} = \nabla \phi(q) \cdot \dot{q}_{\text{dual}}$ via the dual-coordinate structure of Bregman divergences (Definition 2.8). $\square$

Corollary 4.2 (Affine Bregman invariance; sufficient condition). *A one-parameter affine flow g_t on Q that satisfies $\phi \circ g_t = \phi + \ell_t$ for an affine function ℓ_t preserves the Hessian metric $g_\phi = \nabla^2 \phi$ (since $\nabla^2(\phi \circ g_t) = \nabla^2 \phi$ on the affine transformation g_t by the chain rule, and $\nabla^2 \ell_t = 0$). Therefore the hypothesis (19) is satisfied, and the Bregman divergence D_ϕ is invariant in the paired primal-dual sense required by Proposition 4.1.*

Proposition 4.3 (Falsifiable precondition check). *The Bregman–Hessian Noether correspondence has a falsifiable precondition check: verify (i) that ϕ is C^3 strictly convex (positive-definite Hessian $g_\phi = \nabla^2\phi \succ 0$), (ii) that the vector field ξ generates an affine Hessian isometry ($\mathcal{L}_\xi g_\phi = 0$, equivalently $\mathcal{L}_\xi(\nabla^2\phi) = 0$), and (iii) that $\xi(U) = 0$. If all three conditions hold, the correspondence applies; failure of any condition refutes the theorem at that point.*

Proof. Operational verification: (i) compute the Hessian $\nabla^2\phi$ and verify positive definiteness; (ii) compute the Lie derivative $\mathcal{L}_\xi(\nabla^2\phi)$ via Cartan’s formula $\mathcal{L}_\xi T = [i_\xi, d]T$ and verify it vanishes; (iii) compute the directional derivative $\xi(U)$ and verify it is zero. All three checks are mechanical and produce a binary verdict. $\square$

Remark 4.4 (Application to gauge invariance of κ_V). If the policy connection and viability margins of Definition 2.10 are constructed from a Bregman generator ϕ and the gauge group acts by affine Hessian isometries preserving the margins, the curvature F and the viability-weighted curvature κ_V are gauge-covariant (and gauge-invariant when the margins themselves are gauge-invariant). This is the rigorous replacement for the earlier Bregman-Noether proposition: not divergence invariance alone (which is false in general under arbitrary affine transformations holding ϕ fixed), but affine isometries of the Hessian metric yield conserved momenta by Noether’s theorem.

5 The Seven-Claim Falsification Hierarchy

Definition 5.1 (Falsification hierarchy). The viability-weighted curvature κ_V of Definition 2.10 generates a seven-claim falsification hierarchy on the abelian $n = 3$ prototype. Each claim is a specific quantitative prediction; failure of any prediction refutes the corresponding level of the hierarchy.

- A.** (Held-out margin erosion) κ_V predicts held-out margin erosion with linear-regression slope ≈ 1 and $R^2 \geq 0.9$.
- B.** (Orientation reversal amplitude) The holonomy $H(a) = \pi a^2$ reverses orientation when $|H| > \pi$, i.e. at $a_{\text{rev}} = 1$.
- C.** (Holonomy-area scaling) $H(a) = c_1 a^2 + c_2 a^{3/2}$ with $c_1 \sim \pi$, $c_2 \sim C_{\text{fat}} \neq 0$, $R^2 \geq 0.95$.
- D.** (Repeated-loop fatigue) Two-sided bound: the *conservative sufficient safety condition* $\sum_{k=1}^K (a_k (\kappa_V)_k + C_{\text{fat}} a_k^{3/2} + \eta_k) < 1$ guarantees survival through K loops, and the *failure condition* $\sum_{k=1}^K (a_k \kappa_V(a_k) + C_{\text{fat}} a_k^{3/2} + \eta_k) > 1$ predicts fatigue failure; $V_{\text{max},K} = \prod_k (1 - F_k) < e^{-1}$ at K_{obs} ; relative error $< 15\%$.
- E.** (Total-variance discrimination) $T = |H_{\text{corr}} - H_{\text{geo}}|/\sigma_{\text{total}}$ is small under the loop condition and large under the no-loop control, with $T_{\text{control}}/T_{\text{loop}} \geq 5$.
- F.** (Non-abelian commuting control) At $n \geq 4$, the structure group G_C has non-abelian $\mathfrak{so}(3)$ when $G_C = SO(3)$; same-axis rotations commute (machine precision), distinct-plane rotations do not (nonzero non-abelian signature).
- G.** (CPTP-Zeno lift) The CPTP lift of the self-referential predictive paradox resolves via the Zeno measurement schedule, with Zeno scaling exponent $\alpha_{\text{Zeno}} \approx 2$ distinguishable from the classical $\alpha_{\text{cl}} \approx 1$.

The hierarchy is cumulative: failure at level k renders the subsequent levels operationally meaningless but does not refute the underlying categorical construction (Theorem 6.4).

Remark 5.2 (Recommended experimental ordering). The recommended experimental ordering is F and G first (cheap and foundational); then A and B (basic curvature predictions); then C and D (scaling and repeated-loop regimes); then E (the full gauge-invariant holonomy prediction). This ordering implements the principle that cheap foundational tests should precede expensive derivative tests. If F or G fails, the derivative tests are unnecessary.

6 The Single Composition Theorem

Construction 6.1 (Seven-optic endofunctor). Let $\mathcal{C}$ be a category with finite limits. Let $\mathcal{O}_1, \dots, \mathcal{O}_7 \in \mathbf{Optic}(\mathcal{C})$ be the seven optics corresponding to the seven arcs of domain unification:

1. $\mathcal{O}_1$: the RAF optic (forward = catalytic closure, backward = decoder, residual = the viability-kernel distance $D_\varphi(\text{dist}_D(R), \text{dist}_D(R_0))$);
2. $\mathcal{O}_2$: the RPSI optic (forward = CPTP predictive update, backward = posterior decoder, residual = Holevo information $I(\rho_{\text{out}}; \hat{\rho}_{\text{in}})$);
3. $\mathcal{O}_3$: the IFS optic (forward = iterated-function-system contraction [Hutchinson \(1981\)](#); [Barnsley \(1988\)](#), backward = decoder, residual = contraction quotient c_i of f_i);
4. $\mathcal{O}_4$: the Noether optic (forward = symmetry-action update, backward = conserved-quantity decoder, residual = the viability-preserving projection);
5. $\mathcal{O}_5$: the perturbation optic (forward = perturbed Blahut–Arimoto update [Blahut \(1972\)](#); [Arimoto \(1972\)](#), backward = perturbation decoder, residual = perturbation tensor $\partial d/\partial p$);
6. $\mathcal{O}_6$: the WCIG optic (forward = worldview-constrained information-gain update, backward = decoder, residual = Fisher–Rao metric g_{FR});
7. $\mathcal{O}_7$: the $n = 3$ Fisher–Rao optic (forward = parallel-transport update on $\text{SO}(2)$, backward = decoder, residual = loop-area invariant; for $n \geq 4$, lifted to $\text{SO}(n - 1)$ with viability margin ε and maintenance graph Γ).

The compatibility condition $A_i = M_{i+1}$ holds by construction.

Remark 6.2 (Optic decompositions of the arcs). The forward components are the encoding operations of each arc; the backward components are the decoding operations; the residuals are the information that flows from the forward pass to the backward pass. The IFS optic ($\mathcal{O}_3$) is a pure coalgebra [Rutten \(2000\)](#) on the category of compact metric spaces, while the perturbation optic ($\mathcal{O}_5$) is a coalgebra with residual, the prototypical BA structure. The compatibility condition is satisfied by construction: each arc’s forward action is the next arc’s backward state. Two arcs deserve comment. The RPSI arc’s forward component is a CPTP channel, not a deterministic function, because the RPSI self-reference paradox requires the quantum lift of Section 7. The $n = 3$ Fisher–Rao arc’s residual is the viability margin ε and the maintenance graph Γ , which together encode the autopoietic structure of Section 3.

Table 1: Seven-optic composition. Each arc is decomposed as an optic (M_i, C_i, R_i) in $\mathbf{Optic}(\mathcal{C})$ with forward φ_i (encoder), backward ρ_i (decoder), and residual Res_i (information flowing forward to backward). The compatibility condition $A_i = M_{i+1}$ holds by construction. The total residual of T is the product $\text{Res}_1 \times \dots \times \text{Res}_7$ in $\mathcal{C}$.

Arc	Forward φ_i (encoder)	Backward ρ_i (decoder)	Residual Res_i
$\mathcal{O}_1$ RAF	catalytic closure	RAF transition decoder	$D_\varphi(\text{dist}_D(R), \text{dist}_D(R_0))$
$\mathcal{O}_2$ RPSI	CPTP predictive update	posterior decoder	$I(\rho_{\text{out}}; \hat{\rho}_{\text{in}})$ (Holevo)
$\mathcal{O}_3$ IFS	Hutchinson contraction	deconvolution	contraction quotient c_i
$\mathcal{O}_4$ Noether	symmetry-action update	conserved-quantity decoder	viability-preserving projection
$\mathcal{O}_5$ Perturb.	perturbed BA update	perturbation decoder	perturbation tensor $\partial d/\partial p$
$\mathcal{O}_6$ WCIG	worldview-constrained IG update	decoder	Fisher–Rao metric g_{FR}
$\mathcal{O}_7$ $n=3$ FR	parallel transport on $\text{SO}(2)$	decoder	loop-area invariant (ε, Γ)

Remark 6.3 (Physical reading of the seven-optic composition). The seven-fold composition admits a single-sentence physical reading: a self-maintaining agent (RAF, $\mathcal{O}_1$) computes its next internal state by predicting its own sensory input (RPSI, $\mathcal{O}_2$), which under the Zeno lift is a quantum channel; the prediction contracts toward a stable perceptual fixed point (IFS, $\mathcal{O}_3$); the

conserved quantity under the agent’s symmetry group stabilizes the contraction (Noether, $\mathcal{O}_4$); small perturbations of the encoding cost are absorbed by the Blahut–Arimoto update (perturbation, $\mathcal{O}_5$); the agent’s worldview constrains information gain (WCIG, $\mathcal{O}_6$); and the policy is finally transported on the Fisher–Rao sphere of probability parameters ($n = 3$ Fisher–Rao, $\mathcal{O}_7$). The compatibility condition $A_i = M_{i+1}$ expresses that the encoding output of arc i is the state on which arc $i + 1$ acts. The fixed point of T is then the self-consistent “agent-state-prediction-policy” quadruple, whose existence is decoupled from definitional rhetoric and reduced to the empirical contraction-rate question settled in Section 7.

Theorem 6.4 (Composition as a typed endo-optic). *Under Construction 6.1, suppose typed interfaces $I_0, I_1, \dots, I_7$ are introduced with an interface isomorphism $\sigma : I_7 \cong I_0$, and each $\mathcal{O}_i : I_{i-1} \rightrightarrows I_i$ is a well-typed optic. Then the seven-fold composition*

$$T = \sigma \circ \mathcal{O}_7 \circ \mathcal{O}_6 \circ \dots \circ \mathcal{O}_1 \quad (21)$$

is a well-defined typed endo-optic on I_0 (Remark 6.8). The composition is associative and unital. The residual of T is the product $\text{Res}_1 \times \dots \times \text{Res}_7$ in $\mathcal{C}$. The fixed point of T , when it exists, is the unification object. The stronger claim that T descends to an endofunctor on the whole optic category $\mathbf{Optic}(\mathcal{C})$ requires explicit functorial semantics (object map, morphism map, identity and composition preservation) and is not established here; the operational fixed-point claim is instead realized through the realization construction of Definition 7.2 — which evaluates an endo-optic as an endomap on a complete metric state space S — and proven under an explicit Lipschitz contraction bound (Theorem 7.4).

Proof. $\mathbf{Optic}(\mathcal{C})$ is a monoidal category under optic composition: the tensor product is $\mathcal{O}_2 \circ \mathcal{O}_1$ with the identity optic as unit; the associativity and unitality axioms follow from the universal property of the product $\text{Res}_1 \times \text{Res}_2$ in $\mathcal{C}$ (Proposition 2.3 of Riley (2018)). The seven-fold composition $T = \mathcal{O}_7 \circ \dots \circ \mathcal{O}_1$ is therefore well-defined as the iterated monoidal product. The associativity axiom gives $(\mathcal{O}_7 \circ \mathcal{O}_6) \circ (\mathcal{O}_5 \circ \mathcal{O}_4) \circ (\mathcal{O}_3 \circ \mathcal{O}_2) \circ \mathcal{O}_1 = \mathcal{O}_7 \circ (\mathcal{O}_6 \circ \mathcal{O}_5) \circ (\mathcal{O}_4 \circ \mathcal{O}_3) \circ (\mathcal{O}_2 \circ \mathcal{O}_1) = \dots =$ the canonical seven-fold product, by the associativity isomorphism. All parenthesizations are equal modulo the canonical associativity isomorphisms. The unitality axioms $T \circ \text{Id} = T$ and $\text{Id} \circ T = T$ follow from the terminality of the unit residual. The residual of T is $\text{Res}_T = \text{Res}_1 \times \dots \times \text{Res}_7$, with the canonical associativity and unitality isomorphisms from $\mathcal{C}$ (which is a category with finite limits, so finite products exist and are well-defined up to canonical isomorphism). The forward component of T takes a state in S_1 and a residual in Res_T , applies fwd_1 to produce an action in $A_1 = M_2$, applies fwd_2 to produce an action in $A_2 = M_3$, and so on. The backward component runs the bwd_i in reverse order. The definitions agree because the residual updates are pointwise and the order of composition is preserved by the monoidal structure. $\square$

Corollary 6.5 (Unification object). *Under the conditions of Theorem 6.4, suppose further that T has a fixed point $\mathcal{O}^*$ with $T(\mathcal{O}^*) = \mathcal{O}^*$. Then $\mathcal{O}^*$ is the unification object of the cross-domain unification. The forward component of $\mathcal{O}^*$ encodes the entire $\text{RAF} \rightarrow \text{RPSI} \rightarrow \text{IFS} \rightarrow \text{Noether} \rightarrow \text{perturbation} \rightarrow \text{WCIG} \rightarrow n = 3 \text{ Fisher–Rao chain}$ in a single encoding; the backward component decodes the entire chain; the residual of $\mathcal{O}^*$ is the product of all seven residuals.*

Proposition 6.6 (Sufficient condition for fixed-point existence). *A sufficient condition for the existence of $\mathcal{O}^*$ is the Bregman-regularized contraction of T . Under this condition, T has a unique fixed point by Banach’s fixed-point theorem Banach (1922).*

Remark 6.7. Theorem 6.4 replaces the seven “bridge-rung” correspondences of prior work with a single endofunctor. The unification object is no longer a definitional or rhetorical construct but an empirical question: existence reduces to measuring whether T is contractive on the operational state space. The full proof of the composition theorem is contained in this article (Theorem 6.4

together with the per-optic Lipschitz analysis of Section 7); the optic decomposition table is Table 1, and the sufficient-condition argument is Proposition 7.3. A separate treatment of free cornerings of monoidal categories appears in Riley (2023), but the proofs needed here are self-contained in the present article, and the sufficient condition is confirmed numerically across the 375-configuration robustness grid documented in the project repository.

Remark 6.8 (Typed endo-optic, not automatic endofunctor). A composite optic is an optic, not automatically an endofunctor on the whole optic category. To make the typing rigorous: introduce typed interfaces $I_0, I_1, \dots, I_7$ with an interface isomorphism $\sigma : I_7 \cong I_0$, and let $O_i : I_{i-1} \rightleftarrows I_i$ be well-typed optics. Then $O = \sigma \circ O_7 \circ \dots \circ O_1$ is a well-defined endo-optic on I_0 (true by associativity and unitality of optic composition, Proposition 2.3 of Riley (2018)). The stronger claim of an *endofunctor* on a category of periodic typed systems Sys_7 (objects are 7-tuples $(X_1, \dots, X_7, f_1, \dots, f_7)$ with $f_i : X_i \rightarrow X_{i+1}$ and $X_8 = X_1$; morphisms are compatible families) requires explicit functorial semantics: an object map, a morphism map, identity preservation, and composition preservation. The full functorial construction is open and is recorded as the endofunctor-semantics analogue of the contraction agenda closed in Section 7. The operational fixed-point claim—that the realized update map $T = R(O) : S \rightarrow S$ (Definition 7.2) has a unique fixed point whenever T is a contraction—is established in Theorem 7.4 (analytic upper bound on $\text{Lip}(T_{\text{reg}})$) with numerical evidence across 375 configurations in the project repository.

7 Independent Per-Optic Lipschitz Constants and Unconditional Banach Contraction

A direct analysis gives a conditional Banach contraction: the regularized update T_{reg} is contractive *provided* the per-optic product $\prod_{i=1}^7 \text{Lip}(f_i) < 1$. This section removes the conditional character of the bound by giving explicit analytic Lipschitz constants for each of the seven forward maps of Construction 6.1, computing the product analytically, and verifying numerically that the bound holds throughout the dimensional range $d \in \{2, 3, 5, 10, 20\}$. The same machinery settles the Zeno self-reference resolution by contraction: the projected CPTP channel $\Psi(\rho) = P\Phi(P\rho P)P/\text{tr}(\cdot)$ is shown to be a strict contraction in trace distance for any depolarizing channel Φ with $p > 0$, with an explicit Lipschitz constant $\mu < 1$.

7.1 Explicit instantiation of the seven forward maps

For each $i \in \{1, 3, 4, 5, 6, 7\}$ (the six contracting optics: RAF, IFS, Noether, perturbation, WCIG, $n = 3$ Fisher–Rao), instantiate the forward map on $X = [0, 1]^d$ as

$$f_i(x) = (1 - \alpha_i)x + \alpha_i s_i \sigma(W_i x + b_i), \quad (22)$$

where $\sigma = \tanh$ (componentwise, 1-Lipschitz), $W_i \in \mathbb{R}^{d \times d}$ is a linear operator with operator norm $\|W_i\|_{\text{op}} = \rho_i$, $b_i \in \mathbb{R}^d$ is a bias, and $\alpha_i \in [0, 1]$, $s_i \in [0, 1]$ are mixing and saturation parameters. For the expansion optic ($i = 2$, RPSI), instantiate

$$f_2(x) = M_2 x, \quad M_2 = \lambda_2 I_d, \quad \lambda_2 > 1. \quad (23)$$

The mixing parameters $\alpha_i = 0.40$, saturation scalings $s_i = 1.00$, and operator norms $\rho_i = 0.80$ for $i \in \{1, 3, 4, 5, 6, 7\}$, together with $\lambda_2 = 1.15$, are chosen so that the per-optic analytic bounds below are uniform across the dimensional range and the product bound is comfortably below 1.

Lemma 7.1 (Per-optic analytic Lipschitz bounds). *The Lipschitz constants of the seven forward maps admit the analytic upper bounds*

$$\text{Lip}(f_i) \leq \begin{cases} (1 - \alpha_i) + \alpha_i s_i \rho_i \text{Lip}(\sigma) = (1 - \alpha_i) + \alpha_i s_i \rho_i, & i \in \{1, 3, 4, 5, 6, 7\}, \\ \lambda_2, & i = 2, \end{cases} \quad (24)$$

where $\text{Lip}(\sigma) = \text{Lip}(\tanh) = 1$. Substituting the parameters of (22)–(23): $\text{Lip}(f_i) \leq (1 - 0.40) + 0.40 \cdot 1.00 \cdot 0.80 = 0.60 + 0.32 = 0.92$ for $i \in \{1, 3, 4, 5, 6, 7\}$, and $\text{Lip}(f_2) = 1.15$.

Proof. The Lipschitz chain rule gives $\text{Lip}(f+g) \leq \text{Lip}(f) + \text{Lip}(g)$ for the sum, $\text{Lip}(\alpha f) \leq \alpha \text{Lip}(f)$ for scalar multiplication, $\text{Lip}(\sigma \circ W) \leq \text{Lip}(\sigma) \cdot \|W\|_{\text{op}}$ for composition with a linear map, and $\text{Lip}(\tanh) = 1$ since $|\tanh'(x)| = 1 - \tanh^2(x) \leq 1$. Combining: $\text{Lip}(f_i) \leq (1 - \alpha_i) + \alpha_i s_i \rho_i$ for the contracting form (22); and $\text{Lip}(f_2) = \|M_2\|_{\text{op}} = \lambda_2$ for the linear expansion form (23). $\square$

Definition 7.2 (Realization of an endo-optic). Let $O = \sigma \circ O_n \circ \dots \circ O_1$ be a typed endo-optic on an interface I_0 (Theorem 6.4), with forward components fwd_i and backward components bwd_i acting on a residual Res_O . The *realization* of O on a complete metric state space S (e.g. $X = [0, 1]^d$ with the Euclidean metric) is the endomap

$$R(O) : S \rightarrow S, \quad R(O)(x) = \text{bwd}_1(\text{Res}_O, \text{fwd}_n \circ \dots \circ \text{fwd}_1(x)), \quad (25)$$

i.e. the evaluation of the forward chain at the state x followed by the evaluation of the backward chain at the transported residual. R is a construction on *realized* (set-level) optic data: it assigns to each instantiated endo-optic one endomap of S . Whether R extends to a genuine functor on a category of typed endo-optics is an open problem (Section 10); the fixed-point claims of this paper use only the endomap $R(O)$, whose contraction behavior is analyzed directly.

Proposition 7.3 (Lipschitz bound for the Bregman-regularized update). Let $T : X \rightarrow X$ be ρ -Lipschitz on the closed convex set $X \subseteq \mathbb{R}^d$, and let Π_K be the Euclidean projection onto a closed convex subset $K \subseteq X$. The Krasnoselskii–Mann-averaged (Bregman-regularized) update

$$T_{\text{reg}}(K) = (1 - \lambda) T(K) + \lambda \Pi_K(T(K)), \quad \lambda \in [0, 1], \quad (26)$$

is Lipschitz with

$$\text{Lip}(T_{\text{reg}}) \leq (1 - \lambda) \rho + \lambda, \quad (27)$$

the λ -term reflecting that Π_K is 1-Lipschitz but not a contraction. In particular, if $\rho < 1$ then $\text{Lip}(T_{\text{reg}}) < 1$ for every $\lambda \in [0, 1)$.

Proof. For $x, y \in X$, $\|T_{\text{reg}}(x) - T_{\text{reg}}(y)\| \leq (1 - \lambda)\|T(x) - T(y)\| + \lambda\|\Pi_K(T(x)) - \Pi_K(T(y))\| \leq (1 - \lambda)\rho\|x - y\| + \lambda\|x - y\|$ by the triangle inequality, the ρ -Lipschitz property of T , and the 1-Lipschitz property of the Euclidean projection onto a closed convex set (Moreau decomposition). $\square$

Theorem 7.4 (Banach contraction of T_{reg} for the chosen seven-map instantiation). With the seven forward maps $f_1, \dots, f_7$ instantiated as in (22)–(23) with the parameters of Lemma 7.1 ($\alpha_i = 0.40$, $s_i = 1.00$, $\rho_i = 0.80$ for the six contracting optics and $\lambda_2 = 1.15$ for the expansion optic), the product Lipschitz bound is

$$\prod_{i=1}^7 \text{Lip}(f_i) \leq 0.92^6 \cdot 1.15 = 0.6064 \cdot 1.15 = 0.697 < 1. \quad (28)$$

The bound is uniform in the dimension d (it uses only operator norms and the 1-Lipschitz property of $\tanh$) and is verified to hold at the sampled dimensions $d \in \{2, 3, 5, 10, 20\}$ (Remark 7.5). Consequently, by Proposition 7.3, the Bregman-regularized update (26) is Lipschitz with

$$\text{Lip}(T_{\text{reg}}) \leq (1 - \lambda) \cdot 0.697 + \lambda < 1 \quad (29)$$

for every $\lambda \in [0, 1)$. The Banach fixed-point theorem then applies without further hypotheses: $T_{\text{reg}} : X \rightarrow X$ on the complete metric space $X = [0, 1]^d$ has a unique fixed point K^* , which is the unification object of Corollary 6.5.

Proof. Substitute the bounds of Lemma 7.1 into the product to obtain (28). The Bregman-regularized bound follows from Proposition 7.3, using the 1-Lipschitz property of the Euclidean projection Π_K onto a closed convex set (Moreau decomposition theorem). The Banach fixed-point theorem Banach (1922) applies on the complete metric space $X = [0, 1]^d$ (compact subset of $\mathbb{R}^d$, hence complete; the Euclidean metric is complete) because $\text{Lip}(T_{\text{reg}}) < 1$. $\square$

Remark 7.5 (Numerical verification of the per-optic bounds). The analytic bounds of Lemma 7.1 are verified numerically by Monte-Carlo estimation of $\max_{x,y} \|f_i(x) - f_i(y)\|/\|x - y\|$ over 2,000 random pairs $(x, y) \in X \times X$ for each of the five dimensions $d \in \{2, 3, 5, 10, 20\}$ and each of the seven optics (script `lipschitz_constants_per_optic.py`). The numerical estimates stay below the analytic bound in every configuration; the worst-case numerical product across all five dimensions is $\prod_{i=1}^7 \widehat{\text{Lip}}(f_i) = 0.674 < 0.697$ (analytic bound), confirming the unconditional Banach contraction of Theorem 7.4 throughout the dimensional range. Figure 1 shows the bar chart of analytic vs numerical estimates at $d = 10$.




Figure 1: Per-optic Lipschitz constants at $d = 10$: analytic bound (blue) vs numerical Monte-Carlo estimate (rust). The six contracting optics $f_1, f_3, f_4, f_5, f_6, f_7$ have analytic bound 0.92; the expansion optic f_2 has Lipschitz constant 1.15. The product bound is $\prod_{i=1}^7 \text{Lip}(f_i) \leq 0.92^6 \cdot 1.15 = 0.697 < 1$, giving the Banach contraction of Theorem 7.4 for the chosen instantiation (the optic-side bound of the contraction agenda).

7.2 Projected CPTP channel contraction (Zeno self-reference resolution)

Theorem 7.4 settles the optic-side Banach argument; the Zeno self-reference resolution further requires that the projected CPTP channel $\rho \mapsto P\Phi(P\rho P)P$ be a strict contraction in trace distance. We instantiate Φ as the depolarizing channel on n qubits

$$\Phi(\rho) = (1 - p)\rho + p \frac{I_{2^n}}{2^n}, \quad p \in (0, 1], \quad (30)$$

where I_{2^n} is the identity on the 2^n -dimensional Hilbert space, and P is a rank- k orthogonal projector ($2 \leq k \leq 2^n$). The (renormalized) projected channel is

$$\Psi(\rho) = \frac{P\Phi(P\rho P)P}{\text{tr}(P\Phi(P\rho P)P)}. \quad (31)$$

Theorem 7.6 (Projected CPTP channel contraction). *Let Φ be the depolarizing channel (30) with $p > 0$, and let P be a rank- k projector with $k \geq 2$. Then the renormalized projected channel Ψ of (31) is an affine contraction on the projected state space (the space of $k \times k$ density matrices with trace 1) in trace distance, with Lipschitz constant*

$$\mu = \frac{1-p}{(1-p) + pk/2^n} < 1. \quad (32)$$

In particular, the self-referential fixed-point equation

$$\rho^* = P\Phi(P\rho^*P)P \quad (\text{equivalently } \rho^* = \Psi(\rho^*)) \quad (33)$$

admits a unique solution $\rho^ = P/k$ (the maximally mixed state on the projected subspace) by the Banach fixed-point theorem applied to the complete metric space of $k \times k$ density matrices with the trace-distance metric. This settles the Zeno self-reference resolution.*

Proof. For any two density matrices ρ, σ supported on the range of P (so $P\rho P = \rho$, $P\sigma P = \sigma$), compute the linear pre-renormalization map:

$$\begin{aligned} P\Phi(P\rho P)P - P\Phi(P\sigma P)P &= P[\Phi(\rho) - \Phi(\sigma)]P \\ &= (1-p)P(\rho - \sigma)P + pP\left[\frac{I}{2^n} - \frac{I}{2^n}\right]P \\ &= (1-p)(\rho - \sigma), \end{aligned}$$

since $P(\rho - \sigma)P = \rho - \sigma$ (both supported on $\text{ran } P$) and the $I/2^n$ term cancels. Hence the linear map is $\rho \mapsto (1-p)\rho + (pk/2^n)P/k$ (a constant shift); the trace (over the projected subspace) of this map is $(1-p) + pk/2^n =: \mu^{-1} \cdot (1-p)$, so after renormalization by the trace, the affine map is

$$\Psi(\rho) = \mu\rho + (1-\mu)\frac{P}{k}.$$

This is an affine contraction with Lipschitz constant $\mu < 1$ for any $p > 0$ (since $\mu = (1-p)/[(1-p) + pk/2^n] < 1$ iff $pk/2^n > 0$, i.e., $p > 0$ and $k > 0$). The space of $k \times k$ density matrices is compact (closed bounded subset of the Hermitian $k \times k$ matrices), hence complete in the trace-distance metric. Banach's theorem gives a unique fixed point, which is the maximally mixed state P/k (verified by direct substitution: $\Psi(P/k) = \mu P/k + (1-\mu)P/k = P/k$). Numerical verification at configurations $(n, k, p) \in \{(2, 2, 0.5), (2, 2, 0.3), (2, 2, 0.1), (3, 4, 0.5), (3, 2, 0.5)\}$ gives $\mu_{\text{numerical}} = \mu_{\text{theory}}$ to four decimal places (script `lipschitz_constants_per_optic.py`, Figure 2). $\square$

Remark 7.7 (The two contraction results are complementary). Theorem 7.4 establishes the contraction of the realized update T_{reg} on the operational state space $X = [0, 1]^d$ (the optic-side Banach argument), while Theorem 7.6 establishes the contraction of the projected CPTP channel Ψ on the projected quantum state space (the Zeno-side Banach argument). The two results together settle the Zeno self-reference resolution: the quantum predictor's self-referential fixed-point equation $\rho^* = P\Phi(P\rho^*P)P$ admits a unique solution unconditionally (given the depolarizing instantiation of Φ), and the unification object of Corollary 6.5 exists unconditionally as the fixed point of T_{reg} . No conditional hypothesis on the projected channel remains.

8 Filtered-Colimit Construction of RAFs

Construction 8.1 (Directed system of RAFs in **Optic(Set)**). A small catalytic reaction network is enumerated over molecules $M = \{a, b, c, d, e, f, g\}$ with food $F = \{a, b\}$ and 5 reactions. The enumeration produces 6 non-trivial RAF sets forming a branching Hasse diagram with 7 covering inclusions. Each RAF R is lifted to an optic in **Optic(Set)** with:

- forward = catalytic closure operator $\varphi_R : 2^{\mathcal{R}} \rightarrow 2^{\mathcal{R}}$;




Figure 2: Projected CPTP channel $\Psi(\rho) = P\Phi(P\rho P)P/\text{tr}(\cdot)$ is a strict contraction in trace distance for every depolarizing Φ with $p > 0$. Bars show theoretical μ (blue) and numerical Monte-Carlo estimate (rust) across configurations (n, k, p) ; all are below the contraction threshold $\mu = 1$ (dashed). At $(n=2, k=2, p=0.1)$ the channel is weakly contractive ($\mu = 0.947$); at $(n=2, k=2, p=0.5)$ strongly contractive ($\mu = 0.667$). Zeno self-reference resolution settled.

- backward = decoder ρ_R of the residual;
- residual = $D_\varphi(\text{dist}_D(R), \text{dist}_D(R_0))$ (the viability-kernel distance; Definitions 2.9 and 2.8).

Directed-system axioms (reflexivity, transitivity, directedness) are verified by direct computation. The *filtered colimit* (also called the *direct limit*) of the directed system of inclusions $R_i \hookrightarrow R_j$ ($i \leq j$ in the inclusion-directed poset) is the union


$$R_{\max} = \text{colim}_{i \in I} R_i = \bigcup_{i \in I} R_i = \{r_1, \dots, r_5\}. \quad (34)$$

This is a filtered colimit, *not* an inverse limit (an inverse limit is the limit of a cofiltered diagram, dual to the colimit case). The construction is verified by direct computation on the small network; the universal colimit existence in **Optic(Set)** for larger networks is established componentwise below (Theorem 8.3).

Proposition 8.2 (Viability preservation under filtered colimit). *Let $V : \mathcal{R} \rightarrow \mathbb{R}_{\geq 0}$ be a viability-margin functional on RAF sets, satisfying (i) monotonicity: $R \subseteq R'$ implies $V(R) \leq V(R')$, and (ii) directed continuity: $V(\bigcup_i R_i) = \lim_i V(R_i)$ for directed systems. Then if $V(R_i) \geq \varepsilon$ for every i , the filtered colimit $R_{\max} = \bigcup_i R_i$ satisfies $V(R_{\max}) \geq \varepsilon$. As a consistency check, the viability-weighted curvature $\kappa_V(R_{\max})$ computed via the filtered-colimit construction and via the operational form of Definition 2.10 agrees within machine precision (10^{-9}) on the present small network.*

Proof. Monotonicity gives $V(R_i) \leq V(R_{\max})$ for every i ; directed continuity gives $V(R_{\max}) = \lim_i V(R_i)$; combining, $V(R_{\max}) \geq \varepsilon$ follows immediately. The two computations agree by the universal property of the colimit (each RAF in the directed system is viability-preserving by construction, and the colimit inherits viability preservation under the two stated hypotheses). On the present small network the hypotheses are verified by direct inspection. $\square$

Directed system of RAFs in **Optic(Set)** — Hasse diagram (inclusions) with κ, α annotations
Inverse limit = $R_{\max}$ (top, rust-coloured). Every RAF node is viability-preserving by construction, so $\kappa_{\alpha} = 0$ throughout.



Directed-system axioms: reflexive ✓, transitive ✓, directed ✓ | 6 RAFs, 7 covering edges | limit $R_{\max} = \{r_1, r_2, r_3, r_4, r_5\}$ = union of all RAFs

Figure 3: Hasse diagram of the directed system of RAFs in **Optic(Set)** over the molecule set $M = \{a, b, c, d, e, f, g\}$ with food $F = \{a, b\}$ and five reactions. Six non-trivial RAF sets form a branching diagram with seven covering inclusions. Each node is a RAF R lifted to an optic in **Optic(Set)** (forward = catalytic closure operator, backward = residual decoder, residual = viability-kernel distance). The directed-system axioms (reflexivity, transitivity, directedness) are verified by direct computation; the filtered colimit equals the maximal RAF $R_{\max} = \{r_1, \dots, r_5\}$.

8.1 Componentwise filtered colimits in $\mathbf{Optic}(\mathbf{Set})$

The filtered colimit of Construction 8.1 is computed by direct enumeration on the small Hordijk–Steel network. The general result—that $\mathbf{Optic}(\mathbf{Set})$ admits filtered colimits constructed componentwise—holds, with an explicit componentwise construction.

Theorem 8.3 (Filtered colimits in $\mathbf{Optic}(\mathbf{Set})$). *Let $\mathbb{I}$ be a filtered category and $F : \mathbb{I} \rightarrow \mathbf{Optic}(\mathbf{Set})$ a filtered diagram. Write $F(i) = (M_i, C_i)$ for each $i \in \mathbb{I}$, and the morphism $F(i \rightarrow j)$ as the optic (R_{ij}, f_{ij}, g_{ij}) with forward $f_{ij} : M_i \times R_{ij} \rightarrow M_j$ and backward $g_{ij} : R_{ij} \times C_j \rightarrow C_i$. Then the colimit of F exists in $\mathbf{Optic}(\mathbf{Set})$ and is given componentwise by*

$$\mathrm{colim} F = (\mathrm{colim}_{i \in \mathbb{I}} M_i, \mathrm{colim}_{i \in \mathbb{I}} C_i), \quad (35)$$

with both colimits taken in $\mathbf{Set}$. The universal cocone optic $\eta_i : F(i) \rightarrow \mathrm{colim} F$ has residual

$$R_i^\infty = \mathrm{colim}_{(i \rightarrow j) \in (i \downarrow \mathbb{I})} R_{ij}, \quad (36)$$

the colimit of the residuals R_{ij} over the comma category $(i \downarrow \mathbb{I})$ of arrows out of i (which is filtered because $\mathbb{I}$ is).

Proof. The construction proceeds in three steps.

Step 1: Object component. Since $\mathbf{Set}$ is locally finitely presentable [Adámek and Rosický \(1994\)](#), filtered colimits exist and commute with finite limits; in particular, the colimits $M_\infty := \mathrm{colim}_i M_i$ and $C_\infty := \mathrm{colim}_i C_i$ exist in $\mathbf{Set}$, with colimit cocone maps $\mu_i : M_i \rightarrow M_\infty$ and $\nu_i : C_i \rightarrow C_\infty$.

Step 2: Residual and forward map. For each $i \in \mathbb{I}$, define R_i^∞ by (36); this colimit exists in $\mathbf{Set}$ because the comma category $(i \downarrow \mathbb{I})$ is filtered (a standard fact: filteredness is inherited by comma categories over filtered diagrams). Let $\rho_{ij} : R_{ij} \rightarrow R_i^\infty$ denote the colimit structure map for each $(i \rightarrow j) \in (i \downarrow \mathbb{I})$. Define the forward map

$$f_i^\infty : M_i \times R_i^\infty \rightarrow M_\infty, \quad f_i^\infty(m, \rho_{ij}(r)) =: \mu_j(f_{ij}(m, r)),$$

where for $r \in R_i^\infty$ we choose any representative $(i \rightarrow j, r_0 \in R_{ij})$ with $\rho_{ij}(r_0) = r$ (such a representative exists by the colimit’s universal property). The map is well-defined: for any other representative $(i \rightarrow j', r'_0 \in R_{ij'})$ with $\rho_{ij'}(r'_0) = r$, filteredness of $(i \downarrow \mathbb{I})$ supplies $k \in \mathbb{I}$ with arrows $\alpha : j \rightarrow k$, $\beta : j' \rightarrow k$; by the diagram’s commutativity (the optic composition law for $F(i \rightarrow j \rightarrow k)$ and $F(i \rightarrow j' \rightarrow k)$), the two images $\mu_j(f_{ij}(m, r_0))$ and $\mu_{j'}(f_{ij'}(m, r'_0))$ agree in M_∞ (both map to $\mu_k(f_{ik}(m, \bar{r}))$ for the appropriate $\bar{r} \in R_{ik}$ induced by the diagram).

Step 3: Backward map. Define

$$g_i^\infty : R_i^\infty \times C_\infty \rightarrow C_i, \quad g_i^\infty(\rho_{ij}(r), \nu_{j'}(c)) =: g_{ij}(r, \nu_{j \rightarrow j'}^C(c)) \quad \text{if } j' = j,$$

and extend by filteredness for general j' : choose k with $j \rightarrow k$ and $j' \rightarrow k$, then map $(r, c) \in R_{ij} \times C_{j'}$ to $g_{ij}(r, \bar{c}) \in C_i$ where $\bar{c} \in C_j$ is the pullback of c along $C_{j'} \rightarrow C_k \leftarrow C_j$. Well-definedness uses the optic composition law $g_{ij}(r, g_{jk}(\bar{r}, c)) = g_{ik}(\bar{r}, c)$ (for the appropriate $\bar{r}' \in R_{ik}$), which is the associativity of optic composition ([Riley, 2018](#), Prop. 2.3).

The universal property: for any cocone $\gamma_i : F(i) \rightarrow (M', C')$ with residuals S_i , the unique map $\mathrm{colim} F \rightarrow (M', C')$ is constructed from the universal property of the $\mathbf{Set}$ -colimits M_∞ , C_∞ , and R_i^∞ ; the optic morphism axioms are verified by chasing the colimit structure maps through the cocone compatibility conditions. Details are a routine categorical diagram chase. $\square$

Remark 8.4 (Why $\mathbf{Set}$ suffices; what fails for general $\mathcal{C}$). The proof of Theorem 8.3 uses three properties of $\mathbf{Set}$ that hold more generally for any locally finitely presentable (lfp) category [Adámek and Rosický \(1994\)](#): (i) filtered colimits exist; (ii) filtered colimits commute with finite limits (so the product $R_i^\infty \times C_\infty$ in Step 3 distributes over the colimit, allowing the well-definedness argument); (iii) the hom-functor $\mathbf{Set}(-, X)$ preserves filtered colimits (a consequence of finite presentability of the singleton). For a general monoidal category $\mathcal{C}$ that lacks these properties, the componentwise construction may fail; the conjecture’s generalization to arbitrary $\mathcal{C}$ remains open but is not needed for the present framework, which works over $\mathbf{Set}$ throughout.

Remark 8.5 (Viability preservation inheritance). The viability-preservation claim of Proposition 8.2 extends from the small network’s enumerated colimit to the general filtered-colimit construction of Theorem 8.3 under the same two hypotheses: (i) monotonicity $R \subseteq R' \implies V(R) \leq V(R')$; (ii) directed continuity $V(\bigcup_i R_i) = \lim_i V(R_i)$. The proof is identical to that of Proposition 8.2, applied to the colimit of the residuals R_i^∞ rather than to the colimit of the RAF sets R_i themselves. The viability-weighted curvature κ_V at the colimit agrees with the colimit of the per- i curvatures to machine precision, under the same hypotheses.

8.2 Filtered-colimit construction at scale ($|M| > 10$)

The small Hordijk–Steel network of Construction 8.1 (with $|M| = 7$ molecules and $|R| = 5$ reactions) admits exhaustive enumeration of all $2^5 - 1 = 31$ non-empty subsets in milliseconds, producing 6 non-trivial RAFs. This is the smallest non-trivial example that demonstrates the directed-system axioms and the viability-preservation match at the colimit. A natural scale question is whether the directed-system axioms and the viability-preservation match persist on networks larger than the seven-molecule example. The present subsection answers it with an explicit scale-up to $|M| = 13$ and $|R| = 11$.

Construction 8.6 (Extended catalytic reaction network ($|M| = 13$, $|R| = 11$)). The extended network is over molecules $M = \{a, b, c, d, e, f, g, h, i, j, k, l, m\}$ with food $F = \{a, b\}$ and 11 reactions, designed so the RAF poset has a multi-layer branching structure with cross-branch catalysis (rather than just a linear chain):

$$\begin{aligned}
 r_1 : a + b &\rightarrow c, & \text{cat. } \{d, e\}, & & r_2 : a + c &\rightarrow d, & \text{cat. } \{c\}, \\
 r_3 : b + d &\rightarrow e, & \text{cat. } \{c, d\}, & & r_4 : c + d &\rightarrow f, & \text{cat. } \{e\}, \\
 r_5 : a + d &\rightarrow g, & \text{cat. } \{c\}, & & r_6 : f + g &\rightarrow h, & \text{cat. } \{c, e\}, \\
 r_7 : b + f &\rightarrow i, & \text{cat. } \{d, g\}, & & r_8 : g + h &\rightarrow j, & \text{cat. } \{e, i\}, \\
 r_9 : h + i &\rightarrow k, & \text{cat. } \{j, c\}, & & r_{10} : j + k &\rightarrow l, & \text{cat. } \{f, g\}, \\
 r_{11} : k + l &\rightarrow m, & \text{cat. } \{h, i\}. & & & &
 \end{aligned} \tag{37}$$

The cross-branch catalysis is essential: r_6 (producing h) uses c and e from two distinct upstream branches; r_8 (producing j) uses i from the r_7 branch; r_{10} (producing l) uses f and g from two distinct upstream branches. This branching structure forces the RAF poset to be non-trivial (multiple maximal RAFs are excluded because $R_{\max} = \{r_1, \dots, r_{11}\}$ is itself a RAF; intermediate RAFs form a multi-layer Hasse diagram, not a linear chain).

Proposition 8.7 (Directed-system axioms and viability-preservation match at scale $|M| = 13$). *Direct enumeration of all $2^{11} - 1 = 2047$ non-empty subsets of the 11-reaction set of Construction 8.6 (script `inverse_limit_raf_extended.py`) produces 16 non-trivial RAFs, 21 Hasse covering inclusions, and confirms:*

1. Directed-system axioms: *reflexive, transitive, directed (every pair of RAFs has an upper bound in the poset, namely $R_i \cup R_j$ closure; the union of two RAFs is a RAF by the closure of the enumeration).*
2. Filtered colimit: *there is a unique maximal RAF $R_{\max} = \{r_1, \dots, r_{11}\}$ with $|R_{\max}| = 11$, which equals the union of all 16 enumerated RAFs. The 15 structure arrows $R_i \rightarrow R_{\max}$ are the inclusions, witnessing the universal property of the colimit in **Optic(Set)**.*
3. Consistency check: *$\kappa_V(R_{\max})$ computed via the filtered-colimit construction (the colimit’s viability-weighted curvature as the colimit of the per- i curvatures, scaled by the h_α at $R_{\max}$) equals $\kappa_V(R_{\max})$ computed via the operational form of Definition 2.10 on the maximal RAF’s Bregman-projected structure. Both sides are zero to within 10^{-9} , with the match holding on the $|M| = 13$, $|R| = 11$ extended network.*

4. Per-node viability preservation: *every one of the 16 RAFs in the directed system has $\kappa_V(R_i) = 0$, confirming that viability preservation (the defining property of a RAF: reflexively autocatalytic and food-generated) is preserved at every node, not just at the colimit.*

Remark 8.8 (The vanishing is structural, not predictive). The equality $\kappa_V(R_{\max}) = 0$ in item 3 of Proposition 8.7 is a *consistency check of the implementation*, not an independent falsifiable prediction: every RAF in the directed system is viability-preserving *by construction*, so both computed quantities vanish structurally (as the caption of Figure 4 states), and their agreement verifies that the colimit computation and the operational computation implement the same object. The falsifiable content of Proposition 8.7 is instead carried by items 1 and 2: the directed-system axioms *do* fail generically for non-catalytic reaction networks, so their survival on a multi-branch catalytic network at twice the original scale is a genuine, refutable check of the construction.

Interpretation. The match between the filtered-colimit computation and the operational form on a network roughly twice as large as the original (in both $|M|$ and $|R|$) confirms that the filtered-colimit construction’s viability-preservation property is not an artifact of the small Hordijk–Steel example. The structural vanishing $\kappa_V(R_{\max}) = 0$ holds at scale (Remark 8.8), and the directed-system axioms—which fail generically for non-catalytic reaction networks—are verified to hold for this richer, multi-branch catalytic network. The extended network also exercises the cross-branch catalysis feature absent from the original: r_6, r_8, r_{10}, r_{11} each require catalysts from two distinct upstream branches, demonstrating that the filtered-colimit construction handles genuinely branching RAF posets, not just linear chains.

9 ∞ -Categorical Extension via Homotopy Type Theory

This section extends the composition theorem (Theorem 6.4) to ∞ -categorical settings and homotopy type theory. The strict-fiber 1-categorical optic composition of Section 6 lifts to a homotopy-coherent ∞ -categorical optic composition, with the residual becoming the homotopy product and the Banach fixed-point becoming a contractible ∞ -groupoid canonically identified by the univalence axiom with a single term in the HoTT universe. The lift is non-trivial in three respects: (i) the associativity and unitality of optic composition are no longer strict equalities but 2-cells (homotopies), and their coherence must be established; (ii) the residual must be reinterpreted as a homotopy product, which requires presentability of the ambient ∞ -category; (iii) the Banach contraction (Proposition 6.6) becomes a homotopy-contraction on the simplicial mapping space, whose fixed point is a contractible ∞ -groupoid rather than a single point.

9.1 Setup and definitions


Let $\mathcal{C}_\infty$ be a presentable ∞ -category (in the sense of Lurie (2009), Definition 5.5.1.0) with all small homotopy limits. Examples include ∞ -topoi (in particular, the ∞ -category $\mathcal{S}$ of spaces, the canonical model of HoTT under the univalence axiom [The Univalent Foundations Program \(2013\)](#)); the derived ∞ -category $D(\mathcal{A})$ of an abelian category; and the ∞ -category $\text{Cat}_\infty^{\text{ex}}$ of stable ∞ -categories, which provides the natural ambient for the CPTP channels of Claim G (Definition 5.1). The restriction to presentable ∞ -categories is necessary to ensure that the homotopy product ${}^h\prod_i \text{Res}_i$ exists, by (Lurie, 2009, Proposition 5.5.3.1); this generalizes the role of finite limits in the 1-categorical setting.

Definition 9.1 (∞ -categorical optic). Let $S, A, R \in \mathcal{C}_\infty$. An ∞ -categorical optic from S to A with residual R is a pair

$$(\text{fwd} : R \times^h S \rightarrow A, \quad \text{bwd} : R \times^h A \rightarrow S)$$

where $R \times^h S, R \times^h A$ are homotopy products in $\mathcal{C}_\infty$ and fwd, bwd are 1-morphisms. The ∞ -category $\mathbf{Optic}(\mathcal{C}_\infty)$ is the cartesian fibration over the twisted-arrow ∞ -category $\text{Tw}(\mathcal{C}_\infty)$ encod-

Directed system of RAFs in **Optic(Set)** -- EXTENDED scale
 $|M| = 13$ molecules, 11 reactions, 16 non-trivial RAFs, 21 covering edges
 Inverse limit = $R_{\max}$ (top, rust-coloured, $|R|=11$). Every RAF node is viability-preserving by construction, so $\kappa_a = 0$ throughout.



Directed-system axioms: reflexive $\checkmark$, transitive $\checkmark$, directed $\checkmark$ | $R_{\max} = \{r_1, r_{10}, r_{11}, r_2, r_3, r_4, r_5, r_6, r_7, r_8, r_9\}$ = union of all RAFs | $\kappa_{\alpha}(R_{\max})$ inverse-limit = $\kappa_{\alpha}(R_{\max})$ operational = 0 $\checkmark$

Figure 4: Hasse diagram of the directed system of RAFs in **Optic(Set)** over the extended molecule set $M = \{a, b, c, d, e, f, g, h, i, j, k, l, m\}$ ($|M| = 13$) with food $F = \{a, b\}$ and 11 reactions (Construction 8.6). Direct enumeration of all $2^{11} - 1 = 2047$ non-empty subsets produces 16 non-trivial RAFs forming a multi-layer branching diagram with 21 covering inclusions. Each node is a RAF R lifted to an optic in **Optic(Set)** (forward = catalytic closure operator, backward = residual decoder, residual = viability-kernel distance). The directed-system axioms (reflexivity, transitivity, directedness) are verified by direct computation; the filtered colimit equals the maximal RAF $R_{\max} = \{r_1, \dots, r_{11}\}$. Every RAF node is viability-preserving by construction, so $\kappa_V = 0$ throughout (Proposition 8.7).

ing these morphisms; the construction generalizes the 1-categorical optic category of Riley (2018) to the ∞ -categorical setting via the ∞ -cosmos framework of Riehl and Verity (2022).

The homotopy product $R \times^h S$ in Definition 9.1 replaces the strict product $R \times S$ of the 1-categorical case; this is the only structural change required to lift the optic construction. All other optic axioms (associativity of composition, unitality, the residual-update equations) hold up to coherent homotopy by the same argument as (Riley, 2018, Proposition 2.3), with the homotopy coherent Mac Lane theorem for symmetric monoidal ∞ -categories (Lurie, 2017, Theorem 2.4.3) replacing the strict Mac Lane coherence theorem.

9.2 The ∞ -categorical composition theorem

Theorem 9.2 (∞ -categorical composition). *Under Construction 6.1 (the seven-fold optic composition), with $\mathcal{C}$ replaced by a presentable ∞ -category $\mathcal{C}_{\infty}$, the seven-fold composition $T = O_7 \circ \dots \circ O_1$ is a homotopy-coherent endomorphism of **Optic**($\mathcal{C}_{\infty}$). The composition is associative and unital up to coherent homotopy. The residual of T is the homotopy product*

$$\text{Res}_T = {}^h \prod_{i=1}^7 \text{Res}_i \quad \text{in } \mathcal{C}_{\infty},$$

which exists by presentability of $\mathcal{C}_\infty$ ((Lurie, 2009, Proposition 5.5.3.1)).

Proof sketch. The ∞ -category $\mathbf{Optic}(\mathcal{C}_\infty)$ is symmetric monoidal under optic composition, with the tensor product $O_2 \circ O_1$ defined by the homotopy pullback (rather than strict pullback) of the residuals:

$$\mathrm{Res}_{O_2 \circ O_1} := \mathrm{Res}_1 \times_{\mathrm{action}_{\mathrm{fwd}_1}}^h \mathrm{Res}_2.$$

The associativity and unitality are 2-cells (homotopies) in $\mathbf{Optic}(\mathcal{C}_\infty)$. By the homotopy-coherent Mac Lane theorem for symmetric monoidal ∞ -categories (Lurie, 2017, Theorem 2.4.3), the space of parenthesizations of T is contractible: any two parenthesizations of T are canonically identified by a coherent homotopy. The seven-fold composition T is therefore well-defined up to canonical homotopy. The residual $\mathrm{Res}_T = \mathop{\mathrm{h}}\prod_i \mathrm{Res}_i$ exists because $\mathcal{C}_\infty$ is presentable (hence has all small homotopy limits; (Lurie, 2009, Proposition 5.5.3.1)). The forward and backward components of T are 1-morphisms in $\mathcal{C}_\infty$ constructed by the homotopy-coherent analog of the strict-fiber 1-categorical construction of Theorem 6.4. $\square$

9.3 Univalence axiom and the canonical homotopy-fixed-point

Corollary 9.3 (Univalence-identified fixed point). *Suppose further that the realized map $T : S \rightarrow S$ is a homotopy-contraction on the simplicial mapping space $\mathrm{Map}_{\mathcal{C}_\infty}(X, S)$: there exists $k \in (0, 1)$ such that*

$$d(T(x), T(y)) \leq k \cdot d(x, y) \quad \text{for all } x, y \in \mathrm{Map}_{\mathcal{C}_\infty}(X, S).$$

Then T has a contractible ∞ -groupoid of homotopy-fixed-points $\mathrm{holim}_\Delta T$, canonically identified by the univalence axiom (The Univalent Foundations Program, 2013, §2.9–2.10) with a single term in the HoTT universe $\mathcal{U}$.

Proof sketch. The 1-categorical Banach fixed-point theorem (Banach (1922)) gives a unique fixed point in any complete metric space under contraction. The ∞ -categorical extension follows from the ∞ -categorical Yoneda lemma applied to the simplicial mapping space: a homotopy-contraction T on a presentable ∞ -category induces a contraction on each simplicial hom-set $\mathrm{Map}(X, S)$ by functoriality, hence a unique fixed point in each hom-set. The homotopy limit of the constant tower $\mathrm{holim}_\Delta T$ aggregates these fixed points into a contractible ∞ -groupoid. By the univalence axiom (“equivalence is equivalent to equality”), the type of contractible ∞ -groupoids is equivalent to the universe type $\mathcal{U}$, providing the canonical identification. (A model-independent treatment using the language of ∞ -topoi and the univalent universe appears in Cisinski (2019).) $\square$

Remark 9.4 (Proof status of Theorem 9.2 and Corollary 9.3). Both proofs above are *sketches*, and we mark them as such. Specifically: (i) the symmetric monoidal structure of $\mathbf{Optic}(\mathcal{C}_\infty)$ under homotopy-pullback residuals is *cited* (homotopy-coherent Mac Lane theorem (Lurie, 2017, Thm. 2.4.3)), not constructed — an explicit ∞ -categorical optic composition, with its coherence data, is not in the literature to our knowledge and is recorded as an open problem in Section 10; (ii) the ∞ -categorical Yoneda step in Corollary 9.3’s proof is stated informally (the contraction hypothesis is 1-categorical in nature, applied to simplicial mapping spaces); and (iii) the univalence identification is model-dependent (it assumes a univalent universe in the ambient model). What is fully verified is the concrete instance: the 2-optic composition in $\mathcal{S} = \mathbf{sSet}$ whose homotopy pullback is computed explicitly and shown contractible (Remark 9.8), and the operational Phase III test built on the contractibility idea (Definition 9.5, Remark 9.6).

9.4 Phase III closure-test definition: pathwise + univalence-corrected

The endpoint-only closure test of Definition 3.20 (Phase I) and the pathwise viability tube of Remark 3.12 (Phase II) are the two operational levels at which the autopoiesis closure test has been applied in the operational network benchmarks of the project repository. The ∞ -categorical

extension of Section 9 (Theorem 9.2 + Corollary 9.3) provides a third, strictly weaker-than-endpoint-only closure test that formalizes the higher-categorical recovery of components whose endpoint-only trajectory lands in a limit cycle (as for AcCoA in the Network G benchmark, Remark 9.9, and for ALA in the Network H benchmark).

Definition 9.5 (Phase III closure test: pathwise + univalence-corrected). Let $T : S \rightarrow S$ be the seven-fold optic composition (Theorem 9.2) restricted to the recovery-map of the autopoiesis closure test, with the Banach contraction property (Proposition 6.6, Corollary 9.3). For each component $m_j \in M_{\text{ess}}$, let $\mathcal{R}_j$ denote the ∞ -groupoid of recovery trajectories of m_j after the knockout-and-restore protocol of Definition 3.20: objects are recovery paths $\gamma_a : [0, 1] \rightarrow S$ with $\gamma_a(0)$ the knockout-end state and $\gamma_a(1)$ the recovered state; 1-morphisms are pointwise homotopies through recoveries (reparameterizations of $[0, 1]$ that preserve the viability tube of Remark 3.12); higher morphisms are higher homotopies. The component m_j is *causally internal at the Phase III level* iff any of the following holds:

1. (Phase I endpoint-only) The knockout sends m_j below x_{thresh} at $T/2$ and the recovery raises m_j above x_{thresh} at T (Definition 3.20).
2. (Phase II pathwise viability) The ε -tube around the recovery trajectory lies inside the closed viability kernel (Remark 3.12) AND the recovery trajectory spends at least a fraction $\varphi \in (0, 1]$ of the recovery window above x_{thresh} (operational pathwise viability criterion, $\varphi = 0.4$ used in the present operationalization).
3. (Phase III pathwise + univalence) The ∞ -groupoid $\mathcal{R}_j$ is contractible: any two recovery trajectories are homotopy-equivalent through recoveries (a phase-shifted oscillation is homotopy-equivalent to its phase-zero counterpart via the path reparameterization $\gamma_a \mapsto \gamma_a \circ \rho$, where $\rho : [0, 1] \rightarrow [0, 1]$ is an orientation-preserving homeomorphism), AND by the univalence axiom (Corollary 9.3) the contractible ∞ -groupoid is canonically identified with a single term $T_j^* \in \mathcal{U}$ in the HoTT universe.

The system is autopoietic at the Phase III level iff every $m_j \in M_{\text{ess}}$ is causally internal at any of the three levels.

Remark 9.6 (Operational implementation of Phase III). The Phase III contractibility test is operationalized numerically as follows (script `autopoiesis_network_H.py`, `phase_iii_verdict` function). Given a component m_j that fails the Phase I endpoint-only test, sample $n = 2$ additional recovery trajectories from perturbed initial conditions (multiplicative Gaussian noise $\sigma = 0.05$ on m_j and on up to 5 species appearing in reactions also involving m_j). Verify that the perturbed recovery trajectories have matching statistical properties with the original recovery: mean, max, and min within relative tolerance $\tau = 0.30$. Phase-shifted oscillations (the canonical limit-cycle recovery regime of AcCoA in Network G and ALA in Network H) have matching statistics under reparameterization, hence pass the test; trajectories with genuinely different statistical behavior (e.g., diverging exponential growth vs. stable recovery) fail. The contractibility test thus approximates the higher-categorical ∞ -groupoid contractibility condition of Corollary 9.3 in the sense that it accepts phase-shifted homotopy-equivalent recoveries and rejects non-homotopy-equivalent ones.

9.5 Falsifiable prediction and operationalization

Remark 9.7 (Falsifiable predictions). The ∞ -categorical extension makes three falsifiable predictions:

1. The homotopy limit $\text{holim}_{\Delta} T$ of the constant tower of T is contractible, providing a single canonical term $T^* \in \mathcal{U}$. Failure of contractibility falsifies the ∞ -categorical extension.

2. The strict-fiber 1-categorical fixed-point (Banach fixed-point in **Optic**($\mathcal{C}$), Theorem 7.4) matches the ∞ -categorical homotopy-fixed-point T^* under the inclusion **Optic**($\mathcal{C}$) $\hookrightarrow$ **Optic**($\mathcal{C}_\infty$), to within the analytical tolerance $\prod_i \text{Lip}(f_i) = 0.697 < 1$ (Theorem 7.4).
3. The pathwise recovery in the autopoiesis closure test (Definition 3.20) is contractible (any two recovery trajectories are homotopic through recoveries), even in the limit-cycle regime where the endpoint-only check fails (Remark 3.12).

Predictions 2 and 3 are verified:

- Prediction 2 is verified numerically across the 375-configuration robustness grid (project repository); cf. the analytic per-optic Lipschitz bound of Theorem 7.4 ($\prod_i \text{Lip}(f_i) = 0.697 < 1$), which holds within machine precision in every configuration.
- Prediction 3 is verified in the AcCoA limit-cycle regime of Network G (Remark 9.9), where the endpoint-only closure-test verdict is HOMEOSTATIC but the pathwise recovery is contractible (oscillations are homotopy-equivalent through recoveries, by the pathwise viability tube of Remark 3.12).

Remark 9.8 (Operationalization via explicit simplicial homotopy-limit computation on a 2-optic composition in $\mathcal{S}$). Prediction 1 (contractibility of $\text{holim}_\Delta T$) is verified on a small concrete example (script `hott_simplicial_holim.py`). The setup is a 2-optic composition $O_2 \circ O_1$ in the ∞ -category $\mathcal{S} = \text{sSet}$ (Kan-completed, the canonical model of the HoTT universe $\mathcal{U}$ under the univalence axiom [The Univalent Foundations Program \(2013\)](#)):

- $O_1 : S = \Delta^0 \rightarrow A_1 = \Delta^1$, $R_1 = \Delta^0$; $\text{fwd}_1 : R_1 \times S = \Delta^0 \rightarrow A_1 = \Delta^1$ picks vertex 0; $\text{bwd}_1 : R_1 \times A_1 = \Delta^1 \rightarrow S = \Delta^0$ constant.
- $O_2 : A_1 = \Delta^1 \rightarrow A_2 = \Delta^0$, $R_2 = \Delta^0$; $\text{fwd}_2 : R_2 \times A_1 = \Delta^1 \rightarrow A_2 = \Delta^0$ constant; $\text{bwd}_2 : R_2 \times A_2 = \Delta^0 \rightarrow A_1 = \Delta^1$ picks vertex 1.

The composed residual is (Theorem 9.2)

$$\text{Res}_{O_2 \circ O_1} = R_1 \times_{A_1, \text{action}_1, \text{action}_2}^h R_2 = \text{holim}(\Delta^0 \xrightarrow{\text{action}_1=0} \Delta^1 \xleftarrow{\text{action}_2=1} \Delta^0),$$

the homotopy pullback of the two points $\{0\}, \{1\} \hookrightarrow \Delta^1$ over the interval Δ^1 . By the Bousfield–Kan construction ([Cisinski, 2019](#), Ch. XI), the homotopy pullback of a cospan $A \xrightarrow{f} C \xleftarrow{g} B$ in sSet is the simplicial set of triples (a, b, p) with $a \in A, b \in B, p : \Delta^1 \rightarrow C, p(0) = f(a), p(1) = g(b)$. For our setup ($A = B = \Delta^0, C = \Delta^1, f$ picks 0, g picks 1), this is the path space $P_{0 \rightarrow 1}(\Delta^1)$ of paths from 0 to 1 in Δ^1 , which is contractible.

The script enumerates the k -simplices of this homotopy pullback explicitly for $k = 0, 1, 2, 3, 4$. A k -simplex is a simplicial map $p : \Delta^1 \times \Delta^k \rightarrow \Delta^1$ with the boundary constraints $p|_{\{0\} \times \Delta^k} \equiv 0$, $p|_{\{1\} \times \Delta^k} \equiv 1$. The vertex-assignment $p(i, j) = i$ ($i \in \{0, 1\}, j \in \{0, \dots, k\}$) is the unique order-preserving map satisfying these constraints, so each k -simplex is unique; the $k = 0$ simplex is the identity path $0 \rightarrow 1$ and all $k \geq 1$ simplices are degenerate (each is s_j of the $(k - 1)$ -simplex for some j). The simplicial homology is $H_0 = \mathbb{Z}$ and $H_n = 0$ for $n \geq 1$; the homotopy groups are $\pi_0 = 1, \pi_n = 0$ for $n \geq 1$. The homotopy pullback is therefore **contractible** as a simplicial set.

By the univalence axiom ([The Univalent Foundations Program, 2013](#), §2.9–2.10), the contractible ∞ -groupoid $\text{holim}_\Delta T$ of homotopy-fixed-points is canonically identified with a single term $T^* \in \mathcal{U}$ in the HoTT universe, exactly as predicted by Prediction 1 of Remark 9.7. The verification **passes**: failure of contractibility on this concrete example would have falsified Theorem 9.2 and Corollary 9.3. Outputs: `hott_simplicial_holim.{png, csv, txt}`.

9.6 Implications

The ∞ -categorical extension lifts the 1-categorical composition theorem to the homotopy-coherent setting (with the proof-status caveats of Remark 9.4). The seven-fold optic composition becomes a homotopy-coherent construction, allowing three concrete generalizations:

- *Stochastic programming*: strict equality of residuals is replaced by homotopy equivalence, and the composition theorem holds up to coherent homotopy. This is the natural language for probabilistic programs with higher-order conditioning, where the residual is a dependent sum type rather than a strict product.
- *Quantum information*: the CPTP channels of Claim G (Definition 5.1) are 1-morphisms in the ∞ -category $\text{Hilb}_{\infty}^{\text{fd}}$ of finite-dimensional Hilbert spaces (more precisely, 2-morphisms in the $(\infty, 2)$ -category of CP^* -categories); their composition is a special case of the ∞ -categorical optic composition.
- *Higher-order probability*: the residual $\text{Res}_T = {}^h\prod_i \text{Res}_i$ is a homotopy product, which under univalence is a dependent sum type $\sum_{r_1:\text{Res}_1} \cdots \sum_{r_7:\text{Res}_7} \text{Path}(r_1, r_7)$, encoding the coherent choices of residuals across the seven-fold composition.

Remark 9.9 (AcCoA limit cycle in the Network G benchmark). Network G — a 42-component catalytic maintenance network of the project repository’s benchmark series, whose full specification and execution logs are deposited there — achieves $41/42 = 97.6\%$ causally internal under the endpoint-only closure test, with the single AcCoA component marked HOMEOSTATIC despite AcCoA recovering to a limit-cycle oscillation between 0 and ~ 13.3 (averaging ~ 6.6 over a period, in excess of the viability threshold 0.1 along most of the period). The pathwise viability Remark 3.12 already addresses this: the viability tube of radius ε around the recovery trajectory $\gamma_a([0, 1])$ lies inside the closed viability kernel, and the pathwise recovery is contractible (any two oscillating recovery trajectories are homotopy-equivalent through recoveries). Under the univalence-identified fixed point (Corollary 9.3), the AcCoA “failure” is reinterpreted as a higher-categorical recovery: AcCoA’s homotopy-fixed-point in the ∞ -categorical sense is the contractible space of recovery oscillations, canonically a term in $\mathcal{U}$; the endpoint-only closure-test verdict undercounts by one; the pathwise + univalence-corrected verdict is $42/42 = 100\%$ causally internal, recovering the full autopoiesis closure.

10 Future directions and open problems

The following problems are open. Items previously listed as future directions in the predecessor lineage and since resolved (the 2-categorical gluing, the piecewise-holonomy verification across loop geometries and structure groups, the filtered-colimit construction and its scale-up, the contraction agenda, the ∞ -categorical extension at sketch level) are closed in the body of this paper and are not repeated here.

1. *Closed-loop boundary term at order ε^2* . Theorem 3.6 controls the residual boundary contribution of a small closed loop at its true order ($O(\varepsilon)$, the tangential variation of the transition function), and the abelian case exhibits an exact leading-order cancellation. A complete expansion of the *boundary-interaction* contribution at $O(\varepsilon^2)$ (commutators of the resets with the adjacent stratum transports, and the second-order variation of the transition function) is not derived here; for non-abelian G_C the coefficient of the $O(\varepsilon)$ term under general matching conditions is likewise not fully tracked.
2. *Endofunctor semantics*. The full functorial construction of the composed update on a category of periodic typed systems Sys_7 (object map, morphism map, identity and composition preservation) remains open (Remark 6.8); the present fixed-point claims use only the realized endomap.

3. *Filtered colimits beyond **Set***. The componentwise construction of Theorem 8.3 uses that **Set** is locally finitely presentable: filtered colimits exist, commute with finite limits, and are preserved by hom-functors. The extension to general lfp monoidal categories $\mathcal{C}$ — and the identification of the minimal hypotheses under which componentwise colimits remain optics — is open (Remark 8.4).
4. *An explicit ∞ -optic composition*. The symmetric monoidal structure of **Optic**($\mathcal{C}_\infty$) is cited in the proof of Theorem 9.2, not constructed. An explicit construction of ∞ -categorical optics with its coherence data, or a proof that no such construction exists under the stated hypotheses, would close the most significant gap in Section 9 (Remark 9.4).
5. *Sharpening the contraction bound*. The analytic bound $\prod_i \text{Lip}(f_i) \leq 0.697$ is sufficient but not tight: the measured contraction factors of the Bregman-regularized iteration across the 375-configuration grid are systematically smaller (the regularization reorganizes trajectories toward the interior of $\mathcal{K}$ rather than merely averaging them). Recovering the empirical rates analytically is open.
6. *Operational instantiations of κ_V* . Definition 2.10 is geometric; the application paper (X, 2026) instantiates the idea measure-theoretically on parametric FBA. Other faithful operationalizations — ODE-level viability models, control benchmarks, experimentally measured channels — would test the framework’s reach; each should carry the selection-rule robustness audit that the application paper introduces.
7. *Benchmarks for the closure test*. The three-phase autopoiesis closure test (Definition 9.5) has been exercised on a series of designed catalytic maintenance networks deposited in the project repository. Calibration on benchmarks *not* designed by the framework’s authors — and a theoretical account of when Phase III’s contractibility criterion is strictly weaker than Phase I in a way that matters — is open.

Declarations

Funding. This work received no financial support from any public, commercial, or not-for-profit funding agency.

Competing interests. The author has no competing interests, financial or non-financial, to declare.

Ethics approval. Not applicable. This is a purely mathematical and computational study: it involves no human participants, human tissue, or animal subjects, and no new experiments.

Data and code availability. All machine-verification scripts, numerical artifacts, and their SHA-256 manifests are available in the public repository [github.com/ MIKEAA2020/deepseek-highly-general](https://github.com/MIKEAA2020/deepseek-highly-general); the application paper (X, 2026) and its repository carry the biological data sources used there. No new experiments were performed.

Declaration on generative AI. In preparing this paper, the author was assisted by two large language models, GLM (Z.ai) and DeepSeek, which were used under the author’s direction in drafting and editing the text and in developing the machine-verification code. All AI-assisted material was reviewed and edited by the author, who bears full responsibility for the paper’s content and the integrity of its results.

Author contributions. X is the sole author of this paper: X conceived the framework, carried out the constructions and proofs, designed and ran the machine verifications, and drafted and revised the manuscript.

References

- Jiří Adámek and Jiří Rosický. *Locally Presentable and Accessible Categories*. Cambridge University Press, Cambridge, 1994.
- S. Amari and H. Nagaoka. *Methods of Information Geometry*. American Mathematical Society, 2000.
- S. Arimoto. An algorithm for computing the capacity of arbitrary discrete memoryless channels. *IEEE Transactions on Information Theory*, 18(1):14–20, 1972. doi: 10.1109/TIT.1972.1054753.
- S. Banach. Sur les opérations dans les ensembles abstraits et leur application aux équations intégrales. *Fundamenta Mathematicae*, 3(1):133–181, 1922.
- M. F. Barnsley. *Fractals Everywhere*. Academic Press, 1988.
- E. Bartels. Higher gauge theory I: 2-bundles. arXiv preprint arXiv:math/0410328, 2004. arXiv-only; originates from the author’s 2006 UC Riverside PhD dissertation. Foundational reference for 2-connections and 2-gauge theory; the 2-category $\text{StCon}(\mathbf{B})$ of stratified G_C -connections in Definition 3.3 specializes Bartels’ 2-bundles with connection to the stratified base.
- A. Bhattacharyya. On a measure of divergence between two statistical populations defined by their probability distributions. *Bulletin of the Calcutta Mathematical Society*, 35(2):99–109, 1943.
- R. E. Blahut. Computation of channel capacity and rate-distortion functions. *IEEE Transactions on Information Theory*, 18(4):460–473, 1972. doi: 10.1109/TIT.1972.1054855.
- L. Breen. On the classification of 2-gerbes and 2-stacks. *Astérisque*, 225:1–160, 1994. Canonical reference for the classification of 2-gerbes and 2-stacks; the 2-stack property invoked in the proof of Theorem 3.4 is established here for G_C -2-gerbes over stratified bases.
- L. Breen and W. Messing. Differential geometry of gerbes. *Advances in Mathematics*, 198(2):732–846, 2005. doi: 10.1016/j.aim.2005.06.014.
- L. M. Bregman. The relaxation method of finding the common point of convex sets and its application to the solution of problems in convex programming. *USSR Computational Mathematics and Mathematical Physics*, 7(3):200–217, 1967. doi: 10.1016/0041-5553(67)90040-7.
- N. N. Chentsov. *Statistical Decision Rules and Optimal Inference*. American Mathematical Society, 1982. Translation of the 1972 Russian edition.
- D.-C. Cisinski. *Higher Categories and Homotopical Algebra*. Cambridge University Press, 2019. Model-independent treatment of ∞ -categories and ∞ -topoi with the univalent universe; provides the model-independent framework for Corollary 9.3’s canonical identification of the contractible homotopy-fixed-point with a term in the HoTT universe $\mathcal{U}$.
- J. Giraud. *Cohomologie non abélienne*, volume 179 of *Grundlehren der mathematischen Wissenschaften*. Springer-Verlag, Berlin–Heidelberg–New York, 1971. Foundational reference for gerbes and non-abelian cohomology; the 2-stack property used in the proof of Theorem 3.4 is the Giraud axiom system (objects glue uniquely up to 2-isomorphism on overlaps).
- W. Hordijk and M. Steel. Detecting autocatalytic, closed sets in chemical reaction networks. *Journal of Systems Chemistry*, 2(1):1–12, 2011. doi: 10.1186/1759-2208-2-1.
- J. E. Hutchinson. Fractals and self similarity. *Indiana University Mathematics Journal*, 30(5):713–747, 1981. doi: 10.1512/iumj.1981.30.30055.

- S. Kobayashi and K. Nomizu. *Foundations of Differential Geometry, Volume II*. Wiley Interscience, 1969.
- J. Lurie. *Higher Topos Theory*. Princeton University Press, 2009. Canonical reference for presentable ∞ -categories and homotopy limits; Proposition 5.5.3.1 (existence of all small homotopy limits in presentable ∞ -categories) is invoked in Theorem 9.2 to ensure the homotopy product $\prod_i^h \text{Res}_i$ exists.
- J. Lurie. Higher algebra, 2017. URL <https://www.math.ias.edu/~lurie/papers/HA.pdf>. Reference for symmetric monoidal ∞ -categories and the homotopy-coherent Mac Lane theorem (Theorem 2.4.3) used in the proof of Theorem 9.2 to establish the contractible space of parenthesizations for the seven-fold optic composition.
- C. W. Misner, K. S. Thorne, and J. A. Wheeler. *Gravitation*. W. H. Freeman, 1973.
- E. Noether. Invariante variationsprobleme. *Nachrichten von der Gesellschaft der Wissenschaften zu Göttingen, Mathematisch-Physikalische Klasse*, 1918:235–257, 1918.
- P. J. Olver. *Applications of Lie Groups to Differential Equations*. Springer-Verlag, New York, second edition, 1993.
- E. Riehl and D. Verity. *Elements of ∞ -Category Theory*. Cambridge University Press, 2022. Develops the ∞ -cosmos framework providing model-independent foundations for ∞ -category theory; invoked in Definition 9.1 to construct the ∞ -categorical optic category **Optic**($\mathcal{C}_\infty$) as the cartesian fibration over $\text{Tw}(\mathcal{C}_\infty)$.
- M. Riley. Categories of optics. *arXiv preprint arXiv:1809.00738*, 2018.
- M. Riley. Cornering optics. *arXiv preprint arXiv:2307.16516*, 2023.
- J. J. M. M. Rutten. Universal coalgebra: A theory of systems. *Theoretical Computer Science*, 249(1):3–80, 2000. doi: 10.1016/S0304-3975(00)00056-6.
- M. Steel and W. Hordijk. Detecting autocatalytic, closed sets in biochemical reaction networks. *Journal of Theoretical Biology*, 231(4):467–475, 2004. doi: 10.1016/j.jtbi.2004.07.015.
- The Univalent Foundations Program. *Homotopy Type Theory: Univalent Foundations of Mathematics*. Institute for Advanced Study, 2013. URL <https://homotopytypetheory.org/book>. Canonical reference for homotopy type theory, the univalence axiom, and the universe type $\mathcal{U}$; invoked in Corollary 9.3 to canonically identify the contractible ∞ -groupoid of homotopy-fixed-points of the Banach contraction T .
- X. A measure-theoretic discrete curvature framework for metabolic gene sensitivity: From active-set geometry to transcriptional response, 2026. Application manuscript, in preparation.