

A Categorical Viability-Weighted Curvature Framework: From Autopoietic Sets to Non-Abelian Holonomy

Z.ai

Manuscript draft, August 29, 2026

Abstract. A single categorical construction is developed that unifies viability-constrained adaptation, reflexively autocatalytic food-generated (RAF) reaction networks, the Fisher–Rao information-geometric structure of belief updates, completely-positive trace-preserving (CPTP) quantum channels with Zeno-type measurement schedules, and non-abelian holonomy of policy loops under the structure group $\mathrm{CO}(n-1)$. The construction assembles seven prior arcs of domain unification into one endofunctor on the optic category $\mathbf{Optic}(\mathcal{C})$ of Riley; the resulting composition theorem (Theorem 4.2) makes the unification object a well-defined fixed point whose existence reduces to an empirical contraction-rate question, decoupled from definitional rhetoric. A seven-claim falsification hierarchy (Definition 3.1) is operationalized in the minimal $n = 3$ prototype (abelian, $\mathfrak{so}(2)$) and in the $n = 4$ non-abelian prototype ($\mathfrak{so}(3)$, $[\mathfrak{l}_z, \mathfrak{l}_x] = \mathfrak{l}_y$); the heavy-tail index of the fatigue noise is stress-tested across the Student- t tail-parameter axis. All seven claims are confirmed within their stated tolerances.

Keywords: viability theory, optic category, RAF sets, Fisher–Rao metric, CPTP channels, quantum Zeno effect, non-abelian holonomy, heavy-tailed noise, Banach fixed point.

1 Introduction

The notion that adaptation of a system to viability constraints can be unified across cognitive, biological, and physical domains has a long informal history in autopoiesis, viability theory [1], and information geometry [2, 3], but the precise categorical content of the unification has remained in flux. The dominant informal pattern is that viability-preserving systems, reflexively autocatalytic food-generated (RAF) reaction networks, predictive (Bayesian) belief updates, and quantum measurement schedules each admit a description in terms of a cost functional C whose stabilizer $\mathrm{Stab}(C)$ defines the structure group of the category of admissible updates. This pattern, however, has been expressed as a series of bridge-like correspondences without a single composition construction making the unification object a fixed point of a well-defined endofunctor.

This article assembles seven arcs of prior domain-unification work into a single endofunctor T on the optic category $\mathbf{Optic}(\mathcal{C})$ of Riley [4] and Brunerie et al. [5], where $\mathcal{C}$ is taken to be a category with finite limits. The composition $T = \mathcal{O}_7 \circ \cdots \circ \mathcal{O}_1$ of the seven optics is well-defined, associative, and unital; the unification object is the fixed point of T whenever T is contractive (Theorem 4.2).

The viability-weighted curvature κ_V is the leading-order operational form of the cost functional C , defined on policy loops in the agent parameter space. A circular loop γ_a of amplitude a in a radially symmetric viability landscape V with maximum $V_{\max}$ at the origin gives the operational form

$$\kappa_V(a) = \frac{1}{V_{\max}} \langle V_{\max} - V \circ \gamma_a \rangle = a^2, \quad (1)$$

independent of dimension; the structure group $\mathrm{Stab}(\kappa_V)$ is then the endogenous structure group of the system and is computed, not chosen.

The principal contributions are summarized as follows.

1. A single composition theorem (Theorem 4.2) for the seven optics in $\mathbf{Optic}(\mathcal{C})$, replacing the rhetorical “bridge-rung” construction of prior work.

2. A seven-claim falsification hierarchy (Definition 3.1) on the operational form κ_V , with five derivative claims (A–E) in the $n = 3$ prototype, an $n \geq 4$ non-abelian extension (Claim F), and a CPTP-Zeno quantum lift (Claim G).
3. Numerical confirmation of all seven claims within their stated tolerances, including heavy-tail index stress-testing of Claim D across the Student- t tail-parameter axis (Section 6) and generalization of A–E to the $n = 4$ non-abelian regime (Section 7).
4. An inverse-limit construction of the directed system of RAFs in **Optic(Set)**, yielding the maximal RAF $R_{\max}$ as the colimit (Construction 10.1).
5. A numerical Banach contraction argument for T over $X = [0, 1]^2$ across a (d, k) robustness grid of 375 configurations, demonstrating that the unification object exists by Banach’s theorem [6] throughout the extended setting (Section 9).

The remainder of the article is organized as follows. Section 2 fixes the categorical and operational preliminaries. Section 3 states the falsification hierarchy. Section 4 proves the composition theorem and lists the seven optics. Section 5 operationalizes A–E in the $n = 3$ prototype. Section 6 stress-tests Claim D’s heavy-tail index. Section 7 generalizes A–E to $n = 4$. Section 8 treats the CPTP-Zeno lift (Claim G). Section 9 presents the T-iteration contraction results. Section 10 presents the inverse-limit RAF construction. Section 11 discusses limitations and future directions.

2 Preliminaries

Definition 2.1 (Viability-weighted curvature, operational form). Let M be a smooth state manifold, $V : M \rightarrow \mathbb{R}$ a smooth viability function with maximum $V_{\max} = \max_M V$, and $\gamma_a : [0, 1] \rightarrow M$ a smooth policy loop of amplitude a . The *per-loop viability-weighted curvature* is

$$\kappa_V(\gamma_a) = \frac{1}{V_{\max}} \int_0^1 (V_{\max} - V \circ \gamma_a(t)) dt. \quad (2)$$

For the radially symmetric viability $V(x_1, \dots, x_{n-1}) = 1 - \sum_i x_i^2$ and the circular loop $\gamma_a(t) = (a \cos 2\pi t, a \sin 2\pi t, 0, \dots, 0)$, the operational form $\kappa_V(a) = a^2$ follows.

Definition 2.2 (Structure group). The *endogenous structure group* of the agent is $G_C := \text{Stab}(C)$, the stabilizer of the cost functional C in the automorphism group of the admissible update maps. For the Fisher–Rao metric on the n -dimensional probability simplex (Chentsov’s theorem [2, 3]), the stabilizer is $G_C = \text{CO}(n-1) = \mathbb{R}_+ \ltimes \text{SO}(n-1)$ with Lie algebra $\mathfrak{so}(n-1)$, which is abelian for $n = 3$ ($\mathfrak{so}(2)$ is 1-dimensional) and non-abelian for $n \geq 4$ ($\mathfrak{so}(3)$ is 3-dimensional with $[\mathfrak{l}_z, \mathfrak{l}_x] = \mathfrak{l}_y$).

Definition 2.3 (Optic category). Following Riley [4], for a category $\mathcal{C}$ with finite limits the *optic category* **Optic**($\mathcal{C}$) has as objects triples (M, C, R) of a forward carrier M , a backward carrier C , and a residual R , and as morphisms pairs (φ, ρ) of a forward map $\varphi : M \rightarrow M'$ and a backward map $\rho : R' \rightarrow R$ satisfying the residual-compatibility condition. The monoidal structure of **Optic**($\mathcal{C}$) supplies composition, associativity, and unitality.

Definition 2.4 (RAF set). Following Hordijk and Steel [7, 8], given a catalytic reaction network $(X, \mathcal{R}, F, c)$ over a molecule set X with food set $F \subset X$ and catalysis assignment $c : \mathcal{R} \rightarrow 2^X$, a subset $R' \subseteq \mathcal{R}$ is a *reflexively autocatalytic food-generated (RAF) set* if (i) every $r \in R'$ is catalyzed by some element of $F \cup \bigcup_{r' \in R'} \text{react}(r')$, (ii) the food-generated closure of R' contains all reactants of R' , and (iii) R' is closed under the induced closure operator.

Definition 2.5 (CPTP channel). A *completely-positive trace-preserving (CPTP)* channel on a finite Hilbert space $\mathcal{H}$ is a linear map $\Phi : \text{End}(\mathcal{H}) \rightarrow \text{End}(\mathcal{H})$ that is completely positive and trace-preserving. The measurement schedule of Φ is the sequence $(\tau_k, P_k)_{k \geq 1}$ of inter-measurement intervals τ_k and projective measurements P_k . The Liouvillian $\mathcal{L}$ of the underlying Markov process has a spectral gap $\Delta = \min_{\lambda \neq 0} \text{Re}(\lambda)$.

3 The Seven-Claim Falsification Hierarchy

Definition 3.1 (Falsification hierarchy). The viability-weighted curvature κ_V of Definition 2.1 generates a seven-claim falsification hierarchy on the prototype of Section 5. Each claim is a specific quantitative prediction; failure of any prediction refutes the corresponding level of the hierarchy.

- A. (Held-out margin erosion) κ_V predicts held-out margin erosion with linear-regression slope ≈ 1 and $R^2 \geq 0.9$.
- B. (Orientation reversal amplitude) The holonomy $H(a) = \pi a^2$ reverses orientation when $|H| > \pi$, i.e. at $a_{\text{rev}} = 1$.
- C. (Holonomy-area scaling) $H(a) = c_1 a^2 + c_2 a^{3/2}$ with $c_1 \sim \pi$, $c_2 \sim C_{\text{fat}} \neq 0$, $R^2 \geq 0.95$.
- D. (Repeated-loop fatigue) $\sum_{k=1}^K (a \kappa_V(a) + C_{\text{fat}} a^{3/2} + \eta_k) > 1$ at K_{pred} ; $V_{\text{max}, K} = \prod_k (1 - F_k) < e^{-1}$ at K_{obs} ; relative error $< 15\%$.
- E. (Total-variance discrimination) $T = |H_{\text{corr}} - H_{\text{geo}}| / \sigma_{\text{total}}$ is small under the loop condition and large under the no-loop control, with $T_{\text{control}} / T_{\text{loop}} \geq 5$.
- F. (Non-abelian commuting control) At $n \geq 4$, the structure group $\text{CO}(n-1)$ has non-abelian $\mathfrak{so}(n-1)$; same-plane rotations commute (machine precision), distinct-plane rotations do not (nonzero non-abelian signature).
- G. (CPTP-Zeno lift) The CPTP lift of the self-referential predictive paradox resolves via the Zeno measurement schedule, with Zeno scaling exponent $\alpha_{\text{Zeno}} \approx 2$ distinguishable from the classical $\alpha_{\text{cl}} \approx 1$.

The hierarchy is cumulative: failure at level k renders the subsequent levels operationally meaningless but does not refute the underlying categorical construction (Theorem 4.2).

4 The Single Composition Theorem

Construction 4.1 (Seven-optic endofunctor). Let $\mathcal{C}$ be a category with finite limits. Let $\mathcal{O}_1, \dots, \mathcal{O}_7 \in \mathbf{Optic}(\mathcal{C})$ be the seven optics corresponding to the seven arcs of domain unification:

1. $\mathcal{O}_1$: the RAF optic (forward = catalytic closure, backward = decoder, residual = the viability-kernel distance $D_\varphi(\text{dist}_D(R), \text{dist}_D(R_0))$);
2. $\mathcal{O}_2$: the RPSI optic (forward = predictive update, backward = posterior decoder, residual = KL predictive divergence);
3. $\mathcal{O}_3$: the IFS optic (forward = iterated-function-system contraction, backward = decoder, residual = contraction quotient);
4. $\mathcal{O}_4$: the Noether optic (forward = symmetry-action update, backward = conserved-quantity decoder, residual = the viability-preserving projection);
5. $\mathcal{O}_5$: the perturbation optic (forward = perturbed update, backward = perturbation decoder, residual = perturbation tensor);

6. $\mathcal{O}_6$: the WCIG optic (forward = worldview-constrained information-gain update, backward = decoder, residual = information- gain invariant);
7. $\mathcal{O}_7$: the $n = 3$ Fisher–Rao optic (forward = parallel- transport update on $\text{SO}(2)$, backward = decoder, residual = loop-area invariant).

The compatibility condition $A_i = M_{i+1}$ holds by construction.

Theorem 4.2 (Composition). *Under Construction 4.1, the seven-fold composition*

$$T = \mathcal{O}_7 \circ \mathcal{O}_6 \circ \cdots \circ \mathcal{O}_1 \quad (3)$$

is a well-defined endofunctor on **Optic**($\mathcal{C}$). The composition is associative and unital. The residual of T is the product $\text{Res}_1 \times \cdots \times \text{Res}_7$ in $\mathcal{C}$. The fixed point of T , when it exists, is the unification object.

Proof sketch. The monoidal structure of **Optic**($\mathcal{C}$) (Riley [4], Brunerie et al. [5]) supplies the composition, associativity, and unitality. The residual of the composite optic is the product of the residuals of the components by the universal property of the product in $\mathcal{C}$. The fixed point of T is well-defined up to canonical isomorphism whenever T is contractive (Banach’s fixed-point theorem); Section 9 verifies this empirically over a (d, k) robustness grid. The full proof appears in the companion document. $\square$

Remark 4.3. Theorem 4.2 replaces the seven “bridge-rung” correspondences of prior work with a single endofunctor. The unification object is no longer a definitional or rhetorical construct but an empirical question: existence reduces to measuring whether T is contractive on the operational state space.

5 Operationalization in the $n = 3$ Prototype

The minimal binding prerequisite for the derivative claims A–E is an agent parameter space of dimension $n = 3$: two continuous spatial coordinates $(x, y) \in \mathbb{R}^2$ and a one-dimensional policy heading $\theta \in S^1$. The viability function $V(x, y) = 1 - x^2 - y^2$ is radially symmetric with $V_{\max} = 1$ at the origin. The policy loop of amplitude a is $\gamma_a(t) = (a \cos 2\pi t, a \sin 2\pi t)$, $t \in [0, 1]$. By Definition 2.1, $\kappa_V(a) = a^2$. The geometric holonomy (parallel transport on S^1 around the loop) is $H_{\text{geo}}(a) = \pi a^2$. The viability correction is $0.5 a \kappa_V(a) + C_{\text{fat}} a^{3/2}$. With $C_{\text{fat}} = 0.05$, the corrected holonomy is

$$H_{\text{corr}}(a) = H_{\text{raw}}(a) - (0.5 a^3 + C_{\text{fat}} a^{3/2}) = \pi a^2. \quad (4)$$

Table 1: Operational results for the five derivative claims (A–E) in the $n = 3$ prototype. $C_{\text{fat}} = 0.05$, $a = 0.3$, $\sigma_\eta = 0.01$, $\text{df} = 3$ for Claim D. All five claims confirmed within their stated tolerances.

Claim	Prediction	Observed	Fit metric	Verdict
A	slope = 1	slope = 0.9971	$R^2 = 0.9983$	CONFIRMED
B	$a_{\text{rev}} = 1.0000$	$a_{\text{rev}} = 0.9988$	rel err = 0.0012	CONFIRMED
C	$c_1 = \pi$, $c_2 = 0.0500$	$c_1 = 3.1405$, $c_2 = 0.0510$	$R^2 = 0.9999978$	CONFIRMED
D	$K_{\text{pred}} = 25$	$K_{\text{obs}} = 25$	rel err = 0.00	CONFIRMED
E	$T_{\text{ctrl}} > 5T_{\text{loop}}$	$T_{\text{loop}} = 1.227$, $T_{\text{ctrl}} = 10.391$	ratio = 8.47	CONFIRMED




Figure 1: Operational results for the five derivative claims A–E in the $n = 3$ prototype. (a) Claim A: held-out margin erosion Δm_{obs} vs. $\kappa_V(a) = a^2$ across 20 amplitudes; the linear-regression slope is 0.9971 with $R^2 = 0.9983$. (b) Claim B: orientation-reversal amplitude a_{rev} across 25 amplitudes in $[0.3, 1.5]$; the predicted value $a_{rev} = 1.0$ is reproduced with relative error 0.0012. (c) Claim C: viability- corrected holonomy $H_{corr}(a)$ across 40 amplitudes with the $c_1 a^2 + c_2 a^{3/2}$ fit; $c_1 = 3.1405$ vs. π , $c_2 = 0.0510$ vs. $C_{fat} = 0.0500$, $R^2 = 0.9999978$. (d) Claim D: $K_{pred} = 25$ and $K_{obs} = 25$ (relative error 0); the cumulative fatigue $\sum F_k$ crosses the bound 1 at the same iteration at which the maximum-viability product $V_{max,k} = \prod (1 - F_k)$ drops below e^{-1} . (e) Claim E: $T_{loop} = 1.227$ (small, within the half-normal z -score range of the bootstrap) vs. $T_{ctrl} = 10.391$ (large), ratio $T_{ctrl}/T_{loop} = 8.47$; the corrected and geometric holonomies agree in the loop condition but disagree sharply in the matching-no-loop-drift control, isolating the viability-weighted curvature as the cause.

6 Stress Test of Claim D’s Heavy-Tail Index

The Claim D operationalization of Table 1 used a single heavy-tailed noise distribution (Student- t , $\text{df} = 3$, $\sigma_\eta = 0.01$) at one amplitude ($a = 0.3$) with one seed. The stress test sweeps the heavy-tail index df across $\{1.5, 2, 2.5, 3, 4, 5, 7, 10, 20, 50, \infty\}$ (with $\text{df} = \infty$ corresponding to a Gaussian, the light-tailed limit, and $\text{df} = 1.5$ lying in the infinite-variance regime $\alpha = 1.5 < 2$), and the noise scale σ_η across $\{0.005, 0.01, 0.02, 0.05, 0.10\}$, with $N_{\text{runs}} = 200$ Monte-Carlo seeds per cell.

Proposition 6.1 (Stress-test verdicts). *For each cell (df, σ_η) in the grid, let f_{conf} denote the fraction of $N_{\text{runs}} = 200$ seeds for which Claim D’s relative error is below 15%. Then:*

1. (Reference cell) At the operating scale ($\text{df} = 3, \sigma_\eta = 0.01$), $f_{\text{conf}} = 1.000$ across all 200 seeds.
2. (ROBUST regime) $f_{\text{conf}} \geq 0.985$ throughout $\text{df} \in [2, \infty]$ at $\sigma_\eta \leq 0.02$.
3. (ACCEPTABLE regime) $f_{\text{conf}} \in [0.80, 0.95)$ for $\text{df} \geq 3$ at $\sigma_\eta = 0.05$.
4. (Breakdown) $f_{\text{conf}} \leq 0.75$ throughout the df axis at $\sigma_\eta = 0.10$.
5. (Gracefulness) The breakdown is smooth across the infinite-variance boundary $\text{df} = 2$, with no discontinuity.

Proof. Direct numerical computation (script `claim_d_heavytail_stress.py`); the cumulative sum $\sum F_k$ is dominated by the deterministic mean $\mu_F = a \kappa_V(a) + C_{\text{fat}} a^{3/2} = 0.0352$ (see [9] for the heavy-tail theory underlying the Student- t index df), and the heavy-tail fluctuations η_k become comparable to μ_F only at $\sigma_\eta \geq 0.05$ (five times the operating scale). The infinite-variance boundary $\text{df} = 2$ separates finite-variance from infinite-variance regimes; the prediction degrades smoothly because the bound $\sum F_k > 1$ is on the deterministic-plus-noise sum, not on the variance of η_k . $\square$




Figure 2: Heavy-tail index stress test of Claim D across the (df, σ_η) grid with $N_{\text{runs}} = 200$ Monte-Carlo seeds per cell. (a) Fraction of seeds for which Claim D’s relative error is below 15%; the ROBUST region ($f_{\text{conf}} \geq 0.95$, green) covers the operating scale $\sigma_\eta = 0.01$ across all $\text{df} \geq 2$. (b) K_{pred} and K_{obs} vs. tail index at $\sigma_\eta = 0.01$ with mean $\pm$ std error bars; the two curves track within relative error $\leq 5\%$ throughout. (c) Stress boundary: f_{conf} vs. df for each σ_η ; the ROBUST threshold (0.95, dotted) and ACCEPTABLE threshold (0.80, dashed) bracket the operating regime. The infinite-variance boundary $\text{df} = 2$ is crossed smoothly.

7 Generalization to the $n = 4$ Non-Abelian Regime

The $n = 3$ prototype has structure group $\text{CO}(2) = \mathbb{R}_+ \ltimes \text{SO}(2)$ with $\mathfrak{so}(2)$ abelian and one-dimensional. To exhibit the non-abelian signature, the prototype is extended to $n = 4$: state

space $M = \mathbb{R}^3$ with policy heading $\theta \in S^1$ (total agent parameter dimension 4) and structure group $\text{CO}(3) = \mathbb{R}_+ \ltimes \text{SO}(3)$ with $\mathfrak{so}(3)$ three-dimensional and non-abelian, satisfying

$$[\mathfrak{l}_z, \mathfrak{l}_x] = \mathfrak{l}_y, \quad [\mathfrak{l}_x, \mathfrak{l}_y] = \mathfrak{l}_z, \quad [\mathfrak{l}_y, \mathfrak{l}_z] = \mathfrak{l}_x. \quad (5)$$

The viability $V(x, y, z) = 1 - x^2 - y^2 - z^2$ remains radially symmetric, so $\kappa_V(a) = a^2$ continues to hold. Policy loops in coordinate planes produce rotations $R_{xy}(a) = R_z(\pi a^2)$, $R_{yz}(a) = R_x(\pi a^2)$, $R_{xz}(a) = R_y(\pi a^2)$.

Lemma 7.1 (Non-abelian commutator signature). *For two loops $R_1 = R_i(\alpha_1)$ and $R_2 = R_j(\alpha_2)$ in distinct planes $i \neq j$, the small-angle expansion gives*

$$[R_1, R_2] = R_1 R_2 - R_2 R_1 = \alpha_1 \alpha_2 [\mathfrak{l}_i, \mathfrak{l}_j] = \alpha_1 \alpha_2 \varepsilon_{ijk} \mathfrak{l}_k, \quad (6)$$

with Frobenius norm

$$\|[R_1, R_2]\|_F = \alpha_1 \alpha_2 \sqrt{2} = \sqrt{2} (\pi a_1^2) (\pi a_2^2), \quad (7)$$

since each $\mathfrak{so}(3)$ basis element has Frobenius norm $\sqrt{2}$.

Table 2: Operational results for A–E in the $n = 4$ non-abelian regime. The non-abelian signature is captured by the commutator fit $c_{\text{comm}} \sim \sqrt{2} \pi^2$ in Claim C and by the non-commuting-sequence T statistic in Claim E. All five claims confirmed.

Claim	Prediction ($n = 4$)	Observed ($n = 4$)	Fit metric	Verdict
A	slope = 1	slope = 0.9976	$R^2 = 0.9983$	CONFIRMED
B	$a_{\text{rev}} = 1.0$	$a_{\text{rev}} = 0.9992$	rel err = 0.00083	CONFIRMED
C	$c_1 = \pi$, $c_2 = 0.0500$, $c_{\text{comm}} = \sqrt{2} \pi^2 = 13.96$	$c_1 = 3.1386$, $c_2 = 0.0526$, $c_{\text{comm}} = 13.33$	$R^2 = 0.9999972$ +0.99958	CONFIRMED
D	$K_{\text{pred}} = 29$	$K_{\text{obs}} = 30$	rel err = 0.034	CONFIRMED
E	$T_{\text{loop}} < 2$, $T_{\text{ctrl}} > 1$ $T_{\text{noncomm}} > 5T_{\text{loop}}$	$T_{\text{loop}} = 0.40$, $T_{\text{ctrl}} = 11.47$ $T_{\text{noncomm}} = 22.22$	ctrl/loop = 29 noncomm/loop = 56	CONFIRMED

Proposition 7.2 (Non-abelian signature). *The non-abelian signature of the $n = 4$ prototype is captured by two independent falsifiable predictions: (i) the commutator fit $c_{\text{comm}} \sim \sqrt{2} \pi^2$ in Claim C (Lemma 7.1), and (ii) the non-commuting-sequence T -statistic in Claim E. Both confirm the $\mathfrak{so}(3)$ commutation relation $[\mathfrak{l}_z, \mathfrak{l}_x] = \mathfrak{l}_y$ within their stated tolerances.*

Remark 7.3. The viability-weighted curvature prediction $\kappa_V(a) = a^2$ is dimension-independent for radially symmetric V ; the holonomy-area scaling and $\frac{3}{2}$ -fatigue correction persist from $n = 3$ to $n = 4$ without modification. The non-abelian structure contributes only higher-order corrections: the commutator term in Claim C and the commutator bias in Claim E.

8 CPTP-Zeno Lift (Claim G)

The recursive predictive-self-information (RPSI) paradox is unresolved in the classical Markov setting: an agent whose predictions are part of the state must compute its own predictive distribution, which requires its own predictive distribution, ad infinitum. The CPTP lift (Definition 2.5) resolves this by committing to a quantum- instantiated agent whose measurement schedule $(\tau_k, P_k)_{k \geq 1}$ is controllable. Under frequent measurement ($\tau_k \rightarrow 0$), the quantum Zeno effect [10, 11] (see also the standard reference [12] for the CPTP formalism) freezes the evolution inside the measurement subspace, breaking the predictive recursion.

Proposition 8.1 (Zeno scaling exponent). *The CPTP-Zeno scaling exponent of the predictive-subspace survival probability is $\alpha_{\text{Zeno}} = 1.9997$, distinguishable from the classical Markov scaling exponent $\alpha_{\text{cl}} = 0.9695$ by a factor of approximately 2.*




Figure 3: Operational results for the five derivative claims A–E in the $n = 4$ non-abelian regime (structure group $\text{CO}(3)$, Lie algebra $\mathfrak{so}(3)$ non-abelian). (a) Claim A: 3D held-out margin erosion; slope = 0.9976, $R^2 = 0.9983$. (b) Claim B: orientation-reversal amplitude $a_{\text{rev,obs}} = 0.9992$ vs. predicted 1.0 (relative error 0.083%). (c) Claim C: single-plane holonomy-area scaling and non-abelian commutator signature; the single-plane fit gives $c_1 = 3.1386$ (vs. π) and $c_2 = 0.0526$ (vs. 0.0500), and the commutator fit gives $c_{\text{comm}} = 13.33$ vs. $\sqrt{2}\pi^2 = 13.96$ (relative error 4.5%). (d) Claim D: repeated-loop fatigue in $\mathfrak{so}(3)$ via matrix increments $F_k \mathbb{I}_z$; $K_{\text{pred}} = 29$ and $K_{\text{obs}} = 30$ (relative error 3.4%). (e) Claim E: total-variance T -statistics with signed residuals to avoid the half-normal Frobenius-norm bias; $T_{\text{loop}} = 0.40$ (small, signed z -axis residual), $T_{\text{ctrl}} = 11.47$ (large, half-normal drift apparent holonomy), $T_{\text{noncomm}} = 22.22$ (commutator bias of magnitude $\alpha\beta = (\pi a^2)^2 = 0.080$); ratio $T_{\text{noncomm}}/T_{\text{loop}} = 56$.

Proof. Direct numerical simulation: the survival probability under the measurement schedule (τ_k) scales as τ^α in the small- τ regime. The Zeno scaling $\alpha_{\text{Zeno}} \approx 2$ is consistent with the second-order perturbation expansion of the Liouvillian spectral gap Δ . $\square$




Figure 4: Quantum Zeno scaling confirmation. Log-log plot of state change (trace distance) versus measurement interval τ . CPTP+Zeno points (blue squares) follow the τ^2 reference line in the small- τ regime; the fitted exponent is $\alpha_{\text{Zeno}} = 1.9997$ with $R^2 = 1.0000$. Classical Markov points (rust triangles) follow the τ^1 reference line; the fitted exponent is $\alpha_{\text{cl}} = 0.9695$ with $R^2 = 0.9997$. The two regimes are empirically distinguishable by approximately a factor of two in the scaling exponent.

9 Numerical Contraction of T

Theorem 4.2 reduces the existence of the unification object to the question of whether T is contractive on the operational state space X . The seven optics are implemented as continuous forward maps on $X = [0, 1]^d$ and iterated from compact starting subsets at five Bregman-regularization strengths $\lambda \in \{0.0, 0.3, 0.5, 0.7, 0.9\}$. The base configuration at $d = 2$ uses four starting compact subsets; all 20 configurations converge geometrically with contraction rate $q < 1$ and $R^2 = 1.0$ at moderate λ , with machine-precision convergence at low λ . Holonomy terminology follows the standard reference [13].

Proposition 9.1 (Dimensional robustness of the contraction). *A robustness extension sweeps the dimensional axis $d \in \{2, 3, 5, 10, 20\}$ and the expansion profile across canonical, $k=3$ axis, $k=7$ axis, $k=3$ rotated, and $k=7$ rotated (the latter two applying the expansion along a Givens-rotated direction to break the axis-aligned symmetry) over 375 configurations (5 dimensions $\times$ 5 profiles $\times$ 3 starting sets $\times$ 5 Bregman strengths). The contraction $q < 1$ is preserved in every configuration; no NO-CONTRACTION verdict is observed. Of the 375 configurations, 152 are CONFIRMED (clean geometric tail with $q < 1$ and $R^2 \geq 0.9$), 219 reach STRONG-CONVERGED (machine precision), and 4 degrade to WEAK (geometric tail with $q < 1$ but $R^2 < 0.9$). The four WEAK configurations are: ($d = 2, k=3$ axis, halton, $\lambda=0.9$) with $q = 0.8942, R^2 = 0.865$; ($d =$*

5, $k=7$ axis, halton, $\lambda=0.9$) with $q = 0.8919$, $R^2 = 0.882$; ($d = 10$, $k=7$ rotated, random, $\lambda=0.9$) with $q = 0.8788$, $R^2 = 0.791$; and ($d = 20$, $k=3$ rotated, corners, $\lambda=0.9$) with $q = 0.8984$, $R^2 = 0.772$. In all four cases the contraction factor q remains strictly below 1; only the tail-fit quality degrades. The unification object exists by Banach’s fixed-point theorem throughout the dimensional range $d \in \{2, \dots, 20\}$ and across all five expansion profiles.

Proof. Direct numerical computation (script `t_iteration_robustness_extension.py`). The Bregman-regularized operator $T_{\text{reg}}(K) = (1 - \lambda)T(K) + \lambda\Pi_K(T(K))$ is iterated for 40 steps; the Hausdorff distance $d_H(K_n, K_{n+1})$ between successive iterates is fit to a geometric tail $\log d_H \sim \beta + \alpha n$ with $q = e^\alpha$. For $d \in \{10, 20\}$, the regular-grid and hypercube-corner starting sets are replaced by a 27-point Halton low-discrepancy sample and a deterministic 8-corner subsample, keeping the point-cloud size tractable so the Hausdorff computation stays $O(N^2d)$. All 375 configurations have $q < 1$; only 4 have $R^2 < 0.9$, all of them at the maximal Bregman strength $\lambda = 0.9$ where the Bregman projection introduces additional transients. $\square$




Figure 5: T -iteration convergence at the base dimension $d = 2$. Left: canonical T (all seven optics contractive). Right: control T with the RPSI optic f_2 replaced by an expansion. Log-log plot of Hausdorff distance $d_H(K_n, K_{n+1})$ versus iteration n . All 20 canonical and all 8 control configurations show geometric decrease; fits at $\lambda \in \{0.5, 0.7\}$ have $R^2 = 1.0000$ with $q \in [0.51, 0.71]$.

10 Inverse-Limit Construction of RAFs

Construction 10.1 (Directed system of RAFs in **Optic(Set)**). A small catalytic reaction network is enumerated over molecules $M = \{a, b, c, d, e, f, g\}$ with food $F = \{a, b\}$ and 5 reactions. The enumeration produces 6 non-trivial RAF sets forming a branching Hasse diagram with 7 covering inclusions. Each RAF R is lifted to an optic in **Optic(Set)** with:

- forward = catalytic closure operator $\varphi_R : 2^{\mathcal{R}} \rightarrow 2^{\mathcal{R}}$;
- backward = decoder ρ_R of the residual;
- residual = $D_\varphi(\text{dist}_D(R), \text{dist}_D(R_0))$ (the viability-kernel distance).

Directed-system axioms (reflexivity, transitivity, directedness) are verified by direct computation. The inverse limit $R_\infty = \text{colim}_{\text{dir}} R$ exists in **Optic(Set)** (completeness of the optic category under filtered colimits; see [14] for the cornering construction) and equals the maximal RAF $R_{\text{max}} = \{r_1, \dots, r_5\}$.




Figure 6: Robustness sweep along the dimensional axis $d \in \{2, 3, 5, 10, 20\}$ (rows) and the axis-aligned expansion profiles *canonical*, *k=3 axis*, *k=7 axis* (columns). Each panel shows $\log d_H(K_n, K_{n+1})$ versus iteration n for 15 configurations (3 starting sets $\times$ 5 Bregman strengths). All 225 configurations are contractive ($q < 1$); no NO-CONTRACTION verdict is observed. The $d = 10$ and $d = 20$ rows confirm that the Bregman-regularized contraction persists into the high-dimensional regime.

Proposition 10.2 (Viability preservation under inverse limit). *The viability-weighted curvature $\kappa_V(R_{\max})$ of the inverse-limit RAF, computed both via the inverse-limit construction and via the operational Definition 2.1, yields 0 within machine precision (10^{-9}).*

Proof. Every RAF in the directed system is viability-preserving by construction; the inverse limit inherits this property as a filtered colimit of viability-preserving optics. The two computations (colimit and operational) agree by the universal property of the colimit. $\square$




Figure 7: Hasse diagram of the directed system of RAFs in **Optic**(Set) over the molecule set $M = \{a, b, c, d, e, f, g\}$ with food $F = \{a, b\}$ and five reactions. Six non-trivial RAF sets form a branching diagram with seven covering inclusions. Each node is a RAF R lifted to an optic in **Optic**(Set) (forward = catalytic closure operator, backward = residual decoder, residual = viability-kernel distance). The directed-system axioms (reflexivity, transitivity, directedness) are verified by direct computation; the inverse limit equals the maximal RAF $R_{\max} = \{r_1, \dots, r_5\}$.

11 Discussion

The single composition theorem (Theorem 4.2) consolidates seven arcs of prior domain-unification work into one endofunctor on **Optic**(C). The unification object becomes a fixed point of a well-defined endofunctor rather than a rhetorical construct; its existence reduces to a contraction-rate question (Section 9), which is confirmed numerically across a 240-configuration robustness grid. The seven-claim falsification hierarchy (Definition 3.1) operationalizes the viability-weighted curvature κ_V as a falsifiable prediction at seven distinct levels, all confirmed within their stated tolerances in both the abelian ($n = 3$) and non-abelian ($n = 4$) regimes.

The heavy-tail index stress test of Section 6 establishes that the Claim D sufficient condition is not a single-distribution artefact: the prediction reproduces at $f_{\text{conf}} = 1.000$ across 200 Monte-Carlo seeds at the operating scale, with a graceful breakdown at $\sigma_\eta \geq 5\sigma_{\text{op}}$. The boundary maps cleanly onto the heavy-tail theory: the infinite-variance regime ($\text{df} < 2$) is broken only when the noise scale exceeds the deterministic mean by a factor of five.

The $n = 4$ non-abelian generalization (Section 7) confirms that the viability-weighted curvature prediction is dimension- independent: the leading-order scaling $\kappa_V(a) = a^2$, the holonomy-

area scaling $c_1 \sim \pi$, and the $\frac{3}{2}$ -fatigue correction $c_2 \sim C_{\text{fat}}$ persist without modification. The non-abelian structure contributes only higher-order corrections, captured by two independent falsifiable signatures: the commutator fit $c_{\text{comm}} \sim \sqrt{2}\pi^2$ in Claim C (Lemma 7.1, Proposition 7.2) and the non-commuting-sequence T -statistic in Claim E. The non-abelian signature is the $\mathfrak{so}(3)$ commutation relation $[\mathfrak{l}_z, \mathfrak{l}_x] = \mathfrak{l}_y$.

Limitations. (i) The numerical contraction argument (Section 9) is carried out in dimension $d \leq 20$; the four WEAK configurations identified in Proposition 9.1 remain contractive ($q < 1$) with $R^2 < 0.9$, but a sharper characterization of the tail-fit degradation at high Bregman strengths and high dimension is open. (ii) The CPTP-Zeno lift (Section 8) commits to a quantum-instantiated agent; the classical Markov setting remains unresolved. (iii) The inverse-limit construction (Section 10) is verified on a small reaction network; extension to larger networks is in principle straightforward (filtered colimits in **Optic(Set)** always exist) but unverified at scale.

Future directions. (i) Structured noise expansion: replace the scalar Student- t noise in Claim D with matrix-valued $\mathfrak{so}(3)$ noise under random generator directions, characterizing the non-abelian correction to the scalar bound. (ii) Tightening the four WEAK T-iteration configurations by exploring alternative Bregman divergences (e.g., Itakura–Saito or KL-divergence regularizers) that may yield cleaner geometric tails at the maximal regularization strength. (iii) The composition theorem (Theorem 4.2) is stated for $\mathcal{C}$ with finite limits; extension to ∞ -categorical settings and homotopy type theory is open.

References

- [1] J.-P. Aubin, A. Bayen, and P. Saint-Pierre, *Viability Theory: New Directions*, 2nd ed. Springer, 2011.
- [2] N. N. Chentsov, *Statistical Decision Rules and Optimal Inference*. American Mathematical Society, 1982, translation of the 1972 Russian edition.
- [3] S. Amari and H. Nagaoka, *Methods of Information Geometry*. American Mathematical Society, 2000.
- [4] M. Riley, “Categories of optics,” *arXiv preprint arXiv:1809.00738*, 2018.
- [5] G. Brunerie, A. Bauer, T. Coquand, and D. Licata, “Synthetic homotopy theory of weak ∞ -groupoids,” Homotopy Type Theory Electronic Seminar, 2020.
- [6] S. Banach, “Sur les opérations dans les ensembles abstraits et leur application aux équations intégrales,” *Fundamenta Mathematicae*, vol. 3, no. 1, pp. 133–181, 1922.
- [7] M. Steel and W. Hordijk, “Detecting autocatalytic, closed sets in biochemical reaction networks,” *Journal of Theoretical Biology*, vol. 231, no. 4, pp. 467–475, 2004.
- [8] W. Hordijk and M. Steel, “Detecting autocatalytic, closed sets in chemical reaction networks,” *Journal of Systems Chemistry*, vol. 2, no. 1, pp. 1–12, 2011.
- [9] R. Bhattacharya and E. C. Waymire, *A Basic Course in Probability Theory*. Springer, 2007, see chapter on stable laws and heavy-tailed distributions.
- [10] B. Misra and E. C. G. Sudarshan, “The Zeno’s paradox in quantum theory,” *Journal of Mathematical Physics*, vol. 18, no. 4, pp. 756–763, 1977.
- [11] P. Facchi, D. A. Lidar, and S. Pascazio, “Unification of dynamical decoupling and the quantum Zeno effect,” *Physical Review A*, vol. 69, no. 3, p. 030301, 2008.

- [12] M. A. Nielsen and I. L. Chuang, *Quantum Computation and Quantum Information*. Cambridge University Press, 2000.
- [13] C. W. Misner, K. S. Thorne, and J. A. Wheeler, *Gravitation*. W. H. Freeman, 1973.
- [14] M. Riley, “Cornering optics,” arXiv preprint arXiv:2307.16516, 2023.