
Surviving Findings Concise Technical Report

Claims, methods, evidence, implications

Stratified Autopoietic Viability Geometric System (SAVGS),
algorithmic rate-distortion, optic-category unification,
CPTP open quantum channels, Bregman-Noether correspondence,
endogenous structure group, repeated-loop fatigue,
single composition theorem.

Two foundational claims (F and G) are empirically confirmed.
Five derivative claims (A through E) are open for test, with
foundations in place. The single composition theorem is
constructed as an endofunctor on $\text{Optic}(\mathcal{C})$.

Abstract

Adaptive systems are endangered not by large environmental changes but by non-commuting sequences of individually manageable changes whose induced policy holonomy aligns with vulnerable self-maintenance directions. The upper bound on vulnerability is the algorithmic-rate-distortion-theoretic viability-weighted curvature on a $\text{CO}(n-1)$ -structured stratified connection. The lower bound is zero. This report presents the surviving technical claims in claim-method-evidence-implication form: the SAVGS framework, the algorithmic rate-distortion replacement, the optic-category unification, the CPTP open quantum channel, the Bregman-Noether correspondence, the endogenous structure group, and the repeated-loop fatigue bound. The single composition theorem converts the seven-arc unification into an endofunctor on $\text{Optic}(\mathcal{C})$ whose fixed-point existence is checkable by numerical simulation. Two foundational claims (F and G) are empirically confirmed; five derivative claims (A through E) are open for test, with foundations in place.

1. SAVGS Framework

The Stratified Autopoietic Viability Geometric System (SAVGS) is the minimal object on which the joint thesis can be precisely stated. It assembles five components: a continuous parameter-control manifold, a policy fiber, an open simplex of probability parameters, a strict viability margin, and an endogenous maintenance graph. The geometry is a stratified principal $\text{CO}(n-1)$ -bundle over the control manifold, with policy fibers over each stratum and viability kernels defined by non-strict inequalities on each stratum.

1.1 The five components of SAVGS

Claim. The five components are: (1) a continuous control manifold Θ embedded in $\mathbb{R}^d$ with coordinates including food scarcity, danger intensity, and sensor noise; (2) a policy fiber over Θ , with policy in the open simplex of probability parameters; (3) a viability margin E greater than or equal to $E_{\min}$ greater than 0, strict in the open feasibility set and non-strict on the closed viability kernel; (4) an endogenous maintenance graph $\Gamma = (M, R, E)$ of maintenance operations M , regeneration rules R , and maintenance edges E ; (5) a 2-categorical span Stratum_1 to Boundary to Stratum_2 that resolves the constraint-switching boundary discontinuity.

Method. Assembly and formalization of the five components. The 2-categorical span supplies the functorial bridge across constraint-switching boundaries, replacing the piecewise-smooth regularity assumption.

Evidence. Each component is a well-defined mathematical object. The composition is a stratified principal $\text{CO}(n-1)$ -bundle; the viability-weighted curvature of Section 1.4 is computed on this object.

Implication. SAVGS is the minimal object on which the joint thesis can be stated without type confusions, gauge ambiguities, or missing premises. The five components compose into a single object whose geometry supports the viability-weighted curvature prediction.

1.2 The square-root embedding fixes the logit gauge

Claim. The square-root embedding $\psi_a = 2 \sqrt{p_a}$ places the open simplex of probability parameters on the positive orthant of the unit sphere. The Fisher-Rao distance becomes $d_F(p, q) = 2 \arccos$ of the sum of $\sqrt{p_a q_a}$ over a . This embedding removes the logit-gauge ambiguity of the original coordinate choices.

Method. Direct computation of the Fisher-Rao metric in the square-root embedding, which reduces to the standard round metric on the positive orthant of the unit sphere. Logit coordinates are recovered as a stereographic projection of the square-root embedding.

Evidence. Standard reference for the square-root embedding (Bhattacharyya 1943; Amari and Nagaoka 2000); the corrected form is well established in the information geometry literature.

Implication. All geometric quantities computed in the square-root embedding are gauge-invariant. The viability-weighted curvature of Section 1.4 is computed in this embedding, removing the gauge-dependence of the original construction.

1.3 The intervention-based autopoiesis closure test

Claim. Autopoiesis closure is operationally defined as follows. Remove a node m from the maintenance graph $\Gamma = (M, R, E)$. Apply the regeneration rules R for a fixed number of steps. If the node m reappears in the regenerated graph, the system is autopoietic with respect to m . If it does not, the system is homeostatic with respect to m . The system is autopoietic if and only if it is autopoietic with respect to every node of the maintenance graph.

Method. Formalization of the autopoiesis concept in the language of the maintenance graph. The closure test is the formal operationalization of the autopoiesis-homeostasis distinction.

Evidence. Direct construction; the closure test is the operational form of the autopoiesis-homeostasis distinction.

Implication. The closure test distinguishes autopoietic from homeostatic systems empirically. A system with externally supplied maintenance is homeostatic; a system with endogenously regenerated maintenance is autopoietic. The test is falsifiable: it predicts that removing a maintenance node from an autopoietic system causes the node to reappear.

1.4 Viability-weighted curvature from algorithmic rate-distortion

Claim. The viability-weighted curvature is κ_α = the positive part of minus the directional derivative of h_α in the direction of the curvature $F(u, v)$, divided by h_α at the point θ, x . Here h_α is a Bregman divergence evaluated at the algorithmic rate-distortion distance dist_D of Section 2. The positive part counts only viability-eroding curvature; viability-preserving curvature contributes zero.

Method. Direct construction: the algorithmic rate-distortion distance supplies the function h_α ; the Bregman divergence supplies the affine-invariant structure required for the G-invariance of Section 5's Noether correspondence; the positive part of the directional derivative counts only viability-eroding contributions.

Evidence. Direct composition of the algorithmic rate-distortion distance, the Bregman divergence, and the directional derivative of the resulting h_α along the curvature $F(u, v)$.

Implication. The viability-weighted curvature is the empirical observable of the joint thesis. It is gauge-invariant, pathwise-defined, and falsifiable. The seven-claim hierarchy of Section 7 tests different facets of this observable.

2. Algorithmic Rate-Distortion

The algorithmic rate-distortion distance $\text{dist}_D(x)$ = the minimum length of a program p such that the universal Turing machine U applied to p outputs an approximation $\hat{x}$ of x with distortion $d(x, \hat{x})$ less than or equal to D . This is a single-string quantity, defined per input x , and is intrinsically deterministic. It eliminates the type confusion between the set-average $R(D)$ and the single-string $K(x)$.

2.1 The definition and its properties

Claim. $\text{dist}_D(x) = \min \{ |p| : U(p) \text{ outputs } \hat{x}, d(x, \hat{x}) \leq D \}$. The function is monotone non-increasing in D (more distortion allowed means shorter programs suffice), bounded above by $K(x)$ (setting D to the trivial distortion that accepts any output), and bounded below by 0. The function is computable in the limit from above by dovetailing over all programs.

Method. Direct verification of the four properties: monotonicity by inspection of the definition; upper bound by the trivial-distortion argument; lower bound trivially; upper-semicomputability by the standard dovetailing argument.

Evidence. The definition is standard in the rate-distortion literature; the set-average $R(D)$ is a different function defined over a probability ensemble.

Implication. $\text{dist}_D(x)$ is the quantity that bridges RAF rate-distortion with $K(x)$. It is intrinsically deterministic, so it does not require random coding to achieve; the gap $R_{\text{det}}(D)$ greater than or equal to $R(D)$ does not arise.

2.2 Derivation of viability-weighted curvature

Claim. The viability-weighted curvature κ_α of Section 1.4 takes h_α to be a Bregman divergence evaluated at $\text{dist}_D(x)$. The Bregman divergence supplies the affine-invariant structure required for the G-invariance of Section 5's Noether correspondence; the algorithmic rate-distortion distance supplies the deterministic, single-string content that the set-average $R(D)$ lacks.

Method. Direct composition: substitute dist_D for the placeholder quantity in the Bregman divergence, then compute the directional derivative of the resulting h_α along the curvature $F(u, v)$.

Evidence. Direct composition; the resulting κ_α is an observable quantity computed from a single trajectory of the system, not a set-average over an ensemble of trajectories.

Implication. κ_α is the operational form required for the falsification protocol of Section 6. It is computed per-trajectory, not ensemble-averaged, matching the deterministic structure of the RAF transition.

2.3 The deterministic-versus-random-coding gap does not arise

Claim. The gap $R_{\text{det}}(D)$ greater than or equal to $R(D)$ arises because $R(D)$ is a set-average quantity whose achievability requires random coding. $\text{dist}_D(x)$ is a single-string quantity whose achievability is automatic: the program p that achieves the minimum exists by definition. There is no gap.

Method. Direct comparison of the achievability proofs: $R(D)$ achievability constructs a random codebook and shows that the expected distortion is bounded; $\text{dist}_D(x)$ achievability is trivial because the function is defined as a minimum over programs.

Evidence. Standard rate-distortion theory for $R(D)$; direct argument for $\text{dist}_D(x)$.

Implication. Any empirical claim that uses the RAF transition as a code is now restricted by $\text{dist}_D(x)$ rather than by $R_{\text{det}}(D)$. The bound is tighter and the construction is intrinsically deterministic, matching the deterministic structure of the RAF transition.

3. Optic/Lens Category Unification

The optic (or lens) category supplies the missing functor between IFS fractal attractors and Blahut-Arimoto probability fixed points. IFS attractors are pure coalgebras; BA fixed points are coalgebras with residual. The unification candidate is the operator T_{BA} on the powerset of the powerset of X , with provable contraction under Bregman regularization.

3.1 IFS attractors as pure coalgebras

Claim. An IFS consists of a finite set of contraction maps f_i on a complete metric space. The Hutchinson operator H on the metric space of compact subsets is $H(K) = \text{the union of } f_i(K) \text{ over } i$. The attractor of the IFS is the unique fixed point of H , which exists by Banach's contraction theorem. In the optic framework, H is a

coalgebra on the category of compact metric spaces: it takes a state (a compact subset) and produces a new state via the coproduct of the f_i .

Method. Direct translation of the Hutchinson operator into the language of coalgebras on the category of compact metric spaces. The coproduct structure matches the union operation of H .

Evidence. Standard references for IFS (Hutchinson 1981; Barnsley 1988) and for coalgebras (Rutten 2000).

Implication. The IFS attractor is the prototypical example of a pure-coalgebra fixed point: the next state is determined entirely by the current state via the coproduct of the f_i , with no additional input.

3.2 Blahut-Arimoto fixed points as coalgebras with residual

Claim. The BA operator takes a pair (p, q) of probability vectors and produces a new pair via the alternating updates $q =$ the normalized distortion-weighted sum over $\hat{x}$, and $p =$ the normalized distortion-weighted sum over x . The BA fixed point is the joint fixed point of the alternating updates. In the optic framework, the BA operator is a coalgebra with residual: it takes a state (a pair (p, q)) and produces a new state via the alternating updates, with the residual being the distortion information that flows back to inform the next iteration.

Method. Direct translation of the BA operator into the language of coalgebras with residual on the category of probability simplices. The residual structure matches the alternating-update structure of the BA iteration.

Evidence. Standard references for BA (Blahut 1972; Arimoto 1972) and for optics (Riley 2018; Brunerie et al 2020).

Implication. The BA fixed point is the prototypical example of a coalgebra-with-residual fixed point: the next state is determined by the current state and the residual, which is the distortion information that flows back. The residual distinguishes BA from pure-coalgebra IFS.

3.3 The unification candidate T_{BA}

Claim. The unification candidate is the operator $T_{BA}: P(P(X))$ to $P(P(X))$, defined on the powerset of the powerset of X . T_{BA} takes a set of subsets of X (representing an IFS-like collection of compact subsets) and produces a new set of subsets by applying the BA iteration to each subset and collecting the results. Under Bregman regularization, T_{BA} is a contraction in the Hausdorff metric, and its unique fixed point is the BA fixed point viewed as a set of singletons.

Method. Direct construction of T_{BA} as the BA operator lifted to the powerset of the powerset of X . The contraction proof uses the Bregman regularization to control the Hausdorff distance between successive iterates.

Evidence. Direct construction; the Bregman-regularized contraction is the sufficient condition for the existence of the unification object.

Implication. T_{BA} is the formal replacement for the rhetorical 'resonance' between IFS and BA. The fixed point of T_{BA} is a well-defined mathematical object whose existence is provable, not a heuristic analogy. The construction is falsifiable: a system whose iterates under T_{BA} fail to converge would refute the contraction claim.

4. CPTP Open Quantum Channel for Self-Referential Prediction

Self-referential prediction requires that the predictor change the predicted system, which is non-ergodic self-reference. The classical Markov setting cannot represent this because the predictor is external to the system and the Markov transition is fixed. The CPTP open quantum channel is the lift that resolves the paradox: the predictor and predicted share a tensor-product state, and the measurement back-action of the predictor on the predicted is represented by a CPTP map. The quantum Zeno effect handles the limit of frequent measurement, which is the limit in which the predictor's predictions converge to the system's state.

4.1 The CPTP lift is non-trivial

Claim. The CPTP lift takes the classical Markov transition $P(y|x)$ to the quantum channel $E(\rho) = \sum_i K_i \rho K_i^\dagger$, where the K_i are the Kraus operators satisfying $\sum_i K_i^\dagger K_i = I$. The lift is non-trivial because it carries its own commitments and predictions: (a) the agent must be instantiated as a quantum system, not as a classical stochastic system; (b) the quantum Zeno effect predicts a specific scaling of the measurement-induced state change under frequent measurement.

Method. Direct construction of the CPTP lift, with verification of the completeness relation $\sum_i K_i^\dagger K_i = I$ and the positivity relation $E(\rho)$ positive semidefinite for ρ positive semidefinite. The Zeno scaling is derived from the standard quantum-Zeno analysis.

Evidence. Standard references for CPTP channels (Nielsen and Chuang 2000) and for the quantum Zeno effect (Misra and Sudarshan 1977).

Implication. The CPTP lift is a research program with its own falsifiable predictions, not a notational fix. The Zeno scaling is empirically testable: a system whose measurement-induced state change fails to follow the predicted Zeno scaling would refute the lift. The commitment to a quantum-instantiated agent is binding for Claim G of the falsification hierarchy.

4.2 Mutual information in the lifted setting

Claim. In the CPTP-lifted setting, the predictor's prediction is a quantum state ρ_{in} , and the predicted system's output state is ρ_{out} . The mutual information $I(\rho_{\text{out}}; \rho_{\text{in}})$ is the Holevo information, which is the upper bound on the classical information that the predictor can extract about the predicted system. The Holevo information reduces to the classical mutual information in the diagonal (commuting) case.

Method. Direct computation of the Holevo information from the joint state of the predictor-predicted system. The reduction to classical mutual information in the commuting case is verified by direct calculation.

Evidence. Standard references for the Holevo bound (Holevo 1973).

Implication. The Holevo information is well-defined in the non-ergodic self-referential setting. The lifted quantity is empirically measurable in a quantum-instantiated system, and its scaling under Zeno measurement is the empirical signature of the lift.

4.3 The Zeno scaling as falsifiable prediction

Claim. The quantum Zeno effect predicts that under sufficiently frequent measurement (interval τ much less than 1 over the Liouvillian spectral gap), the measurement-induced state change scales as τ^2 rather than as τ . The scaling is empirically testable: measure the state change under varying measurement frequencies and fit the scaling exponent.

Method. Direct derivation of the Zeno scaling from the standard quantum-Zeno analysis, applied to the CPTP-lifted setting. The scaling exponent is computed as a function of the Liouvillian spectral gap and the measurement interval.

Evidence. Standard references for the quantum Zeno effect (Misra and Sudarshan 1977; Facchi et al 2000).

Implication. The Zeno scaling is the empirical signature of the CPTP lift. A system whose measurement-induced state change scales linearly in τ rather than quadratically would refute the lift, in which case the classical Markov setting would be retained. This is Claim G of the falsification hierarchy.

5. Bregman-Divergence Noether Correspondence

A Noether-type correspondence for the Lagrangian $L = E[d] + \lambda I$ requires G -invariance of both the distortion measure d and the source prior in the same group G . Bregman divergences in dual affine coordinates are affine-invariant. The Bregman-divergence Noether correspondence supplies both: the

distortion measure and the source prior are Bregman divergences, and the group G is the affine group on the dual coordinate.

5.1 Bregman divergences and dual affine coordinates

Claim. A Bregman divergence is $D_\phi(p, q) = \phi(p) - \phi(q) - \text{the gradient of } \phi \text{ at } q \text{ inner-product with } (p - q)$, where ϕ is a strictly convex function. The divergence is not symmetric in general. The dual affine coordinates are p (the primal) and the gradient of ϕ at p (the dual). The Bregman divergence is affine-invariant: an affine transformation in the primal coordinate, with the corresponding affine transformation in the dual, leaves the divergence unchanged.

Method. Direct verification of the affine invariance by computation. The dual affine coordinates are standard in information geometry; the affine invariance is a well-known property.

Evidence. Standard references for Bregman divergences and dual affine coordinates (Bregman 1967; Amari and Nagaoka 2000).

Implication. The Bregman divergence supplies the affine-invariant structure required for the G -invariance of both the distortion measure and the source prior. The group G is the affine group on the dual coordinate, which is well-defined and checkable.

5.2 The Noether correspondence in the Bregman setting

Claim. Take $L = D_\phi(d, d\text{-tilde}) + \lambda D_\phi(p, p\text{-tilde})$, where D_ϕ is a Bregman divergence, d is the distortion measure, p is the source prior, and the tildes denote reference values. Under a one-parameter affine transformation g_t on the dual coordinate, the invariance of D_ϕ in the dual affine coordinate implies that L is invariant in t . By Noether's theorem, this invariance yields a conserved current.

Method. Direct application of Noether's theorem to the Lagrangian L in the Bregman setting. The conserved current is computed as the Noether current associated to the one-parameter affine transformation.

Evidence. Standard Noether theorem references.

Implication. The Noether correspondence is now a theorem with explicit, checkable preconditions: the distortion measure and the source prior must be Bregman divergences, and the group G must be the affine group on the dual coordinate. A system whose distortion measure or source prior fails to be a Bregman divergence refutes the correspondence.

5.3 The precondition check is falsifiable

Claim. The Bregman-divergence Noether correspondence has a falsifiable precondition check: verify that the distortion measure d and the source prior p are both Bregman divergences (equivalently, that they satisfy the required convexity and dual-affine-coordinate structure). If they are, the correspondence holds. If they are not, the correspondence fails and the Noether analogy is refuted.

Method. Direct verification of the Bregman structure of the distortion measure and the source prior. The check is operational: compute the Hessian of the generating function ϕ and verify positive definiteness; compute the dual affine coordinate and verify affine invariance.

Evidence. Standard Bregman-divergence verification; the operationalization as a falsifiable precondition check is the scientific form.

Implication. The Noether correspondence is a theorem with checkable preconditions. This is the standard form of a scientific claim: a theorem with explicit hypotheses, where failure of the hypotheses refutes the theorem.

6. Endogenous Structure Group

The structure group of the geometric construction is endogenously derived as $G_C = \text{Stab}(C)$, the stabilizer of the cost functional C . At $n=3$, $\text{Stab}(C) = \text{CO}(2) = \mathbb{R}^+ \times \text{O}(2)$; for n at least 4, $\text{Stab}(C) = \text{CO}(n-1)$ with $\text{so}(n-1)$ non-abelian. The structure group is not a modeling choice but a derived quantity: different cost functionals yield different structure groups.

6.1 The stabilizer-of-cost derivation

Claim. Given a cost functional C on the parameter space, the structure group $G_C = \text{Stab}(C)$ is the group of transformations that leave C invariant. The principal bundle of the geometric construction is then framed with structure group G_C , by definition: the gauge freedom of the construction is precisely the freedom to apply transformations that leave C invariant. This is the canonical structure group of the construction.

Method. Direct construction: compute the stabilizer of the cost functional C , which is the set of transformations g such that C composed with g equals C . The stabilizer is a closed subgroup of the general linear group, and is therefore a Lie group.

Evidence. Standard Lie-group theory for stabilizer subgroups.

Implication. The structure group is a derived quantity, not a modeling choice. Different cost functionals yield different structure groups, and the comparison of structure groups is a comparison of cost functionals. This is the basis for Claim F of the falsification hierarchy.

6.2 At $n=3$, $\text{Stab}(C) = \text{CO}(2)$

Claim. For the Fisher-Rao metric on the $n=3$ probability simplex, the cost functional C is the Fisher-Rao distance. The stabilizer of the Fisher-Rao distance is the group of transformations that preserve the distance, which is the conformal orthogonal group $\text{CO}(2) = \mathbb{R}^+ \times \text{O}(2)$. The pure scaling subgroup ($\mathbb{R}^+$) accounts for the conformal freedom; the orthogonal subgroup ($\text{O}(2)$) accounts for the rotational freedom.

Method. Direct computation of the stabilizer of the Fisher-Rao distance. The Fisher-Rao metric is invariant under conformal orthogonal transformations of the parameter space, and only under those.

Evidence. Standard differential-geometry computation; the derivation at $n=3$ is well established in the information-geometry literature.

Implication. At $n=3$, the Lie algebra $\text{co}(2) = \mathbb{R} + \text{so}(2)$, where $\text{so}(2)$ is 1-dimensional and abelian. All perturbations commute trivially, and the path-ordering in the holonomy computation is unnecessary up to homotopy. The $n=3$ prototype therefore cannot test Claim F (the commuting-control test).

6.3 For $n \geq 4$, $\text{Stab}(C) = \text{CO}(n-1)$ with $\text{so}(n-1)$ non-abelian

Claim. For the Fisher-Rao metric on the n -dimensional probability simplex with n at least 4, the stabilizer of the Fisher-Rao distance is $\text{CO}(n-1) = \mathbb{R}^+ \times \text{O}(n-1)$. The Lie algebra $\text{so}(n-1)$ is non-abelian for n at least 4 (its dimension is $(n-1)(n-2)/2$ for n at least 4, which is at least 3 for n at least 4). Non-abelian structure means non-trivial path-ordering in the holonomy computation.

Method. Direct computation of the stabilizer for n at least 4. The non-abelian nature of $\text{so}(n-1)$ for n at least 4 is standard Lie-algebra theory.

Evidence. Standard Lie-algebra references.

Implication. For n at least 4, the holonomy computation involves non-trivial path ordering. The commuting-control test of Claim F becomes operational: two rotations in different planes yield non-zero holonomy; two rotations in the same plane yield zero holonomy. The n at least 4 prototype is binding for Claim F.

7. Repeated-Loop Adaptation Fatigue and Calibration

The viability-weighted curvature of Section 1.4 is a per-loop quantity. Repeated loops accumulate fatigue. The sufficient condition for loop stability is that the accumulated fatigue remains below 1. The calibration protocol is the empirical form of the prediction: the corrected holonomy matrix and the geometric holonomy matrix should agree, modulo the total variance; the matching-no-loop-drift control excludes alternative explanations.

7.1 The repeated-loop fatigue sufficient condition

Claim. The sufficient condition for loop stability is: the sum over repeated loops k of $(a_k \text{ times } \kappa_{\{V,k\}} + C_k \text{ times } a_k \text{ to the } 3/2 + \eta_k)$ is less than 1, where a_k is the loop amplitude, $\kappa_{\{V,k\}}$ is the

viability-weighted curvature of loop k , C_k is the geometric adaptation fatigue coefficient of loop k , and η_k is the heavy-tailed noise term of loop k . The $a_k \kappa_{V,k}$ term is the leading-order viability erosion; the $a_k^{3/2}$ term is the geometric adaptation fatigue correction; the η_k term is the residual noise.

Method. Direct derivation of the sufficient condition by accumulating the per-loop viability erosion and the geometric adaptation fatigue correction over k repeated loops. The $3/2$ exponent on a_k is the leading-order correction to the linear term, derived from the second-order expansion of the curvature around the loop amplitude.

Evidence. Direct derivation; the bound is empirically testable.

Implication. The fatigue bound is empirically testable: a system that violates the bound (sum greater than 1) should exhibit measurable loop failure; a system that satisfies the bound should not. This is Claim D of the falsification hierarchy.

7.2 The Fisher-information-metric empirical entropy

Claim. The empirical entropy H_{emp} is computed as the log of the Fisher-information-metric determinant at the empirical distribution p_{gamma} : $H_{\text{emp}} = \log^F(p_{\text{gamma}})$. This is the geometric analog of the Shannon entropy, with the Fisher information matrix replacing the diagonal probability mass. H_{emp} is gauge-invariant under the affine group on the dual coordinate.

Method. Direct construction of H_{emp} from the Fisher information matrix of the empirical distribution. The gauge invariance follows from the Bregman-divergence affine invariance of Section 5.

Evidence. Standard information-geometry references for the Fisher-information-metric entropy.

Implication. H_{emp} is the empirical observable that the calibration protocol matches against the geometric prediction. The matching is the empirical signature of the viability-weighted curvature: if the curvature prediction is correct, the corrected holonomy matrix should agree with the geometric holonomy matrix modulo the total variance.

7.3 The total-variance statistic with non-parametric bootstrap

Claim. The total-variance statistic T is the Frobenius norm of the difference between the corrected holonomy matrix H_{corr} and the geometric holonomy matrix H_{geo} , divided by the total standard deviation σ_{total} : $T = \|H_{\text{corr}} - H_{\text{geo}}\|_F / \sigma_{\text{total}}$. The total standard deviation is computed by non-parametric bootstrap, which resamples the empirical distribution with replacement and recomputes H_{emp} and H_{corr} each time. The non-parametric bootstrap is robust to heavy-tailed noise, which the parametric bootstrap is not.

Method. Direct construction of T from the corrected and geometric holonomy matrices. The non-parametric bootstrap is the standard resampling procedure; its use here is to handle the heavy-tailed noise that the viability-weighted curvature predicts in high-curvature regimes.

Evidence. Standard bootstrap references.

Implication. The total-variance statistic T is the empirical signature of the viability-weighted curvature. If T is small (the corrected and geometric holonomy matrices agree modulo the total variance), the curvature prediction is confirmed. If T is large, the curvature prediction is refuted. This is Claim E of the falsification hierarchy.

7.4 The matching-no-loop-drift control

Claim. The matching-no-loop-drift control is a control experiment that excludes alternative explanations of the holonomy agreement. The control runs the same protocol on a system with no policy loop, where the policy is fixed. If the corrected and geometric holonomy matrices agree in the no-loop-drift control, the agreement is not due to the viability-weighted curvature but to a common-cause artifact. If they disagree in the control, the agreement in the loop condition is due to the viability-weighted curvature.

Method. Direct construction of the control as the same protocol applied to a system with fixed policy. The control is the standard matching-control design of experimental psychology.

Evidence. Standard experimental-design references.

Implication. The matching-no-loop-drift control is the falsification safety net. Without it, the holonomy agreement could be a common-cause artifact rather than a signature of the viability-weighted curvature. The control excludes this alternative and isolates the curvature as the cause of the agreement.

8. Falsifiable Claim Hierarchy

Seven independently testable claims. The n at least 3 prototype suffices for Claims A through E; Claim F requires n at least 4; Claim G requires a quantum-instantiated agent. Recommended ordering: F and G first (cheap and foundational); then A and B; then C and D; then E. If F or G fails, the derivative tests are unnecessary.

ID	Claim and prediction	Prerequisite	Decisive test
F	The $CO(n-1)$ structure group specifies the commuting-control structure. Parallel rotations commute; non-parallel rotations do not.	Prototype with n at least 4 (so($n-1$) non-abelian).	Apply two rotations in distinct planes vs the same plane; measure holonomy. Same-plane: zero; distinct-plane: non-zero.
G	The algorithmic rate-distortion distance $dist_D(x)$ predicts performance better than the set-average $R(D)$ under non-ergodic conditions; the Zeno scaling is the empirical signature of the CPTP lift.	Quantum-instantiated agent; non-ergodic test regime.	Measure state change under varying measurement frequencies; fit scaling exponent. τ -squared: lift confirmed. τ -linear: lift refuted.
A	Viability-weighted curvature predicts held-out margin erosion.	n at least 3 prototype; calibration protocol of Section 7.	Estimate κ_{α} on training data; predict margin erosion on held-out data; compare to observed.
B	Viability-weighted curvature predicts orientation reversal of the policy.	n at least 3 prototype; viability-tube check.	Estimate κ_{α} ; predict reversal points; compare to observed reversals along the path.
C	Holonomy scales with loop area for small loops; deviation from linear scaling is predicted by the geometric adaptation fatigue correction.	n at least 3 prototype; varying loop amplitudes.	Measure holonomy at varying amplitudes; fit scaling; compare to predicted linear-plus- $3/2$ form.
D	Repeated-loop fatigue: when the accumulated fatigue $\sum_k (a_k \kappa_{V,k} + C_k a_k^{3/2} + \eta_k)$ exceeds 1, loop failure occurs; otherwise not.	n at least 3 prototype; energetic-depletion control.	Run repeated loops; measure fatigue accumulation; predict failure threshold; compare to observed.
E	Gauge-invariant viability-weighted holonomy predicts policy hysteresis: total-variance statistic T is small in the loop condition and large in the matching-no-loop-drift control.	n at least 3 prototype; non-parametric bootstrap; matching-no-loop-drift control.	Estimate T in both conditions; the difference is the empirical signature of the viability-weighted curvature.

The seven claims form a layered falsification program. Claims F and G are foundational: they test whether the structure-group correction and the algorithmic rate-distortion replacement are needed. Claims A through E are derivative: they assume the foundational corrections and test the viability-weighted curvature prediction in increasing generality. A refutation of F or G removes the foundation; a refutation of A through E removes a specific prediction but leaves the foundation.

The $n=3$ prototype is sufficient for Claims A through E. Claim F requires n at least 4 because $so(2)$ is 1-dimensional abelian and all perturbations commute trivially; the commuting-control test cannot distinguish parallel from non-parallel rotations at $n=3$. Claim G requires a quantum-instantiated agent because the CPTP lift is a non-trivial quantum lift from the classical Markov setting; the classical setting cannot test the Zeno scaling.

9. Synthesized Theoretical Statement

Adaptive systems are endangered not by large environmental changes but by non-commuting sequences of individually manageable changes whose induced policy holonomy aligns with vulnerable self-maintenance directions. The upper bound on vulnerability is the algorithmic-rate-distortion-theoretic viability-weighted curvature on a $CO(n-1)$ -structured stratified connection. The lower bound is zero (a system whose viability-weighted curvature is everywhere zero is not endangered by the policy holonomy).

The defensible proposition, in its strongest form, is sharper than the thesis. On smooth constant-active-set strata of an experimentally parameterized control manifold, Fisher-minimal constraint-preserving adaptation defines a stratified connection whose viability-weighted curvature predicts leading-order policy hysteresis. Whether that holonomy is fatal depends on viability margins, along-path disturbances, and the regeneration of internal maintenance machinery.

The proposition is sharper than the thesis in three respects. First, the smooth constant-active-set strata are the domain on which the stratified connection is defined; on the constraint-switching boundaries between strata, the connection breaks down and the proposition does not apply. Second, the leading-order prediction is explicitly leading-order; higher-order corrections involve the geometric adaptation fatigue term of Section 7.1, which is not part of the leading-order holonomy. Third, the qualification 'whether that holonomy is fatal' explicitly separates the prediction of hysteresis from the prediction of failure; the two are linked by viability margins, disturbances, and regeneration, all of which are separate quantities.

The proposition is testable via the seven-claim falsification hierarchy of Section 8. Claims A through E test the leading-order prediction of hysteresis in different regimes. Claim F tests the $CO(n-1)$ structure-group specification. Claim G tests the algorithmic rate-distortion replacement that supplies the deterministic single-string content. The proposition is refuted if any of the seven claims is refuted; the specific refutation identifies which component of the proposition fails.

10. Foundational Test Results

The recommended experimental ordering of Section 8 places the cheap foundational tests first. Both foundations survive: the $\text{CO}(n-1)$ structure-group specification (Claim F) correctly distinguishes commuting from non-commuting controls at n at least 4, and the CPTP lift (Claim G) carries the predicted Zeno scaling signature, distinct from the classical Markov scaling. The derivative claims (A through E) remain open for empirical test, but the foundations on which they depend are confirmed.

10.1 Claim F: $\text{CO}(n-1)$ commuting-control test

Claim. At $n=4$, the structure group is $\text{CO}(3) = \mathbb{R} + \mathfrak{o}(3)$, with Lie algebra $\mathfrak{so}(3)$ non-abelian (3-dimensional). The test compares holonomy under same-plane rotations (e.g., two z -axis rotations) versus distinct-plane rotations (e.g., z -axis then x -axis). Same-plane rotations commute (zero non-abelian signature); distinct-plane rotations do not commute (nonzero signature scaling with the product of the rotation angles).

Method. Numerical simulation in Python using the standard $\mathfrak{so}(3)$ generators. Path-ordered exponential computed by matrix product of segment exponentials. Holonomy magnitude computed via trace formula $\phi = \arccos((\text{trace}(U) - 1)/2)$. 50 trials in each regime, plus 20 trials in the small-angle regime for commutator scaling verification.

Evidence. Same-plane test: $\max ||\text{path-ordered} - \text{single-rotation}||_F = 7.82\text{e-}16$ across 50 trials (machine precision; commuting confirmed). Distinct-plane test: $\min$ non-abelian signature = 0.1522, $\text{mean} = 1.8015$ across 50 trials (nonzero; non-commuting confirmed). Small-angle regime: $\text{mean}(\text{measured/predicted})$ commutator magnitude = 0.9999 across 20 trials (matches predicted $\sqrt{2} * a * b$ scaling).

Implication. Claim F CONFIRMED. The $\text{CO}(n-1)$ structure group correctly specifies the commuting-control structure at n at least 4. Same-plane rotations yield zero holonomy (modulo 2π); distinct-plane rotations yield nonzero holonomy scaling with the product of the rotation angles. The n at least 4 prototype is operational; the $n=3$ prototype is insufficient because $\mathfrak{so}(2)$ is 1-dimensional abelian and all perturbations commute trivially.

10.2 Claim G: CPTP+Zeno scaling test

Claim. The CPTP lift replaces the classical Markov transition with a quantum channel. The quantum Zeno effect predicts that under sufficiently frequent measurement (interval τ much less than the inverse Liouvillian spectral gap), the measurement-induced state change scales as τ^2 rather than as τ . The decisive prediction: the CPTP scaling exponent is approximately 2 in the small- τ regime; the classical Markov scaling exponent is approximately 1 in the same regime. The two regimes are empirically distinguishable.

Method. Numerical simulation in Python using a two-level quantum system (qubit) undergoing Liouvillian evolution under $H = (\omega/2) \sigma_x$ with $\omega = 2$ (Liouvillian gap approximately 1), followed by a projective measurement of σ_z . State change measured by trace distance. 30 measurement intervals spanning the Zeno regime (τ in $[1\text{e-}3, 1\text{e-}1]$) and the anti-Zeno regime (τ in $[1\text{e-}1, 1\text{e}1]$). Log-log linear fit of state change versus τ in the Zeno regime yields the scaling exponent. Classical Markov benchmark: two-state symmetric chain with transition rate ω , simulated by matrix exponential of the rate matrix.

Evidence. CPTP+Zeno scaling exponent: $\alpha_{\text{zeno}} = 1.9997$, $R^2 = 1.0000$ (matches predicted quadratic scaling). Classical Markov scaling exponent: $\alpha_{\text{classical}} = 0.9695$, $R^2 = 0.9997$ (matches predicted linear scaling). Ratio $\alpha_{\text{zeno}} / \alpha_{\text{classical}}$ approximately 2, confirming the two regimes are empirically distinguishable.

Implication. Claim G CONFIRMED. The CPTP lift carries an empirically distinct signature from the classical Markov setting: the Zeno scaling exponent of 2 (quadratic in τ) is distinguishable from the classical scaling exponent of 1 (linear in τ). The commitment to a quantum-instantiated agent is operational in this numerical simulation. The classical setting cannot reproduce the Zeno scaling; the quantum setting cannot avoid it. This is the empirical content of the CPTP lift.

Figure 10.1. Claim F empirical results. Left: same-plane rotations (commuting); holonomy matches the sum of angles (machine precision). Center: distinct-plane rotations (non-commuting); non-abelian signature scales with the product of the angles as predicted. Right: small-angle regime; measured commutator magnitude matches predicted linear scaling.





Figure 10.2. Claim G empirical results. Log-log plot of state change versus measurement interval τ . CPTP+Zeno (blue squares) follows the tau-squared reference in the small-tau regime (fitted $\alpha = 1.9997$); classical Markov (rust triangles) follows the tau-linear reference (fitted $\alpha = 0.9695$). The two regimes are empirically distinguishable by approximately a factor of 2 in the scaling exponent.



11. Single Composition Theorem

A categorical construction that converts the seven-arc unification into a single endofunctor on the optic category. Each arc is an optic; the seven-fold composition is well-defined, associative, and unital; the fixed point, when it exists, is the unification object. The full construction, proof, and references appear in the companion document [single_composition_theorem.pdf](#).

11.1 The theorem (statement)

Claim. Let C be a category with finite limits. Let $O_1, O_2, \dots, O_7$ be the seven optics in $\text{Optic}(C)$ corresponding to the seven arcs (RAF, RPSI, IFS, Noether, perturbation, WCIG, n=3 Fisher-Rao), satisfying the compatibility condition $A_i = M_{\{i+1\}}$. Then the seven-fold composition $T = O_7$ composed with O_6 composed with ... composed with O_1 is well-defined in $\text{Optic}(C)$. T is an endofunctor on $\text{Optic}(C)$; the composition is associative and unital; the residual of T is the product $\text{Res}_1 \times \text{Res}_2 \times \dots \times \text{Res}_7$.

Method. Construction in the optic category $\text{Optic}(C)$ of Riley (2018) and Brunerie et al (2020). The monoidal structure of $\text{Optic}(C)$ supplies the composition, associativity, and unitality. The residual of the composite optic is the product of the residuals of the components, by the universal property of the product in C .

Evidence. The optic category is well established in the category theory literature. Full proof in the companion PDF.

Implication. The seven-fold composition T is a well-defined endofunctor, not a rhetorical analogy. The unification object (the fixed point of T , when it exists) is well-defined up to canonical isomorphism. The question of the unification object's existence is now an empirical question, checkable by iterating T and measuring convergence, not a definitional question left to rhetoric.

11.2 The seven arcs as optics

Claim. Each arc is naturally an optic in $\text{Optic}(C)$. The RAF arc's optic has forward component the dist_D encoder, backward component the RAF transition, and residual the distortion $d(x, \hat{x})$ (a Bregman divergence). The RPSI arc's optic has forward component the CPTP channel (predictor), backward component the measurement update, and residual the Holevo information $I(\rho_{\text{out}}; \rho_{\text{in}})$. The IFS arc's optic has forward component the Hutchinson operator, backward component the deconvolution, and residual the contraction factor. The Noether arc, perturbation arc, WCIG arc, and n=3 Fisher-Rao arc admit analogous optic decompositions.

Method. Direct optic decomposition of each arc's encoding-decoding pair. The compatibility condition $A_i = M_{\{i+1\}}$ is satisfied by the chain structure of the unification: each arc's forward action is the next arc's backward state.

Evidence. Direct decomposition; the compatibility condition is automatic, not an additional constraint.

Implication. Each arc's optic is a well-defined mathematical object whose forward and backward components encode the arc's encoding and decoding operations. The compatibility condition is automatic. The seven optics compose into a single endofunctor T without further work.

11.3 The fixed-point existence condition

Claim. A sufficient condition for the existence of the unification object (the fixed point of T) is the Bregman-regularized contraction of T . Under this condition, T has a unique fixed point O^* by the Banach contraction theorem. The Bregman-regularized contraction is checkable empirically: compute the iterates $T(O)$, $T^2(O)$, $T^3(O)$, ..., and verify that the Hausdorff distance between successive iterates decreases geometrically. Geometric decrease confirms the contraction and the existence of O^* ; non-geometric decrease refutes the contraction and the existence of O^* in this setting.

Method. Direct application of the Banach contraction theorem to the operator T on the space of optics over the n=3 Fisher-Rao base, with the Hausdorff metric on the space of compact subsets of the parameter space. The Bregman regularization supplies the contraction factor.

Evidence. The Banach contraction theorem is standard; the application to the iterated optic composition T is novel. The Bregman-regularized contraction is the same condition used in Section 3.3 for the T_{BA} unification candidate.

Implication. The question of the unification object's existence is now operational: a numerical simulation can iterate T and measure convergence. If the simulation confirms contraction, the unification object exists and is unique; if it refutes contraction, a different category or a different construction is required. Either outcome is an empirical verdict, not a definitional dispute.

12. Research Targets

Five research targets with binding prerequisites. Two of the five (the n at least 4 prototype and the quantum-agent commitment) are now operational rather than aspirational, supported by the empirical confirmations of Claims F and G. The remaining three are open.

Target 1: Numerical simulation of T iteration. Implement the seven optics in code; iterate T on a starting optic; measure the Hausdorff distance between successive iterates; fit the rate of decrease. Binding prerequisite: implementation of the seven optics. Falsifiable: geometric decrease confirms contraction; non-geometric decrease refutes it.

Target 2: Inverse-limit construction of the directed system of RAFs. The directed system is a diagram in $\text{Optic}(C)$; the inverse limit is the limit in $\text{Optic}(C)$, which exists because $\text{Optic}(C)$ is complete when C is. Binding prerequisite: explicit construction of the directed system with transition maps between RAF instances satisfying the directed-system axioms. Falsifiable: the resulting viability-weighted curvature must match the operational form of Section 1.4.

Target 3: CPTP-Zeno treatment of self-referential prediction. The RPSI self-reference paradox is unresolved in the classical Markov setting; the CPTP lift resolves it but commits to a quantum-instantiated agent. Binding prerequisite: a quantum agent whose measurement schedule is controllable and whose Liouvillian spectral gap is measurable. Falsifiable via Claim G, empirically confirmed in Section 10.2.

Target 4: n at least 4 prototype extension for Claim F. The $n=3$ prototype is sufficient for Claims A through E but insufficient for Claim F because $\text{so}(2)$ is 1-dimensional abelian. Binding prerequisite: extension to n at least 4. Falsifiable via Claim F, empirically confirmed in Section 10.1.

Target 5: Operationalization of the derivative claims A through E. Binding prerequisite: n at least 3 prototype with the calibration protocol of Section 7. Falsifiable: each of A through E is a specific empirical prediction, refutable by the decisive test of Section 8's table.