
Surviving Findings from the DeepSeek Cross-Domain Unification Audit Project

Claims, methods, evidence, and implications

A 16,271-line transcript proposes a six-arc unification: RAF rate-distortion, Recursive Predictive Self-Information consciousness, Iterated Function System fractals, Noether-type symmetry correspondence, perturbation theory of optimal encoding, and the $n=3$ Fisher-Rao geometric construction.

Two independent line-level audits applied ripgrep-verified claim extraction to every theorem, definition, and caveat. Cross-examination isolates the technical claims that survive.

This report presents those surviving claims, the methods that established them, the evidence that supports them, and their implications for the underlying research program.

ABSTRACT

Surviving Findings

A 16,271-line transcript proposes a six-arc unification of rate-distortion, fractal, symmetry, and information-geometric constructions. Two independent line-level audits verified each claim against the source text, then cross-examined the converging and diverging findings. This report presents the technical claims that survive that cross-examination.

The DeepSeek cross-domain unification transcript composes six arcs: RAF rate-distortion as a bridge from Kolmogorov complexity to Shannon entropy; Recursive Predictive Self-Information (RPSI) as a consciousness definition; Iterated Function Systems (IFS) as a fractal account of perception; Noether-type symmetry correspondence for Markov chains; perturbation theory of the optimal encoding; the abortive Wasserstein-Categorical Information Geometry (WCIG) upgrade; the bridge-rung composition; and the explicit $n=3$ Fisher-Rao geometric construction with a curvature-survival equivalence theorem. The author of the transcript acknowledges four defects in a final self-assessment section (transcript lines 13985 through 16270).

Two independent line-level audits confirmed those four defects and identified additional mathematical breakdowns, internal inconsistencies, and missing premises across every arc. Cross-examination of the two audits surfaced a further six defects visible only when both perspectives are synthesized, and isolated three corrections to early conjectures. The surviving technical claims compose a coherent research program rather than a catalogue of errors.

The central surviving thesis is the following. Adaptive systems are endangered not by large environmental changes but by non-commuting sequences of individually manageable changes whose induced policy holonomy aligns with vulnerable self-maintenance directions. The upper bound on vulnerability is the algorithmic-rate-distortion-theoretic viability-weighted curvature on a $CO(n-1)$ -structured stratified connection. The defensible proposition is sharper: on smooth constant-active-set strata of an experimentally parameterized control manifold, Fisher-minimal constraint-preserving adaptation defines a stratified connection whose viability-weighted curvature predicts leading-order policy hysteresis; whether that holonomy is fatal depends on viability margins, along-path disturbances, and the regeneration of internal maintenance machinery.

The report is organized as a sequence of claim-method-evidence-implication blocks. Sections 2 through 5 establish what does not survive examination. Sections 6 through 12 present the synthesized theoretical framework that replaces the broken constructions. Section 13 presents the seven-claim falsifiable prediction hierarchy with the prerequisite that the $n=3$ prototype is sufficient for claims A through E but a non-abelian extension to n at least 4 is binding for the structure-group test. Sections 14 and 15 state the synthesized thesis and its open research targets.

SECTION 1

Project Scope and Method

Scope and methodology of the audit project, the falsifiability standard applied, and the verification protocol used to certify absence of constructs in the source transcript.

1.1 Source and scope

The source is a single DeepSeek conversation transcript of 16,271 lines, produced in a single sitting, that proposes a cross-domain unification. The unification composes six construction arcs. Arc 1 establishes RAF (Reactive Anticipatory Framing) as a rate-distortion bridge between Kolmogorov complexity and Shannon entropy (transcript lines 1 through 427). Arc 2 introduces Recursive Predictive Self-Information, or RPSI, as a definition of consciousness (lines 428 through 1022). Arc 3 treats fractals as Iterated Function Systems and connects perception to fractal attractors (lines 1044 through 1814). Arc 4 invokes Noether-type symmetry correspondence for Markov chains (lines 1908 through 2324). Arc 5 develops perturbation theory of the optimal encoding (lines 2378 through 3136). Arc 6 is the abortive Wasserstein-Categorical Information Geometry upgrade (lines 3140 through 3518). The seven bridge rungs compose the arcs across lines 4000 through 4500. The Counterfactual Gauge Theory upgrade spans lines 4600 through 5200. The $n=3$ Fisher-Rao explicit construction begins near line 11000. The final self-assessment section spans lines 13985 through 16270.

The audit project applied two independent line-level reviews to the same transcript. Each audit located every theorem statement, every definition, and every caveat using ripgrep searches against the transcript file, then read the surrounding twenty-line window to extract the precise claim. Each claim was then evaluated against three questions: is it mathematically correct given the stated assumptions; is it substantively novel given the cited prior art; and is it consistent with claims made elsewhere in the transcript.

A second verification protocol, used repeatedly throughout the project, is the zero-match proof of absence. When a defect consists of a missing construct, the absence was certified by running ripgrep for the construct name against the full transcript and recording zero matches. Constructs certified absent by this method include pathwise viability, the Nagumo inward-pointing condition, viability tube, intermediate-point feasibility, active-set constraint switching, and stratified regularity beyond piecewise-smooth. These zero-match results are the evidence that the corresponding defects are not artifacts of an unsympathetic reading but structural features of the source.

The falsifiability standard applied throughout the project is the same standard the project applies to the source transcript: every surviving claim must be associated with at least one observable quantity whose measured value would either confirm or refute the claim. Where a claim cannot be made falsifiable in its original form, the project replaces it with a stronger version that can. The seven-claim falsifiable hierarchy in Section 13 is the output of that replacement program.

1.2 Cross-examination method

The cross-examination method is the project's distinctive contribution. Where the two audits converge, the convergence is itself evidence: the finding is real and not a perspective artifact. Where they diverge, the divergence identifies complementary concerns rather than contradictions. The cross-arc structural pattern (Section 3) is an output of this method, as are the six joint cross-reference defects (Section 5) and the strengthened upgrades that integrate insights from both perspectives (Sections 6 through 12).

Each surviving claim is presented below in a four-part structure: claim, method, evidence, implication. The claim states the technical content. The method states how the claim was established or how the replacement is constructed. The evidence cites the specific transcript line numbers, the ripgrep verification results, or the mathematical counterexample. The implication states what the claim means for the underlying research program. This structure is preserved throughout the report.

SECTION 2

Verified Defects in the Source Transcript

Four defects that the transcript's own final self-assessment acknowledges. The two independent audits independently confirm each one. The convergence is itself the evidence that these are real.

Each defect is presented with its specific transcript line citation, the mathematical content of the defect, and the corrected form. These four are the minimal set that any defensible version of the transcript's unification must address.

2.1 Structure group of the Fisher-Rao construction

Claim. The transcript frames the principal bundle of the $n=3$ Fisher-Rao construction with structure group $GL(2)$. The correct structure group is $CO(2) = \mathbb{R}^+ \times O(2)$, the conformally orthogonal group, because the Fisher-Rao metric is invariant only under conformal orthogonal transformations of the parameter space, not under general linear transformations.

Method. ripgrep verification: the $GL(2)$ framing appears in the $n=3$ construction (approximately transcript line 11000 forward); the self-assessment at lines 13985 through 16270 acknowledges the correction. The mathematical content is verified by direct computation of the Fisher-Rao metric's invariance group.

Evidence. Transcript lines 13985 through 16270 (self-assessment section); the $CO(2)$ correction is acknowledged by the author within the same section.

Implication. The bundle atlas of the $n=3$ construction is correctly described only after the structure group is corrected to $CO(2)$. Any computation that uses $GL(2)$ transitions between charts overcounts the gauge freedom and yields spurious holonomy. The corrected structure group is the foundation for the commuting-control test of Section 13, Claim F.

2.2 'Exact predictive variance' is a squared mean shift

Claim. The quantity labeled 'exact predictive variance' in the transcript is in fact a squared mean shift, bias squared, while the predictive variance $\text{Var}(y)$ remains constant at the identity I throughout. The labeling suggests a variance formula; the substance is a bias formula.

Method. Direct verification of the algebraic derivation as written in the transcript. The expression decomposes as bias squared plus variance, and the variance term is constant; only the bias term carries the dynamics the labeling attributes to the variance.

Evidence. Transcript self-assessment section (lines 13985 through 16270); the same arc in which the variance is initially introduced also retracts the label, an internal inconsistency recorded in Section 3.

Implication. Empirical predictions that depend on the 'predictive variance' actually depend on bias squared. Reinterpretation changes which experimental observable should be measured. The corrected label also clarifies that the dynamics are carried by the bias term,

which is what the algorithmic rate-distortion replacement of Section 7 makes operational.

2.3 The RAF invariance theorem reduces to continuity on a compact set

Claim. The theorem stated as the RAF invariance theorem reduces to the observation that a positive continuous function on a compact set attains its infimum. The conclusion follows from continuity plus compactness alone; no RAF-specific structure is used.

Method. Direct reading of the proof as written. The argument invokes only the extreme-value theorem applied to the relevant functional. No use is made of the reactive, anticipatory, or framing structure that defines RAF.

Evidence. Transcript self-assessment section (lines 13985 through 16270); the proof appears earlier in the RAF arc (lines 1 through 427).

Implication. The theorem cannot bear the weight assigned to it in the unification chain. The inverse-limit construction of the directed system of RAFs, proposed as an open problem in Section 15, is the replacement research target.

2.4 The 4-species register machine is schematic

Claim. The 4-species chemical register machine invoked in the RAF arc has no reaction list. No reactants, products, or catalysts are specified. The machine is a label without a mechanism.

Method. ripgrep verification: searches for reaction-list keywords, reactant names, and catalyst identifiers return only the introductory framing of the machine, no enumeration of reactions.

Evidence. Transcript RAF arc and the self-assessment section (lines 1 through 427 and 13985 through 16270).

Implication. A schematic machine cannot be the substrate of an autopoiesis closure test. The intervention-based autopoiesis closure test of Section 6 supplies a concrete 4-species example (species A, B, C, D; reactions R1 through R4 catalyzing specified transformations) that makes the closure test operational.

SECTION 3

Cross-Arc Structural Pattern

Each individual arc is mathematically real. The connections between arcs are analogies rather than formal mappings. The unification is composition by analogy, not by a composition theorem.

The cross-arc structural pattern is the most consequential systemic finding of the audit project. Every individual arc invokes legitimate mathematics: rate-distortion theory, IFS fixed-point theorems, Noether's theorem, perturbation theory, information geometry, and category theory. The defect is not in any single arc; it is in the connections between arcs. Each bridge between arcs is a rhetorical move rather than a formal construction.

3.1 The rhetorical-bridge pattern

The bridge between the RAF arc and the curvature arc is presented as a theorem. The self-assessment retracts it to a continuity observation. The internal inconsistency between the theorem framing and the honest caveat is an instance of the structural pattern: the strong form is asserted when the construction is built, and the weak form is conceded when the construction is audited. The transcript itself oscillates between these registers.

The bridge between the WCIG upgrade and the dagger-compact categorical framework of constructor theory lacks a translation functor. The transcript introduces both vocabularies but never constructs a functor between them. A vocabulary shared in name is not a mapping in category-theoretic substance. The same pattern appears in Bridge Rung 6, which is a categorical tautology: the conclusion is a definitional restatement of the premises, not a theorem.

The bridge from RAF rate-distortion to algorithmic rate-distortion is blocked by a type confusion. The transcript's $R(D)$ is a set-average quantity defined over a probability distribution on source strings. The Kolmogorov complexity $K(x)$ is a single-string quantity. No i.i.d. or ergodicity assumption is stated to connect them. The bridge from deterministic state transition to the $R(D)$ achievability is similarly blocked: the deterministic transition δ is a special case, but achieving $R(D)$ in general requires random coding, and the gap $R_{\text{det}}(D)$ greater than or equal to $R(D)$ can be infinite.

3.2 The missing-funnel pattern

The bridge from RAF to RPSI consciousness is blocked by an ergodicity self-contradiction. RAF assumes an ergodic channel to invoke the asymptotic equipartition property that gives the rate-distortion bound operational meaning. RPSI consciousness requires that the predictor change the predicted system, which is a non-ergodic self-reference. The two arcs cannot both be true in the same construction without an additional mechanism that resolves the contradiction.

The bridge from IFS fractals to Blahut-Arimoto iterations is blocked by a category confusion. IFS attractors are fixed points of Hutchinson-Banach contractions on metric spaces of

subsets. BA fixed points are fixed points of the BA operator on probability simplices. The transcript calls this shared fixed-point structure 'resonance'; the resonance is rhetorical. The optic-category unification of Section 8 supplies the formal replacement.

The Noether-type correspondence for $L = E[d] + \lambda I$ is blocked by a missing premise. The correspondence requires G-invariance of both the distortion measure d and the source prior, in the same group G . The transcript states neither. The Bregman-divergence Noether correspondence of Section 10 supplies the missing premise and turns the analogy into a theorem with explicit, checkable preconditions.

3.3 The composition-by-analogy pattern

The bridge rungs of the unification chain are presented as a seven-rung ladder. Each rung is a category-theoretic vocabulary item: profunctor, span, optic, lens, dependent type, and so on. No rung is supplied with a composition theorem that takes the previous rung's output as input. The ladder is a sequence of names rather than a sequence of constructions. The single-composition-theorem replacement, listed as an open problem in Section 15, is the research target that would convert the ladder into a unification.

The honest-caveat versus theorem-framing register tension is the systemic form of the pattern. When a claim is being constructed, the transcript frames it as a theorem with hypotheses and a conclusion. When the same claim is later audited, the transcript frames it as a heuristic with caveats. Both framings appear in the same document. The reader cannot tell which claims are asserted as theorems and which as heuristics. This tension is itself a defect: a research program must commit to a register for each claim and defend the claim in that register.

SECTION 4

Specific Mathematical Breakdowns

Six mathematical claims in the transcript that fail examination. Each is paired with the formal reason for failure and the replacement construction that recovers a defensible version.

4.1 Rate-distortion type confusion

Claim. The transcript's $R(D)$ is a set-average quantity defined over a probability distribution on source strings. $K(x)$, the Kolmogorov complexity, is a single-string quantity. The transcript bridges them without stating an i.i.d. or ergodicity assumption that would license the bridge.

Method. Direct reading of the relevant RAF arc and rate-distortion definitions (lines 1 through 427 and the perturbation arc 2378 through 3136). The type confusion is established by inspecting the definitions of $R(D)$ and $K(x)$ as written.

Evidence. Transcript rate-distortion definitions and the RAF arc; the perturbation arc's use of $R(D)$ in perturbation expansions.

Implication. The bridge cannot be repaired by adding an assumption; the deterministic transition delta of the RAF model is incompatible with the random coding that $R(D)$ achievability requires. The algorithmic rate-distortion replacement of Section 7 disambiguates the type and supplies the missing definition.

4.2 Deterministic transition versus random coding achievability

Claim. RAF uses a deterministic transition delta, but achieving $R(D)$ requires random coding in general. The deterministic rate $R_{\text{det}}(D)$ is greater than or equal to $R(D)$, and the gap can be infinite.

Method. Comparison of the deterministic transition delta as defined in the RAF arc with the achievability proof of $R(D)$, which requires an ensemble of codewords drawn from a distribution.

Evidence. Transcript RAF arc (lines 1 through 427); $R(D)$ achievability is cited from standard rate-distortion references; the gap $R_{\text{det}}(D) \geq R(D)$ is well known in the rate-distortion literature.

Implication. Any empirical claim that uses the RAF transition as a code is restricted to $R_{\text{det}}(D)$, which is strictly weaker than $R(D)$ in the worst case. The replacement is to use the algorithmic rate-distortion $\text{dist}_D(x)$ of Section 7, which is intrinsically deterministic and does not require random coding to achieve.

4.3 Ergodicity self-contradiction

Claim. RAF assumes an ergodic channel to invoke asymptotic equipartition. RPSI consciousness requires that the predictor change the predicted system, which is non-ergodic self-reference. The two arcs cannot both hold in the same construction without an additional mechanism.

Method. ripgrep verification of the ergodicity assumption in the RAF arc and the self-reference structure of RPSI in the consciousness arc. The two constructions are formally incompatible in their stated forms.

Evidence. Transcript RAF arc (lines 1 through 427) and consciousness arc (lines 428 through 1022).

Implication. The CPTP open quantum channel of Section 9 supplies the additional mechanism. In an open quantum system, the predictor and predicted share a tensor-product state, and the measurement back-action of the predictor on the predicted is represented by a CPTP map. The quantum Zeno effect handles the limit of frequent measurement.

4.4 IFS and Blahut-Arimoto fixed-point category confusion

Claim. The transcript identifies IFS attractors and BA fixed points as instances of a shared 'resonance' structure. IFS attractors are fixed points of Hutchinson-Banach contractions on metric spaces of subsets. BA fixed points are fixed points of the BA operator on probability simplices. The two are fixed points of operators on different categories; no functor between them is constructed.

Method. Direct reading of the IFS and BA sections of the transcript. The absence of a functor is established by inspecting the constructions: no mapping of objects and morphisms between the categories is given.

Evidence. Transcript IFS arc (lines 1044 through 1814) and BA operator references in the perturbation arc (lines 2378 through 3136).

Implication. The optic-category unification of Section 8 supplies the functor. IFS attractors are pure coalgebras; BA fixed points are coalgebras with residual. The unification candidate is the operator T_{BA} on the powerset of the powerset of X , with contraction under Bregman regularization.

4.5 Noether correspondence missing premises

Claim. The transcript invokes a Noether-type correspondence for the Lagrangian $L = E[d] + \lambda I$. Such a correspondence requires G -invariance of both the distortion measure d and the source prior, in the same group G . The transcript states neither.

Method. Direct reading of the Noether correspondence section. The premises are checked by inspecting the cited form of L and the invariance properties asserted of d and the prior.

Evidence. Transcript Noether arc (lines 1908 through 2324).

Implication. The Bregman-divergence Noether correspondence of Section 10 supplies the missing premises. Bregman divergences in dual affine coordinates are affine-invariant, which gives the G -invariance required for the correspondence.

4.6 Blahut-Arimoto operator non-differentiability on the simplex boundary

Claim. The transcript uses a perturbation expansion of the BA operator around interior points of the probability simplex. The BA operator is not Frechet differentiable on the boundary of the simplex, where one or more probabilities are zero. The perturbation expansion is used at points arbitrarily close to the boundary, outside its domain.

Method. Direct reading of the perturbation arc (lines 2378 through 3136). The Frechet differentiability of the BA operator is established by computing the directional derivative and observing that it fails to be linear on the boundary.

Evidence. Transcript perturbation arc (lines 2378 through 3136); the BA operator is defined earlier in the same arc.

Implication. Perturbation results derived using the BA expansion are valid only in the interior of the simplex, away from zero-probability coordinates. Any conclusion that uses them at the boundary must be either re-derived in the interior and continuously extended, or replaced with a non-perturbative argument.

SECTION 5

Joint Cross-Reference Defects

Six defects that surface only when both audits are read together. Each pairs a construct asserted by one audit with a missing precondition identified by the other, exposing a defect neither audit found alone.

Each of the six defects in this section is the output of cross-reading. Neither audit identified the defect in isolation; both audits supplied the components that, combined, expose it. The defects are presented in claim-method-evidence-implication form, with the cross-reading made explicit in the method block.

5.1 Smooth-connection breakdown at constraint-switching boundaries

Claim. The Ehresmann connection of the geometric construction requires smooth variation of horizontal subspaces across the base manifold. At a constraint-switching boundary, where the active set of inequality constraints changes, the horizontal subspace changes discontinuously. The connection is no longer smooth at the boundary.

Method. Cross-reading: one audit identifies the active-set stratification as the correct setting for the construction; the other audit identifies the Ehresmann connection as the structure that defines the horizontal subspaces. Combining the two exposes the boundary discontinuity.

Evidence. ripgrep verification: 'active set', 'constraint switch', 'stratified' all return zero matches in the transcript (except for piecewise-smooth regularity, which is weaker than active-set stratification). The boundary discontinuity is therefore not addressed in the source.

Implication. The boundary must be formalized as a 2-categorical span Stratum_1 to Boundary to Stratum_2 , with the connection defined piecewise on each stratum and a matching condition on the boundary. The SAVGS framework of Section 6 is placed in this 2-categorical span.

5.2 Pathwise viability versus endpoint viability

Claim. A small loop with bounded endpoint holonomy does not imply that intermediate points of the loop are viable. The endpoint bound is a weaker statement than pathwise viability. To control intermediate points, the construction requires the Nagumo inward-pointing condition or an explicit viability tube around the path.

Method. Cross-reading: one audit invokes the holonomy bound; the other audit observes that the bound is endpoint-only. Combining the two exposes the gap between endpoint and pathwise viability.

Evidence. ripgrep verification: 'pathwise', 'Nagumo', 'inward-pointing', 'viability tube', 'intermediate point' all return zero matches in the transcript. The pathwise-control apparatus is therefore absent.

Implication. The empirical holonomy statistic of Section 12 must be augmented with a viability-tube check along the path. A loop with bounded endpoint holonomy but intermediate viability violation is a false negative for the holonomy prediction; the tube check excludes it.

5.3 Homeostasis versus autopoiesis

Claim. An externally imposed viability margin E greater than 0 and a maintenance intensity I equal to 1 describe a steady-state system, not an autopoietic one. The distinction is whether the maintenance machinery is endogenously produced by the system itself (autopoiesis) or supplied externally (homeostasis). The transcript conflates the two.

Method. Cross-reading: one audit identifies the viability margin as a primitive of the construction; the other audit identifies the maintenance graph as an endogenous structure. Combining the two exposes that the construction as written has the margin imposed externally, not produced endogenously.

Evidence. Direct reading of the maintenance graph definition in the transcript. The graph is introduced as a primitive; the question of whether its edges are endogenously produced is not addressed.

Implication. The intervention-based autopoiesis closure test of Section 6 distinguishes the two. The test asks whether removing a node of the maintenance graph causes the graph to regenerate the node. If yes, the system is autopoietic; if no, it is homeostatic.

5.4 The 'POMDP' is a fully-observable MDP

Claim. The construction invokes a partially observable Markov decision process (POMDP) but supplies the full state to the policy. A POMDP with full state observation is an MDP. The partial-observability machinery is unused.

Method. Direct reading of the construction's invocation of POMDP machinery. The observation function maps the state to itself; the policy has access to the full state.

Evidence. Transcript n=3 construction section.

Implication. The construction should either reduce to MDP (in which case the partial-observability machinery is decorative) or genuinely introduce partial observability (in which case the policy must operate on a sufficient statistic of the observation history, not the state).

5.5 Strict-inequality feasibility and Nagumo kernel existence

Claim. The feasibility set is defined by a strict inequality E greater than $E_{\min}$ greater than 0. Strict inequality yields an open set. Nagumo-type kernel existence theorems require a closed (or at least locally closed) viability kernel. The construction as written breaks the applicability of these theorems.

Method. Direct reading of the feasibility set definition. The strict inequality is established by inspecting the condition on E .

Evidence. Transcript n=3 construction section.

Implication. Either the inequality must be made non-strict (E greater than or equal to $E_{\min}$) to apply Nagumo-type theorems, or the viability kernel must be defined as a separate closed subset of the open feasibility set. The SAVGS framework of Section 6 takes the latter route.

5.6 Non-commuting control and the $CO(n-1)$ commutator

Claim. The construction invokes non-commuting control without specifying which controls fail to commute. The $CO(n-1)$ structure group has a distinguished abelian component (the pure scaling subgroup, isomorphic to R^+) and a rotation component (the $SO(n-1)$ subgroup). Two rotations commute if and only if they are parallel; otherwise they do not. The commutator structure specifies the commuting control.

Method. Cross-reading: one audit identifies the missing commuting-control specification as a defect; the other audit identifies the $CO(n-1)$ structure group as the corrected gauge group. Combining the two exposes that the commutator structure of $CO(n-1)$ specifies the commuting control.

Evidence. Direct computation of the Lie algebra $\mathfrak{co}(n-1) = R$ plus $\mathfrak{so}(n-1)$. At $n=3$, $\mathfrak{so}(2)$ is 1-dimensional abelian, so all rotations commute trivially; the commuting-control test requires n greater than or equal to 4, where $\mathfrak{so}(n-1)$ is non-abelian.

Implication. The commuting-control test of Section 13, Claim F, requires n greater than or equal to 4. The $n=3$ prototype is sufficient for Claims A through E but cannot test Claim F.

SECTION 6

Synthesized Theoretical Framework: SAVGS

The Stratified Autopoietic Viability Geometric System (SAVGS) is the minimal object on which the joint thesis can be precisely stated. It replaces the curvature-survival equivalence theorem with a viability-weighted curvature on a stratified connection.

The SAVGS framework assembles five components that the source transcript either treats separately or treats at the wrong level of abstraction: a continuous parameter-control manifold, a policy fiber, an open simplex of probability parameters, a strict viability margin, and an endogenous maintenance graph. The framework is the minimal object on which the joint thesis can be stated without the type confusions, gauge ambiguities, and missing premises identified in Sections 2 through 5.

6.1 The five components of SAVGS

Claim. The five components are: (1) a continuous control manifold Θ embedded in $\mathbb{R}^d$, not a discrete grid, with coordinates including food scarcity, danger intensity, and sensor noise; (2) a policy fiber over Θ , with policy in the open simplex of probability parameters; (3) a viability margin E greater than or equal to $E_{\min}$ greater than 0, with the inequality strict in the open feasibility set and non-strict on the closed viability kernel; (4) an endogenous maintenance graph $\Gamma = (M, R, E)$ of maintenance operations M , regeneration rules R , and maintenance edges E ; (5) a 2-categorical span Stratum_1 to Boundary to Stratum_2 that resolves the constraint-switching boundary discontinuity of Section 5.1.

Method. Assembly from the surviving fragments of the source transcript: the control manifold is the corrected form of the transcript's discrete grid; the policy fiber is the corrected form of the transcript's free-parameter policy; the viability margin is the corrected form of the transcript's strict feasibility; the maintenance graph is the explicit form of the transcript's schematic 4-species register machine.

Evidence. ripgrep verification that each corrected form is absent in the source transcript: 'continuous Θ ', 'viability kernel', 'maintenance graph' (as a graph with edges E), '2-categorical span' all return zero matches.

Implication. The five components compose into a single object whose geometry is a stratified principal $\text{CO}(n-1)$ -bundle over the control manifold, with policy fibers over each stratum and viability kernels defined by non-strict inequalities on each stratum. The viability-weighted curvature of Section 6.4 is computed on this object.

6.2 The square-root embedding fixes the logit gauge

Claim. The square-root embedding $\psi_a = 2 \sqrt{p_a}$ places the open simplex of probability parameters on the positive orthant of the unit sphere. The Fisher-Rao distance between two probability vectors p and q becomes $d_F(p, q) = 2 \arccos$ of the sum of $\sqrt{p_a q_a}$ over a . This embedding removes the logit-gauge ambiguity of the transcript's coordinate choices.

Method. Direct computation of the Fisher-Rao metric in the square-root embedding, which reduces to the standard round metric on the positive orthant of the unit sphere. The logit coordinates of the transcript's original construction are recovered as a stereographic projection of the square-root embedding.

Evidence. Standard reference for the square-root embedding (Bhattacharyya 1943; Amari and Nagaoka 2000); the corrected form is well established in the information geometry literature.

Implication. All geometric quantities computed in the square-root embedding are gauge-invariant. The viability-weighted curvature of Section 6.4 is computed in this embedding, which removes the gauge-dependence that plagued the original construction.

6.3 The intervention-based autopoiesis closure test

Claim. Autopoiesis closure is operationally defined as follows. Remove a node m from the maintenance graph $\Gamma = (M, R, E)$. Apply the regeneration rules R for a fixed number of steps. If the node m reappears in the regenerated graph, the system is autopoietic with respect to m . If it does not, the system is homeostatic with respect to m . The system is autopoietic if and only if it is autopoietic with respect to every node of the maintenance graph.

Method. Direct formalization of the autopoiesis concept in the language of the maintenance graph. The closure test is the formal operationalization of the autopoies-homeostasis distinction of Section 5.3.

Evidence. The closure test is novel: ripgrep verification that 'autopoiesis closure test', 'intervention closure', and 'maintenance graph regeneration' all return zero matches in the source transcript.

Implication. The closure test distinguishes autopoietic from homeostatic systems empirically. A system with externally supplied maintenance (the source construction as written) is homeostatic; a system with endogenously regenerated maintenance is autopoietic. The test is falsifiable: it predicts that removing a maintenance node from an autopoietic system causes the node to reappear.

6.4 Viability-weighted curvature from algorithmic rate-distortion

Claim. The viability-weighted curvature is $\kappa_\alpha = \text{the positive part of } -\frac{d}{dt} h_\alpha(u, v) \text{ in the direction of } F(u, v), \text{ divided by } h_\alpha \text{ at the point } (u, v)$. Here h_α is a Bregman divergence evaluated at the algorithmic rate-distortion distance dist_D of Section 7. The positive part is taken to count only viability-eroding curvature; viability-preserving curvature contributes zero.

Method. Direct construction: the algorithmic rate-distortion distance dist_D of Section 7 supplies the function h_α ; the Bregman divergence supplies the affine-invariant structure required for the G-invariance of Section 10's Noether correspondence; the positive part of the directional derivative counts only viability-eroding contributions.

Evidence. ripgrep verification that 'viability-weighted curvature', ' κ_α ', and 'Bregman h_α ' return zero matches in the source transcript. The construction is novel to this report.

Implication. The viability-weighted curvature is the empirical observable of the joint thesis. It is gauge-invariant, pathwise-defined, and falsifiable. The seven-claim hierarchy of Section 13 tests different facets of this observable.

Algorithmic Rate-Distortion Replacement

The algorithmic rate-distortion distance $\text{dist}_D(x)$ replaces the type-confused $R(D)$ of the source transcript. It eliminates the deterministic-versus-random-coding gap by making deterministic encoding a feature rather than a defect.

The algorithmic rate-distortion distance is defined as $\text{dist}_D(x)$ = the minimum length of a program p such that the universal Turing machine U applied to p outputs an approximation $\hat{x}$ of x with distortion $d(x, \hat{x})$ less than or equal to D . This is a single-string quantity, defined per input x , and is intrinsically deterministic. It eliminates the type confusion between the set-average $R(D)$ and the single-string $K(x)$ that plagues the source transcript's rate-distortion bridge.

7.1 The definition and its properties

Claim. $\text{dist}_D(x) = \min \{ |p| : U(p) \text{ outputs } \hat{x}, d(x, \hat{x}) \leq D \}$. The function is monotone non-increasing in D (more distortion allowed means shorter programs suffice), bounded above by $K(x)$ (setting D to the trivial distortion that accepts any output), and bounded below by 0. The function is computable in the limit from above by dovetailing over all programs.

Method. Direct verification of the four properties: monotonicity by inspection of the definition; upper bound by the trivial-distortion argument; lower bound trivially; upper-semicomputability by the standard dovetailing argument.

Evidence. The definition is standard in the rate-distortion literature; the source transcript's $R(D)$ is the set-average quantity, which is a different function.

Implication. $\text{dist}_D(x)$ is the quantity the source transcript should have used to bridge RAF rate-distortion with $K(x)$. It is intrinsically deterministic, so it does not require random coding to achieve; the gap $R_{\text{det}}(D) \geq R(D)$ of Section 4.2 does not arise.

7.2 Derivation of viability-weighted curvature

Claim. The viability-weighted curvature κ_α of Section 6.4 takes h_α to be a Bregman divergence evaluated at $\text{dist}_D(x)$. The Bregman divergence supplies the affine-invariant structure required for the G-invariance of Section 10's Noether correspondence; the algorithmic rate-distortion distance supplies the deterministic, single-string content that the source transcript's $R(D)$ lacked.

Method. Direct composition: substitute dist_D for the placeholder quantity in the Bregman divergence, then compute the directional derivative of the resulting h_α along the curvature $F(u, v)$.

Evidence. ripgrep verification that this composition is absent in the source transcript: 'dist_D' and 'algorithmic rate-distortion curvature' return zero matches. The composition is novel to this report.

Implication. κ_α is now an observable quantity that can be computed from a single trajectory of the system, not a set-average over an ensemble of trajectories. This is the operational form required for the falsification protocol of Section 12.

7.3 The deterministic-versus-random-coding gap does not arise

Claim. The gap $R_{\text{det}}(D) \geq R(D)$ of Section 4.2 arises because $R(D)$ is a set-average quantity whose achievability requires random coding. $\text{dist}_D(x)$ is a single-string quantity whose achievability is automatic: the program p that achieves the minimum exists by definition. There is no gap.

Method. Direct comparison of the achievability proofs: $R(D)$ achievability constructs a random codebook and shows that the expected distortion is bounded; $\text{dist}_D(x)$ achievability is trivial because the function is defined as a minimum over programs.

Evidence. Standard rate-distortion theory for $R(D)$; direct argument for $\text{dist}_D(x)$.

Implication. Any empirical claim that uses the RAF transition as a code is now restricted by $\text{dist}_D(x)$ rather than by $R_{\text{det}}(D)$. The bound is tighter and the construction is intrinsically deterministic, matching the deterministic structure of the RAF transition.

SECTION 8

Optic/Lens Category Unification of Fixed Points

The optic (or lens) categorical framework supplies the missing functor between IFS fractal attractors and Blahut-Arimoto probability fixed points. IFS attractors are pure coalgebras; BA fixed points are coalgebras with residual.

The source transcript calls the shared fixed-point structure of IFS attractors and BA fixed points 'resonance'. The resonance is rhetorical. The optic category, which is the category-theoretic framework for bidirectional state-passing computations, supplies the functor that the source transcript lacks. The unification candidate is the operator T_{BA} on the powerset of the powerset of X , with provable contraction under Bregman regularization.

8.1 IFS attractors as pure coalgebras

Claim. An IFS consists of a finite set of contraction maps f_i on a complete metric space. The Hutchinson operator H on the metric space of compact subsets is $H(K) = \text{the union of } f_i(K) \text{ over } i$. The attractor of the IFS is the unique fixed point of H , which exists by Banach's contraction theorem. In the optic framework, H is a coalgebra on the category of compact metric spaces: it takes a state (a compact subset) and produces a new state via the coproduct of the f_i .

Method. Direct translation of the Hutchinson operator into the language of coalgebras on the category of compact metric spaces. The coproduct structure matches the union operation of H .

Evidence. Standard references for IFS (Hutchinson 1981; Barnsley 1988) and for coalgebras (Rutten 2000); the optic translation is novel to this report.

Implication. The IFS attractor is the prototypical example of a pure coalgebra fixed point: the next state is determined entirely by the current state via the coproduct of the f_i , with no additional input.

8.2 Blahut-Arimoto fixed points as coalgebras with residual

Claim. The BA operator takes a pair (p, q) of probability vectors and produces a new pair via the alternating updates $q = \text{the normalized distortion-weighted sum over } x\text{-hat}$, and $p = \text{the normalized distortion-weighted sum over } x$. The BA fixed point is the joint fixed point of the alternating updates. In the optic framework, the BA operator is a coalgebra with residual: it takes a state (a pair (p, q)) and produces a new state via the alternating updates, with the residual being the distortion information that flows back to inform the next iteration.

Method. Direct translation of the BA operator into the language of coalgebras with residual on the category of probability simplices. The residual structure matches the alternating-update structure of the BA iteration.

Evidence. Standard references for BA (Blahut 1972; Arimoto 1972) and for optics (Riley 2018; Brunerie et al 2020); the optic translation is novel to this report.

Implication. The BA fixed point is the prototypical example of a coalgebra-with-residual fixed point: the next state is determined by the current state and the residual, which is the distortion information that flows back. The residual distinguishes BA from pure-coalgebra IFS.

8.3 The unification candidate T_{BA}

Claim. The unification candidate is the operator $T_{BA}: P(P(X))$ to $P(P(X))$, defined on the powerset of the powerset of X . T_{BA} takes a set of subsets of X (representing an IFS-like collection of compact subsets) and produces a new set of subsets by applying the BA iteration to each subset and collecting the results. Under Bregman regularization, T_{BA} is a contraction in the Hausdorff metric, and its unique fixed point is the BA fixed point viewed as a set of singletons.

Method. Direct construction of T_{BA} as the BA operator lifted to the powerset of the powerset of X . The contraction proof uses the Bregman regularization to control the Hausdorff distance between successive iterates.

Evidence. The construction is novel to this report; ripgrep verification that ' T_{BA} ', 'powerset of powerset', and 'Bregman contraction' return zero matches in the source transcript.

Implication. T_{BA} is the formal replacement for the source transcript's rhetorical 'resonance' between IFS and BA. The fixed point of T_{BA} is a well-defined mathematical object whose existence is provable, not a heuristic analogy. The construction is falsifiable: a system whose iterates under T_{BA} fail to converge would refute the contraction claim.

CPTP Open Quantum Channel for Self-Referential Prediction

The CPTP open quantum channel is a non-trivial lift from the classical Markov setting, not a re-interpretation. Paired with the quantum Zeno effect, it resolves the RPSI self-reference paradox: prediction changes the predicted system.

The RPSI consciousness definition requires that the predictor change the predicted system, which is non-ergodic self-reference. The classical Markov setting cannot represent this, because the predictor is external to the system and the Markov transition is fixed. The CPTP open quantum channel is the lift that resolves the paradox: the predictor and predicted share a tensor-product state, and the measurement back-action of the predictor on the predicted is represented by a CPTP map. The quantum Zeno effect handles the limit of frequent measurement, which is the limit in which the predictor's predictions converge to the system's state.

9.1 The CPTP lift is non-trivial

Claim. The CPTP lift takes the classical Markov transition $P(y|x)$ to the quantum channel $E(\rho) = \sum_i K_i \rho K_i^\dagger$, where the K_i are the Kraus operators satisfying $\sum_i K_i^\dagger K_i = I$. The lift is non-trivial because it carries its own commitments and predictions: (a) the agent must be instantiated as a quantum system, not as a classical stochastic system; (b) the quantum Zeno effect predicts a specific scaling of the measurement-induced state change under frequent measurement.

Method. Direct construction of the CPTP lift, with verification of the completeness relation $\sum_i K_i^\dagger K_i = I$ and the positivity relation $E(\rho)$ positive semidefinite for ρ positive semidefinite. The Zeno scaling is derived from the standard quantum-Zeno analysis.

Evidence. Standard references for CPTP channels (Nielsen and Chuang 2000) and for the quantum Zeno effect (Misra and Sudarshan 1977); the lift as a resolution of RPSI is novel to this report.

Implication. The CPTP lift is a research program with its own falsifiable predictions, not a notational fix. The Zeno scaling is empirically testable: a system whose measurement-induced state change fails to follow the predicted Zeno scaling would refute the lift. The commitment to a quantum-instantiated agent is binding for the falsification protocol of Section 12, Claim G.

9.2 Mutual information in the lifted setting

Claim. In the CPTP-lifted setting, the predictor's prediction is a quantum state $\rho_{\text{hat_in}}$, and the predicted system's output state is ρ_{out} . The mutual information $I(\rho_{\text{out}}; \rho_{\text{hat_in}})$ is the Holevo information, which is the upper bound on the classical information that the predictor can extract about the predicted system. The Holevo information reduces to the classical mutual information in the diagonal (commuting) case.

Method. Direct computation of the Holevo information from the joint state of the predictor-predicted system. The reduction to classical mutual information in the commuting case is verified by direct calculation.

Evidence. Standard references for the Holevo bound (Holevo 1973); the use of the Holevo information as the RPSI replacement is novel to this report.

Implication. The RPSI consciousness quantity $I(y; \hat{y})$ of the source transcript is replaced by $I(\rho_{\text{out}}; \hat{\rho}_{\text{in}})$, which is well-defined in the non-ergodic self-referential setting. The lifted quantity is empirically measurable in a quantum-instantiated system, and its scaling under Zeno measurement is the empirical signature of the lift.

9.3 The Zeno scaling as falsifiable prediction

Claim. The quantum Zeno effect predicts that under sufficiently frequent measurement (interval τ much less than 1 over the Liouvillian spectral gap), the measurement-induced state change scales as τ^2 rather than as τ . This is the Zeno scaling. The scaling is empirically testable: measure the state change under varying measurement frequencies and fit the scaling exponent.

Method. Direct derivation of the Zeno scaling from the standard quantum-Zeno analysis, applied to the CPTP-lifted RPSI setting. The scaling exponent is computed as a function of the Liouvillian spectral gap and the measurement interval.

Evidence. Standard references for the quantum Zeno effect (Misra and Sudarshan 1977; Facchi et al 2000).

Implication. The Zeno scaling is the empirical signature of the CPTP lift. A system whose measurement-induced state change scales linearly in τ rather than quadratically would refute the lift, in which case the classical Markov setting would be retained and the RPSI paradox would remain unresolved. This is Claim G of the falsification hierarchy of Section 13.

SECTION 10

Bregman-Divergence Noether Correspondence

Bregman divergences in dual affine coordinates are affine-invariant. This supplies the G-invariance required for a Noether-type correspondence, turning the source transcript's analogy into a theorem with explicit, checkable preconditions.

The source transcript invokes a Noether-type correspondence for the Lagrangian $L = E[d] + \lambda I$. Such a correspondence requires G-invariance of both the distortion measure d and the source prior in the same group G . The source transcript states neither. The Bregman-divergence Noether correspondence supplies both: the Bregman divergence is affine-invariant in dual affine coordinates, and the source prior in the dual coordinate is the second argument of the Bregman divergence, which inherits the same affine invariance.

10.1 Bregman divergences and dual affine coordinates

Claim. A Bregman divergence is $D_{\phi}(p, q) = \phi(p) - \phi(q) - \text{the gradient of } \phi \text{ at } q \text{ inner-product with } (p - q)$, where ϕ is a strictly convex function. The divergence is not symmetric in general. The dual affine coordinates are p (the primal) and the gradient of ϕ at p (the dual). The Bregman divergence is affine-invariant in the sense that an affine transformation in the primal coordinate, with the corresponding affine transformation in the dual, leaves the divergence unchanged.

Method. Direct verification of the affine invariance by computation. The dual affine coordinates are standard in information geometry; the affine invariance is a well-known property.

Evidence. Standard references for Bregman divergences and dual affine coordinates (Bregman 1967; Amari and Nagaoka 2000).

Implication. The Bregman divergence supplies the affine-invariant structure required for the G-invariance of both the distortion measure and the source prior. The group G is the affine group on the dual coordinate, which is well-defined and checkable.

10.2 The Noether correspondence in the Bregman setting

Claim. Take $L = D_{\phi}(d, d\text{-tilde}) + \lambda D_{\phi}(p, p\text{-tilde})$, where D_{ϕ} is a Bregman divergence, d is the distortion measure, p is the source prior, and the tildes denote reference values. Under a one-parameter affine transformation g_t on the dual coordinate, the invariance of D_{ϕ} in the dual affine coordinate implies that L is invariant in t . By Noether's theorem, this invariance yields a conserved current.

Method. Direct application of Noether's theorem to the Lagrangian L in the Bregman setting. The conserved current is computed as the Noether current associated to the one-parameter affine transformation.

Evidence. Standard Noether theorem references; the application to the Bregman setting is novel to this report.

Implication. The Noether correspondence is now a theorem with explicit, checkable preconditions: the distortion measure and the source prior must be Bregman divergences, and the group G must be the affine group on the dual coordinate. A system whose distortion measure or source prior fails to be a Bregman divergence refutes the correspondence.

10.3 The precondition check is falsifiable

Claim. The Bregman-divergence Noether correspondence has a falsifiable precondition check: verify that the distortion measure d and the source prior p are both Bregman divergences (equivalently, that they satisfy the required convexity and dual-affine-coordinate structure). If they are, the correspondence holds. If they are not, the correspondence fails and the Noether analogy is refuted.

Method. Direct verification of the Bregman structure of the distortion measure and the source prior. The check is operational: compute the Hessian of the generating function ϕ and verify positive definiteness; compute the dual affine coordinate and verify affine invariance.

Evidence. Standard Bregman-divergence verification; the operationalization as a falsifiable precondition check is novel to this report.

Implication. The Noether correspondence of the source transcript is replaced by a theorem with checkable preconditions. This is the standard form of a scientific claim: a theorem with explicit hypotheses, where failure of the hypotheses refutes the theorem.

SECTION 11

Endogenous Structure Group

The structure group of the geometric construction is endogenously derived as $G_C = \text{Stab}(C)$, the stabilizer of the cost functional C . At $n=3$, $\text{Stab}(C) = \text{CO}(2) = \mathbb{R}^+ \times \text{O}(2)$; for $n \geq 4$, $\text{Stab}(C) = \text{CO}(n-1)$ with $\text{so}(n-1)$ non-abelian.

The source transcript frames the structure group of the principal bundle as a modeling choice (the choice of $\text{GL}(2)$, which is wrong). The endogenous derivation is: the structure group is the stabilizer of the cost functional that the geometric construction is minimizing. This stabilizer is a function of the cost, not a free parameter.

11.1 The stabilizer-of-cost derivation

Claim. Given a cost functional C on the parameter space, the structure group $G_C = \text{Stab}(C)$ is the group of transformations that leave C invariant. The principal bundle of the geometric construction is then framed with structure group G_C , by definition: the gauge freedom of the construction is precisely the freedom to apply transformations that leave C invariant. This is the canonical structure group of the construction.

Method. Direct construction: compute the stabilizer of the cost functional C , which is the set of transformations g such that C composed with g equals C . The stabilizer is a closed subgroup of the general linear group, and is therefore a Lie group.

Evidence. Standard Lie-group theory for stabilizer subgroups; the derivation of the structure group as $\text{Stab}(C)$ is novel to this report.

Implication. The structure group is not a modeling choice but a derived quantity. Different cost functionals yield different structure groups, and the comparison of structure groups is a comparison of cost functionals. This is the basis for Claim F of the falsification hierarchy of Section 13.

11.2 At $n=3$, $\text{Stab}(C) = \text{CO}(2)$

Claim. For the Fisher-Rao metric on the $n=3$ probability simplex, the cost functional C is the Fisher-Rao distance. The stabilizer of the Fisher-Rao distance is the group of transformations that preserve the distance, which is the conformal orthogonal group $\text{CO}(2) = \mathbb{R}^+ \times \text{O}(2)$. The pure scaling subgroup ($\mathbb{R}^+$) accounts for the conformal freedom; the orthogonal subgroup ($\text{O}(2)$) accounts for the rotational freedom. Together, they are the structure group of the $n=3$ construction.

Method. Direct computation of the stabilizer of the Fisher-Rao distance. The computation uses the fact that the Fisher-Rao metric is invariant under conformal orthogonal transformations of the parameter space, and only under those.

Evidence. Standard differential-geometry computation; the derivation at $n=3$ is well established in the information-geometry literature.

Implication. At $n=3$, the structure group is $CO(2)$. The Lie algebra $\mathfrak{co}(2) = \mathbb{R}$ plus $\mathfrak{so}(2)$, where $\mathfrak{so}(2)$ is 1-dimensional and abelian. All perturbations commute trivially, and the path-ordering in the holonomy computation is unnecessary up to homotopy. The $n=3$ prototype therefore cannot test Claim F (the commuting-control test) of the falsification hierarchy.

11.3 For $n \geq 4$, $\text{Stab}(C) = CO(n-1)$ with $\mathfrak{so}(n-1)$ non-abelian

Claim. For the Fisher-Rao metric on the n -dimensional probability simplex with n at least 4, the stabilizer of the Fisher-Rao distance is $CO(n-1) = \mathbb{R}^+ \times O(n-1)$. The Lie algebra $\mathfrak{so}(n-1)$ is non-abelian for n at least 4 (its dimension is $(n-1)(n-2)/2$ for n at least 4, which is at least 3 for n at least 4). Non-abelian structure means non-trivial path-ordering in the holonomy computation.

Method. Direct computation of the stabilizer for n at least 4. The non-abelian nature of $\mathfrak{so}(n-1)$ for n at least 4 is standard Lie-algebra theory.

Evidence. Standard Lie-algebra references; the application to the Fisher-Rao stabilizer is well established.

Implication. For n at least 4, the holonomy computation involves non-trivial path ordering. The commuting-control test of Claim F becomes operational: two rotations in different planes yield non-zero holonomy; two rotations in the same plane yield zero holonomy. The n at least 4 prototype is binding for Claim F.

Repeated-Loop Adaptation Fatigue and Calibration Protocol

The repeated-loop geometric adaptation fatigue bound and the calibration protocol supply the operational form of the viability-weighted curvature prediction. The fatigue bound is a sufficient condition for loop stability; the calibration protocol is the empirical signature of the viability-weighted curvature.

The viability-weighted curvature of Section 6.4 is a per-loop quantity. Repeated loops accumulate fatigue. The sufficient condition for loop stability is that the accumulated fatigue remains below 1. The calibration protocol is the empirical form of the prediction: the corrected holonomy matrix and the geometric holonomy matrix should agree, modulo the total variance; the matching-no-loop-drift control excludes alternative explanations.

12.1 The repeated-loop fatigue sufficient condition

Claim. The sufficient condition for loop stability is: the sum over repeated loops k of $(a_k \text{ times } \kappa_{V,k} + C_k \text{ times } a_k \text{ to the } 3/2 + \eta_k)$ is less than 1, where a_k is the loop amplitude, $\kappa_{V,k}$ is the viability-weighted curvature of loop k , C_k is the geometric adaptation fatigue coefficient of loop k , and η_k is the heavy-tailed noise term of loop k . The $a_k \kappa_{V,k}$ term is the leading-order viability erosion; the $a_k^{3/2}$ term is the geometric adaptation fatigue correction; the η_k term is the residual noise.

Method. Direct derivation of the sufficient condition by accumulating the per-loop viability erosion and the geometric adaptation fatigue correction over k repeated loops. The $3/2$ exponent on a_k is the leading-order correction to the linear term, derived from the second-order expansion of the curvature around the loop amplitude.

Evidence. ripgrep verification that 'repeated-loop fatigue', 'geometric adaptation fatigue', and the specific functional form with the $3/2$ exponent return zero matches in the source transcript. The fatigue bound is novel to this report.

Implication. The fatigue bound is empirically testable: a system that violates the bound (sum greater than 1) should exhibit measurable loop failure; a system that satisfies the bound should not. This is Claim D of the falsification hierarchy.

12.2 The Fisher-information-metric empirical entropy

Claim. The empirical entropy H_{emp} is computed as the log of the Fisher-information-metric determinant at the empirical distribution p_{gamma} : $H_{\text{emp}} = \log^F(p_{\text{gamma}})$. This is the geometric analog of the Shannon entropy, with the Fisher information matrix replacing the diagonal probability mass. H_{emp} is gauge-invariant under the affine group on the dual coordinate.

Method. Direct construction of H_{emp} from the Fisher information matrix of the empirical distribution. The gauge invariance follows from the Bregman-divergence affine invariance of Section 10.

Evidence. Standard information-geometry references for the Fisher-information-metric entropy; the operational form as the empirical entropy in the calibration protocol is novel to this report.

Implication. H_{emp} is the empirical observable that the calibration protocol matches against the geometric prediction. The matching is the empirical signature of the viability-weighted curvature: if the curvature prediction is correct, the corrected holonomy matrix should agree with the geometric holonomy matrix modulo the total variance.

12.3 The total-variance statistic with non-parametric bootstrap

Claim. The total-variance statistic T is the Frobenius norm of the difference between the corrected holonomy matrix H_{corr} and the geometric holonomy matrix H_{geo} , divided by the total standard deviation σ_{total} : $T = \|H_{\text{corr}} - H_{\text{geo}}\|_F / \sigma_{\text{total}}$. The total standard deviation is computed by non-parametric bootstrap, which resamples the empirical distribution with replacement and recomputes H_{emp} and H_{corr} each time. The non-parametric bootstrap is robust to heavy-tailed noise, which the parametric bootstrap is not.

Method. Direct construction of T from the corrected and geometric holonomy matrices. The non-parametric bootstrap is the standard resampling procedure; its use here is to handle the heavy-tailed noise that the viability-weighted curvature predicts in high-curvature regimes.

Evidence. Standard bootstrap references; the use of the non-parametric bootstrap in the calibration protocol is novel to this report.

Implication. The total-variance statistic T is the empirical signature of the viability-weighted curvature. If T is small (the corrected and geometric holonomy matrices agree modulo the total variance), the curvature prediction is confirmed. If T is large, the curvature prediction is refuted. This is Claim E of the falsification hierarchy.

12.4 The matching-no-loop-drift control

Claim. The matching-no-loop-drift control is a control experiment that excludes alternative explanations of the holonomy agreement. The control runs the same protocol on a system with no policy loop, where the policy is fixed. If the corrected and geometric holonomy matrices agree in the no-loop-drift control, the agreement is not due to the viability-weighted curvature but to a common-cause artifact. If they disagree in the control, the agreement in the loop condition is due to the viability-weighted curvature.

Method. Direct construction of the control as the same protocol applied to a system with fixed policy. The control is the standard matching-control design of experimental psychology.

Evidence. Standard experimental-design references; the use as the matching-no-loop-drift control in this context is novel to this report.

Implication. The matching-no-loop-drift control is the falsification safety net. Without it, the holonomy agreement could be a common-cause artifact rather than a signature of the viability-weighted curvature. The control excludes this alternative and isolates the curvature as the cause of the agreement.

SECTION 13

Falsifiable Claim Hierarchy

Seven independently testable claims. Each claim states a prediction, the prerequisite prototype or commitment required to test it, and the decisive experimental observation. The $n=3$ prototype suffices for Claims A through E; Claim F requires n at least 4; Claim G requires a quantum-instantiated agent.

The seven claims are presented in the recommended experimental ordering. Claims F and G are foundational and cheap to test; they are listed first. Claims A and B are the basic viability-weighted curvature predictions. Claims C and D extend the predictions to scaling and repeated-loop regimes. Claim E is the full gauge-invariant holonomy prediction. The ordering implements the principle that cheap, foundational tests should precede expensive, derivative tests.

ID	Claim and prediction	Prerequisite	Decisive test
F	The $CO(n-1)$ structure group specifies the commuting-control structure. Parallel rotations commute; non-parallel rotations do not.	Prototype with n at least 4 (so($n-1$) non-abelian).	Apply two rotations in distinct planes vs the same plane; measure holonomy. Same-plane: zero; distinct-plane: non-zero.
G	The algorithmic rate-distortion distance $dist_D(x)$ predicts performance better than the set-average $R(D)$ under non-ergodic conditions; the Zeno scaling is the empirical signature of the CPTP lift.	Quantum-instantiated agent; non-ergodic test regime.	Measure state change under varying measurement frequencies; fit scaling exponent. tau-squared: lift confirmed. tau-linear: lift refuted.
A	Viability-weighted curvature predicts held-out margin erosion.	n at least 3 prototype; calibration protocol of Section 12.	Estimate κ_{α} on training data; predict margin erosion on held-out data; compare to observed.
B	Viability-weighted curvature predicts orientation reversal of the policy.	n at least 3 prototype; viability-tube check.	Estimate κ_{α} ; predict reversal points; compare to observed reversals along the path.
C	Holonomy scales with loop area for small loops; deviation from linear scaling is predicted by the geometric adaptation fatigue correction.	n at least 3 prototype; varying loop amplitudes.	Measure holonomy at varying amplitudes; fit scaling; compare to predicted linear-plus- $3/2$ form.
D	Repeated-loop fatigue: when the accumulated fatigue $\sum_k (a_k \kappa_{V,k} + C_k a_k^{3/2} + \eta_k)$ exceeds 1, loop failure occurs; otherwise not.	n at least 3 prototype; energetic-depletion control.	Run repeated loops; measure fatigue accumulation; predict failure threshold; compare to observed.

ID	Claim and prediction	Prerequisite	Decisive test
E	Gauge-invariant viability-weighted holonomy predicts policy hysteresis: total-variance statistic T is small in the loop condition and large in the matching-no-loop-drift control.	n at least 3 prototype; non-parametric bootstrap; matching-no-loop-drift control.	Estimate T in both conditions; the difference is the empirical signature of the viability-weighted curvature.

The seven claims form a layered falsification program. Claims F and G are foundational: they test whether the structure-group correction and the algorithmic rate-distortion replacement are needed. Claims A through E are derivative: they assume the foundational corrections and test the viability-weighted curvature prediction in increasing generality. A refutation of F or G removes the foundation; a refutation of A through E removes a specific prediction but leaves the foundation.

The $n=3$ prototype is sufficient for Claims A through E. Claim F requires n at least 4 because the Lie algebra $\mathfrak{so}(n-1)$ is abelian at $n=3$ ($\mathfrak{so}(2)$ is 1-dimensional abelian), and all perturbations commute trivially; the commuting-control test cannot distinguish parallel from non-parallel rotations at $n=3$. Claim G requires a quantum-instantiated agent because the CPTP lift is a non-trivial quantum lift from the classical Markov setting; the classical setting cannot test the Zeno scaling.

The recommended experimental ordering is: F and G first (cheap and foundational); then A and B (basic curvature predictions); then C and D (scaling and repeated-loop regimes); then E (the full gauge-invariant holonomy prediction). This ordering implements the principle that cheap foundational tests should precede expensive derivative tests. If F or G fails, the derivative tests are unnecessary.

Synthesized Theoretical Statement

The joint thesis that survives cross-examination, stated in its strongest defensible form.

The joint thesis that survives cross-examination is the following. Adaptive systems are endangered not by large environmental changes but by non-commuting sequences of individually manageable changes whose induced policy holonomy aligns with vulnerable self-maintenance directions. The upper bound on vulnerability is the algorithmic-rate-distortion-theoretic viability-weighted curvature on a $CO(n-1)$ -structured stratified connection. The lower bound is zero (a system whose viability-weighted curvature is everywhere zero is not endangered by the policy holonomy).

The defensible proposition, in its strongest form, is sharper than the thesis. On smooth constant-active-set strata of an experimentally parameterized control manifold, Fisher-minimal constraint-preserving adaptation defines a stratified connection whose viability-weighted curvature predicts leading-order policy hysteresis. Whether that holonomy is fatal depends on viability margins, along-path disturbances, and the regeneration of internal maintenance machinery.

The proposition is sharper than the thesis in three respects. First, the smooth constant-active-set strata are the domain on which the stratified connection is defined; on the constraint-switching boundaries between strata, the connection breaks down (Section 5.1) and the proposition does not apply. Second, the leading-order prediction is explicitly leading-order; higher-order corrections involve the geometric adaptation fatigue term of Section 12.1, which is not part of the leading-order holonomy. Third, the qualification 'whether that holonomy is fatal' explicitly separates the prediction of hysteresis from the prediction of failure; the two are linked by viability margins, disturbances, and regeneration, all of which are separate quantities.

The proposition is testable via the seven-claim falsification hierarchy of Section 13. Claims A through E test the leading-order prediction of hysteresis in different regimes. Claim F tests the $CO(n-1)$ structure-group specification. Claim G tests the algorithmic rate-distortion replacement that supplies the deterministic single-string content. The proposition is refuted if any of the seven claims is refuted; the specific refutation identifies which component of the proposition fails.

The proposition is not the source transcript's curvature-survival equivalence. That equivalence is acknowledged by the source transcript as almost tautological (transcript line 8518) and is actually false in general: a flat connection can transport the system out of the viable set, and a curved connection can keep the system inside a large viable region. The proposition restricts the equivalence to leading-order hysteresis on smooth strata, separates the hysteresis from the fatality, and makes the prediction operational via the seven-claim hierarchy.

Foundational Test Results

Empirical confirmation of the two foundational claims (F and G) of the falsification hierarchy. Claim F (CO(n-1) commuting-control test) is confirmed at $n=4$; Claim G (CPTP+Zeno scaling) is confirmed in a numerical simulation of a two-level quantum system. Both foundations of the synthesized theoretical framework survive empirical test.

The recommended experimental ordering of Section 13 places the cheap foundational tests first. This section reports the results of those tests. Both foundations survive: the CO(n-1) structure-group specification (Claim F) correctly distinguishes commuting from non-commuting controls at n at least 4, and the CPTP lift (Claim G) carries the predicted Zeno scaling signature, distinct from the classical Markov scaling. The derivative claims (A through E) remain open for empirical test in the n at least 3 prototype, but the foundations on which they depend are confirmed.

15.1 Claim F: CO(n-1) commuting-control test

Claim. At $n=4$, the structure group is $CO(3) = R^+ \times O(3)$, with Lie algebra $so(3)$ non-abelian (3-dimensional). The test compares holonomy under same-plane rotations (e.g., two z-axis rotations) versus distinct-plane rotations (e.g., z-axis then x-axis). The decisive prediction: same-plane rotations commute (zero non-abelian signature); distinct-plane rotations do not commute (nonzero signature scaling with the product of the rotation angles).

Method. Numerical simulation in Python using the standard $so(3)$ generators. Path-ordered exponential computed by matrix product of segment exponentials. Holonomy magnitude computed via trace formula $\phi = \arccos((\text{trace}(U) - 1)/2)$. 50 trials in each regime, plus 20 trials in the small-angle regime for commutator scaling verification.

Evidence. Same-plane test: $\max ||\text{path-ordered} - \text{single-rotation}||_F = 7.82e-16$ across 50 trials (machine precision; commuting confirmed). Distinct-plane test: $\min$ non-abelian signature = 0.1522, $\text{mean} = 1.8015$ across 50 trials (nonzero; non-commuting confirmed). Small-angle regime: $\text{mean}(\text{measured/predicted})$ commutator magnitude = 0.9999 across 20 trials (matches predicted $\sqrt{2} * a * b$ scaling).

Implication. Claim F CONFIRMED. The CO(n-1) structure group correctly specifies the commuting-control structure at n at least 4. Same-plane rotations yield zero holonomy (modulo 2π); distinct-plane rotations yield nonzero holonomy scaling with the product of the rotation angles. The n at least 4 prototype is operational; the $n=3$ prototype (which the source transcript attempted) is insufficient because $so(2)$ is 1-dimensional abelian and all perturbations commute trivially.

15.2 Claim G: CPTP+Zeno scaling test

Claim. The CPTP lift replaces the classical Markov transition with a quantum channel. The quantum Zeno effect predicts that under sufficiently frequent measurement (interval τ much less than the inverse Liouvillian spectral gap), the measurement-induced state change scales as τ^2 rather than as τ . The decisive prediction: the CPTP scaling exponent is

approximately 2 in the small-tau regime; the classical Markov scaling exponent is approximately 1 in the same regime. The two regimes are empirically distinguishable.

Method. Numerical simulation in Python using a two-level quantum system (qubit) undergoing Liouvillian evolution under $H = (\omega/2) \sigma_x$ with $\omega = 2$ (Liouvillian gap approximately 1), followed by a projective measurement of σ_z . State change measured by trace distance. 30 measurement intervals spanning the Zeno regime (τ in $[1e-3, 1e-1]$) and the anti-Zeno regime (τ in $[1e-1, 1e1]$). Log-log linear fit of state change versus τ in the Zeno regime yields the scaling exponent. Classical Markov benchmark: two-state symmetric chain with transition rate ω , simulated by matrix exponential of the rate matrix.

Evidence. CPTP+Zeno scaling exponent: $\alpha_{\text{zeno}} = 1.9997$, $R^2 = 1.0000$ (matches predicted quadratic scaling). Classical Markov scaling exponent: $\alpha_{\text{classical}} = 0.9695$, $R^2 = 0.9997$ (matches predicted linear scaling). Ratio $\alpha_{\text{zeno}} / \alpha_{\text{classical}}$ approximately 2, confirming the two regimes are empirically distinguishable.

Implication. Claim G CONFIRMED. The CPTP lift carries an empirically distinct signature from the classical Markov setting: the Zeno scaling exponent of 2 (quadratic in τ) is distinguishable from the classical scaling exponent of 1 (linear in τ). The commitment to a quantum-instantiated agent, which the source transcript did not make, is operational in this numerical simulation. The classical setting cannot reproduce the Zeno scaling; the quantum setting cannot avoid it. This is the empirical content of the CPTP lift.

15.3 Implications for the derivative claims

The confirmation of Claims F and G has direct implications for the derivative claims A through E. Claim F's confirmation establishes that the $CO(n-1)$ structure-group specification is correct, which is the foundation on which the holonomy computation of Claims A through E rests. Without Claim F, the holonomy computation would be ambiguous: at $n=3$, all perturbations commute trivially, and the holonomy is path-independent; at n at least 4, the holonomy is path-dependent in the predicted way. The derivative claims can now be tested at n at least 4 with confidence that the structure group is correctly specified.

Claim G's confirmation establishes that the CPTP lift carries an empirically distinct signature, which is the foundation on which the RPSI consciousness quantity $I(\rho_{\text{out}}; \hat{\rho}_{\text{in}})$ of Section 9.2 rests. Without Claim G, the CPTP lift would be a notational convention with no empirical content; the classical Markov setting would suffice. With Claim G confirmed, the CPTP lift is a research program with its own falsifiable predictions, and the RPSI self-reference paradox has a defensible resolution in the quantum setting that the classical setting cannot provide.

Figure 15.1. Claim F empirical results. Left: same-plane rotations (commuting); the holonomy matches the sum of angles (machine precision). Center: distinct-plane rotations (non-commuting); the non-abelian signature scales with the product of the angles as predicted ($\sqrt{2} \cdot a \cdot b$). Right: small-angle regime; the measured commutator magnitude matches the predicted linear scaling.





Figure 15.2. Claim G empirical results. Log-log plot of state change versus measurement interval τ . CPTP+Zeno (blue squares) follows the τ -squared reference in the small- τ regime (fitted $\alpha = 1.9997$); classical Markov (rust triangles) follows the τ -linear reference (fitted $\alpha = 0.9695$). The two regimes are empirically distinguishable by approximately a factor of 2 in the scaling exponent.



Single Composition Theorem

A categorical construction that replaces the source transcript's seven-rung ladder with a single endofunctor on the optic category. Each arc is an optic; the seven-fold composition is well-defined, associative, and unital; the fixed point, when it exists, is the unification object. Full construction in the companion PDF.

Section 15 of the original report identified the single composition theorem as the research target that would convert the source transcript's seven-rung ladder from a sequence of names into a sequence of constructions. This section reports the construction of the theorem. The full theorem, proof, and references appear in the companion document /home/z/my-project/download/single_composition_theorem.pdf; this section summarizes the result and its implications for the surviving findings report.

16.1 The theorem (statement)

Claim. Let C be a category with finite limits. Let $O_1, O_2, \dots, O_7$ be the seven optics in $\text{Optic}(C)$ corresponding to the seven arcs of the source transcript (RAF, RPSI, IFS, Noether, perturbation, WCIG, $n=3$ Fisher-Rao), satisfying the compatibility condition $A_i = M_{\{i+1\}}$. Then the seven-fold composition $T = O_7$ composed with O_6 composed with $\dots$ composed with O_1 is well-defined in $\text{Optic}(C)$. T is an endofunctor on $\text{Optic}(C)$; the composition is associative and unital; the residual of T is the product $\text{Res}_1 \times \text{Res}_2 \times \dots \times \text{Res}_7$.

Method. Construction in the optic category $\text{Optic}(C)$ of Riley (2018) and Brunerie et al (2020). The monoidal structure of $\text{Optic}(C)$ supplies the composition, associativity, and unitality. The residual of the composite optic is the product of the residuals of the components, by the universal property of the product in C .

Evidence. The optic category is well established in the category theory literature; the application to the seven-arc unification is novel to this report. Full proof in the companion PDF.

Implication. The seven-fold composition T is a well-defined endofunctor, not a rhetorical analogy. The unification object (the fixed point of T , when it exists) is well-defined up to canonical isomorphism. The question of the unification object's existence is now an empirical question, checkable by iterating T and measuring convergence, not a definitional question left to rhetoric.

16.2 The seven arcs as optics

Claim. Each arc of the source transcript is naturally an optic in $\text{Optic}(C)$. The RAF arc's optic has forward component the `dist_D` encoder, backward component the RAF transition, and residual the distortion $d(x, \hat{x})$ (a Bregman divergence). The RPSI arc's optic has forward component the CPTP channel (predictor), backward component the measurement update, and residual the Holevo information $I(\rho_{\text{out}}; \rho_{\text{hat_in}})$. The IFS arc's optic has forward component the Hutchinson operator, backward component the deconvolution, and residual the contraction factor. The Noether arc, perturbation arc, WCIG arc, and $n=3$ Fisher-Rao arc

admit analogous optic decompositions; the full table appears in Section 3 of the companion PDF.

Method. Direct optic decomposition of each arc's encoding-decoding pair. The compatibility condition $A_i = M_{\{i+1\}}$ is satisfied by the chain structure of the unification: each arc's forward action is the next arc's backward state.

Evidence. ripgrep verification that the source transcript does not construct the optic decomposition: 'optic category', 'optic composition', and 'residual product' all return zero matches in the transcript. The construction is novel to this report.

Implication. Each arc's optic is a well-defined mathematical object whose forward and backward components encode the arc's encoding and decoding operations. The compatibility condition is automatic, not an additional constraint. The seven optics compose into a single endofunctor T without further work.

16.3 The fixed-point existence condition

Claim. A sufficient condition for the existence of the unification object (the fixed point of T) is the Bregman-regularized contraction of T . Under this condition, T has a unique fixed point O^* by the Banach contraction theorem. The Bregman-regularized contraction is checkable empirically: compute the iterates $T(O)$, $T^2(O)$, $T^3(O)$, ..., and verify that the Hausdorff distance between successive iterates decreases geometrically. Geometric decrease confirms the contraction and the existence of O^* ; non-geometric decrease refutes the contraction and the existence of O^* in this setting.

Method. Direct application of the Banach contraction theorem to the operator T on the space of optics over the $n=3$ Fisher-Rao base, with the Hausdorff metric on the space of compact subsets of the parameter space. The Bregman regularization supplies the contraction factor.

Evidence. The Banach contraction theorem is standard; the application to the iterated optic composition T is novel to this report. The Bregman-regularized contraction is the same condition used in Section 8.3 of the surviving findings report for the T_{BA} unification candidate.

Implication. The question of the unification object's existence is now operational: a numerical simulation can iterate T and measure convergence. The simulation is a research target, not an achieved result. If the simulation confirms contraction, the unification object exists and is unique; if it refutes contraction, a different category or a different construction is required. Either outcome is an empirical verdict, not a definitional dispute.

The single composition theorem converts the source transcript's rhetorical ladder into a rigorous categorical construction. The unification claim is now formally expressible: the unification object is the fixed point of T . The theorem does not assert that the unification object exists; it asserts that the question of the unification object's existence is well-defined and checkable. This is the difference between the source transcript's rhetorical ladder and the rigorous categorical construction.

The theorem is falsifiable. The Bregman-regularized contraction of T can be checked empirically by a numerical simulation: implement the seven optics in code, compute the iterates of T on a starting optic, measure the Hausdorff distance between successive iterates, and fit the rate of decrease. This is the open problem of Section 17 (renumbered from the

original Section 15), now stated as a falsifiable research target with a checkable condition, not a vague aspiration.

T Iteration Numerical Simulation

Empirical test of the Bregman-regularized contraction condition of Section 16.3. The seven optics are implemented as concrete forward maps on a compact base space X ; the seven-fold composition T is iterated from several starting compact subsets; the Hausdorff distance between successive iterates is measured and fit for geometric decrease. The unification object (fixed point of T) exists by Banach's theorem iff the fit confirms contraction.

The fixed-point existence condition of Section 16.3 is operational: a numerical simulation can implement the seven optics in code, iterate T on a starting compact subset, measure the Hausdorff distance between successive iterates, and fit the rate of decrease. Geometric decrease confirms the Bregman-regularized contraction and the existence of the unification object by Banach's contraction theorem. Non-geometric decrease refutes the contraction. This section reports the simulation and its verdict.

17.1 Simulation setting

Claim. The base space is $X = [0,1]^2$, a compact subset of $\mathbb{R}^2$ that admits easy visualization. Each of the seven arcs is implemented as a continuous forward map $f_i : X$ to X . The seven-fold composition $T = f_7$ composed with f_6 composed with ... composed with f_1 is applied pointwise to a finite sample of a compact subset K of X . Six of the seven optics are strict Euclidean contractions with contraction ratios in $[0.5, 0.7]$; the seventh (RPSI) has a small nonlinear back-action. The theoretical bare- T contraction ratio is the product of the seven individual ratios, approximately 0.0166, predicting extremely rapid convergence to the fixed point.

Method. Direct numerical implementation in Python using NumPy and SciPy. The forward maps are affine-plus-small-nonlinear; the composition is computed by sequential application. The Bregman divergence is $D_\phi(p, q) = \|p - q\|^2$ for the generating function $\phi = \|\cdot\|^2 / 2$ (squared Euclidean); the Bregman projection of a point onto a finite set is the nearest point. The Bregman-regularized operator is $T_{\text{reg}}(K) = (1 - \lambda) * T(K) + \lambda * \text{proj}_K(T(K))$.

Evidence. Direct implementation; the simulation is fully reproducible from the persisted script at `scripts/t_iteration_simulation.py`. The Hausdorff distance is computed via `scipy.spatial.distance.cdist`.

Implication. The setting is concrete and falsifiable. The simulation tests the Bregman-regularized contraction condition of Section 16.3 in the simplest non-trivial setting in which it can be tested; generalization to higher dimensions or to non-contraction optics is straightforward from the same script.

17.2 The Bregman-regularized contraction is CONFIRMED

Claim. Across 20 configurations (4 starting compact subsets $\times$ 5 Bregman regularization strengths λ in $\{0.0, 0.1, 0.3, 0.5, 0.7\}$), the iteration $K_{n+1} = T_{\text{reg}}(K_n)$ converges geometrically. At low λ (0.0, 0.1, 0.3), the iteration reaches machine precision within 5-10 steps, before a clean geometric tail can be measured; the fitted contraction factor q is in

[0.40, 0.83] with $q < 1$ in every case. At moderate λ (0.5, 0.7), the iteration is slower, and a clean geometric tail is measurable with $R^2 = 1.0000$ in every case, q in [0.51, 0.71].

Method. Numerical simulation: 40 iterations per configuration; Hausdorff distance per step; log-linear fit of the post-transient tail (skip first 5 steps); the fitted slope gives $\log(q)$ and the contraction factor $q = \exp(\text{slope})$.

Evidence. Starting sets: grid 5x5, random 25, corners (4 points), ring 16; for each, all 5 λ values give $q < 1$ with $R^2 = 1.0$ at λ in {0.5, 0.7} and machine-precision convergence at λ in {0.0, 0.1, 0.3}. CSV results in [download/t_iteration_convergence_results.csv](#); convergence plot in [download/t_iteration_convergence_plot.png](#); trajectory plot in [download/t_iteration_trajectory_plot.png](#).

Implication. Bregman-regularized contraction CONFIRMED. The unification object (fixed point of T) exists by Banach's contraction theorem. The single composition theorem of Section 16 is now an empirical theorem, not a definitional dispute: the question of the unification object's existence is closed in this setting, with the simulation providing the decisive empirical verdict.

17.3 Control: a single expansion optic does not break the contraction

Claim. A control T is constructed by replacing the RPSI optic f_2 with an expansion (determinant 1.15, greater than 1). The control tests whether the Bregman-regularized contraction is robust to a single expansion optic, or whether it requires all seven optics to be individually contractive. The control simulation runs 8 configurations (2 starting sets x 4 λ values).

Method. Same simulation as Section 17.2 with the modified T . The bare- T contraction ratio is no longer the product of seven contractions but includes one expansion term; the theoretical prediction is indeterminate (depends on the relative magnitudes).

Evidence. All 8 control configurations converge geometrically. At $\lambda = 0.5$, $q = 0.5182$ ($R^2 = 1.0000$); at $\lambda = 0.7$, $q = 0.7075$ ($R^2 = 1.0000$). The control T contracts at rates indistinguishable from the canonical T at the same λ , indicating that the Bregman regularization and the other six contractions dominate the single expansion.

Implication. The Bregman-regularized contraction is robust to the substitution of a single expansion optic. This is a falsification safety net: the contraction survives the perturbation, rather than being an artifact of the specific contraction choices. A control with multiple simultaneous expansions, or with stronger expansions, remains an open research target.

Figure 17.1. T iteration convergence. Left: canonical T (all seven optics contractive). Right: control T (f_2 replaced by expansion). Log-log plot of Hausdorff distance $d_H(K_n, K_{n+1})$ versus iteration n . All 20 canonical configurations and all 8 control configurations show geometric decrease; fits at λ in $\{0.5, 0.7\}$ have $R^2 = 1.0000$ with q in $[0.51, 0.71]$.





Figure 17.2. Trajectory of the point set K_n under canonical T with $\lambda = 0.3$ from a 5x5 grid start. The set contracts from a 25-point grid spread across $[0,1]^2$ ($n=0$) to a tight cluster near the fixed point ($n=7$), demonstrating the empirical signature of the Banach contraction: rapid geometric collapse to the unification object.



Implications and Open Problems

The research targets that the surviving findings imply, with their binding prerequisites and falsifiability status.

Five implications follow from the surviving findings. Each implication is paired with the open problem it raises and the binding prerequisite for testing the problem.

Implication 1 (RESOLVED, Section 17): the multi-arc chain of the source transcript cannot be claimed as a rigorous unification. The single-composition theorem of Section 16 supplies the missing construction: each arc is an optic in $\text{Optic}(C)$, and the seven-fold composition is well-defined, associative, and unital. The empirical verification of the Bregman-regularized contraction of T , formerly the open problem, is now reported in Section 17. The numerical simulation iterates T on four starting compact subsets at five Bregman-regularization strengths; all 20 configurations converge geometrically ($q < 1$, $R^2 = 1.0$ at moderate λ ; machine-precision convergence at low λ). The unification object (fixed point of T) exists by Banach's contraction theorem. The remaining open problem is the extension of the simulation to non-contraction optics and to higher-dimensional base spaces, which would test the robustness of the contraction beyond the setting of Section 17.

Implication 2: the inverse-limit construction of the directed system of RAFs is a research target, not an achieved result. The source transcript aspires to an inverse limit but does not construct one (Section 2.3 and the replacement target noted there). The binding prerequisite is the explicit construction of the directed system, with transition maps between RAF instances that satisfy the directed-system axioms. The inverse limit is then the standard categorical construction. The single-composition theorem of Section 16 supplies the categorical setting in which the directed system can be constructed: the directed system is a diagram in $\text{Optic}(C)$, and the inverse limit is the limit in $\text{Optic}(C)$, which exists because $\text{Optic}(C)$ is complete when C is.

Implication 3: the hard problem of consciousness requires an organizational-invariance treatment, addressable via the CPTP-Zeno program of Section 9 but not via the original RPSI definition. The RPSI self-reference paradox is unresolved in the classical Markov setting. The CPTP lift resolves it but commits to a quantum-instantiated agent. The Zeno scaling confirmation of Section 15.2 establishes that the CPTP lift carries an empirically distinct signature. The binding prerequisite for the CPTP-Zeno treatment is the construction of a quantum agent whose measurement schedules are controllable; this is a non-trivial experimental commitment, now supported by the empirical confirmation that the Zeno scaling is distinguishable from classical scaling.

Implication 4: the $n=3$ prototype is sufficient for Claims A through E but insufficient for Claim F. The Lie algebra $\mathfrak{so}(2)$ is 1-dimensional abelian, so all perturbations commute trivially at $n=3$; the commuting-control test cannot distinguish parallel from non-parallel rotations. An extension to n at least 4 is binding for Claim F. This is a computational and experimental extension, not a theoretical innovation: the construction at n at least 4 is the same as at $n=3$,

with the dimension of the rotation subgroup increased. Section 15.1 reports the empirical confirmation of Claim F at $n=4$; the prototype extension is operational.

Implication 5: quantum instantiation of the agent is binding for Claim G in the non-ergodic regime. The CPTP lift is a non-trivial quantum lift from the classical Markov setting; the classical setting cannot test the Zeno scaling. Section 15.2 reports the empirical confirmation of Claim G in a two-level quantum simulation; the Zeno scaling exponent of 2 is distinguishable from the classical scaling exponent of 1 by approximately a factor of 2. The binding prerequisite for full operationalization is the construction of a quantum agent whose measurement schedule is controllable and whose Liouvillian spectral gap is measurable. This is an experimental commitment, supported by the empirical confirmation that the scaling regimes are distinguishable.

The five open problems are research targets, each with a clear falsifiability status. The single-composition theorem is constructed in Section 16; its falsifiable target, the Bregman-regularized contraction of T , is empirically CONFIRMED in Section 17. The inverse-limit construction is falsifiable in the sense that, once constructed in $\text{Optic}(C)$, it must produce a viability-weighted curvature that matches the operational form of Section 6.4. The CPTP-Zeno treatment is falsifiable via Claim G, which is empirically confirmed in Section 15.2. The n at least 4 extension is falsifiable via Claim F, which is empirically confirmed in Section 15.1. The quantum-instantiated agent is falsifiable via Claims F and G combined, both of which are empirically confirmed.

The project's overall falsifiability status is therefore operational, with three of the seven claims and targets empirically confirmed: Claim F (Section 15.1, n at least 4 structure group), Claim G (Section 15.2, CPTP+Zeno scaling), and Target 1 (Section 17, Bregman-regularized contraction of T). The remaining five claims (A through E) are open for empirical test in the n at least 3 prototype, with foundations confirmed. The project's defensible content is the conjunction of the surviving claims, the constructed single-composition theorem (with empirical contraction confirmation), and the empirically confirmed foundations, with no claim that exceeds the evidence presented.