
MANUSCRIPT REVIEW · REPAIR GUIDE AND CROSS-CHECK

Eight Concrete Fixes and the Application-Paper Cross-Check

Drafted replacement text for the companion paper's eight repair-set items, each with targets, verification, and application-paper interaction - plus the line-level re-check of all nineteen repaired items against journal_manuscript_v3.tex, the newly-provided application paper (zai2026measure), ten cross-paper findings, and five new application-paper findings.

Editorial Assessment Unit

Full read of the application paper (1,956 lines) · item-by-item re-examination of the verified repairs · joint submission-readiness actions for the two-paper pair

Table of Contents

1. The Application Paper, Identified	1
2. Eight Concrete Fixes for the Companion Paper	2
Fix 1 - One optic formalism (Riley pairs); retype the unification corollary	2
Fix 2 - Rewrite (or demote) the optic-colimit Step 3	3
Fix 3 - One catalysis condition (Hordijk-Steel closure) across definition, operator, and appendix	4
Fix 4 - Align the gluing hypothesis; relabel the remainder	4
Fix 5 - Re-normalize the Lévy-area term; restate the fatigue law in fluctuation language	5
Fix 6 - Enlarge the state box; instantiate K; drop the dangling hypotheses	5
Fix 7 - Make the AcCoA remark match the table; implement or amend step (v)	6
Fix 8 - Restrict the envelope domination to active surrogates; pose the ratio form as open	6
3. Re-Checking the Repaired Items Against the Application Paper	7
4. Cross-Paper Consistency Findings (C1-C10)	9
5. New Findings in the Application Paper	11
6. Joint Submission Readiness	13

1. The Application Paper, Identified

The upload **journal_manuscript_v3.tex** (1,956 lines, 107.5 KB, read in full for this cross-check) is the previously-unseen **application paper** that the companion theory paper cites under the key `zai2026measure`. The identification is not in doubt: the file's header comment states the repository was renamed to `metabolic-curvature-measure` (the exact repository the companion names), the paper is "revision 3" based on a frozen v2 with complete proofs, the target journal is the Bulletin of Mathematical Biology, and the two papers cross-cite each other (the application paper cites the companion as `zai2026categorical` for "constructions, proofs, and machine verifications"; the companion cites the application paper as `zai2026measure`). Both are authored by X, dated September 2026, and both carry an AI-use declaration (GLM and DeepSeek, under the author's direction) in their declarations blocks.

The paper's title is "A Measure-Theoretic Discrete Curvature Framework for Metabolic Gene Sensitivity: From Active-Set Geometry to Transcriptional Response." Its mathematics is classical and self-contained: the distributional second derivative of the piecewise-affine optimal-flux map of parametric FBA is a matrix-valued Radon measure concentrated on the codimension-one strata of the active-set complex (Definition μ); the two-layer structure (value/flux) with an exact coupling identity $D^2\Phi = \sum_r c_r D^2v_r^*$ (Theorem coupling); a refinement-resolution bridge with weak convergence, an exact counterexample calculus for total variation, and an L^1 reconstruction (Theorem Bprime); a regime dichotomy - curvature mass $O(\epsilon^2)$ for smooth maps versus $O(\epsilon)$ for piecewise-affine maps (measured slopes 2.000 and 1.00); and the measure-derived gene-sensitivity metric κ^μ (Definition $\kappa\mu$). The paper explicitly states: "No new measure-theoretic theorem is claimed."

The empirical core: on E. coli (iJO1366; M3D microarrays, $n = 424$ genes with nonzero κ^μ), κ^μ predicts carbon-depletion transcriptional response (Pearson $r = +0.395$, $p = 2.6 \times 10^{-17}$; partial $r = +0.269$; metric invariance $\rho(\kappa^\mu, \kappa_V^{\text{lex}}) = 0.99998$); the association is tie-break-robust across five stage-3 selection rules ($r \in [+0.386, +0.396]$, partial constant to three decimals); it generalizes to oxygen-limitation and acetate-switch paths ($r = +0.32, +0.22$); the event measures stabilize at the Glivenko-Cantelli rate across random panels; epistasis aligns with active-set footprint overlap (Spearman $\rho_S = 0.865$); genotype loops close with phenotypic memory (66% non-reverting); and the association does NOT propagate to the protein layer (protein-level $r = -0.083$ against transcript $+0.419$ on the same genes) - the protein-layer null, exactly as the companion describes it.

The division of labor is declared on both sides: the application paper carries only a brief adapted statement of the categorical reading (its section 2.5, "The categorical reading, in brief"), which summarizes - in prose - the companion's stratified base, the gauge matching condition, the gluing theorem, the Fisher-minimal transport law, the small-loop viability-holonomy theorem, the optic-composition semantics with its contraction analysis, the filtered-colimit construction of autocatalytic sets, and the homotopy-type-theoretic extension, and states that "nothing below is a premise of the empirical results." The two papers share no verbatim passages and, per their own statements, neither depends on the other's claims - a structure this cross-check confirms, with the exceptions noted in findings C1 and C2 below.

This document does two things. Chapter 2 drafts the concrete fixes for the companion's eight repair-set items - each with the exact target lines, replacement text ready to paste, the verification argument, and the application-paper interaction. Chapters 3-6 then re-check every repaired item (and the eight fixes) against the application paper read at line level: the nineteen verified repairs are re-examined for cross-paper consistency, ten cross-paper findings are recorded, five new application-paper findings are logged, and a joint submission-readiness list closes the report.

Companion cross-claim about zai2026measure	Application-paper status	Verdict
Atomic curvature measure on the parametric-FBA active-set complex	Definition μ + Theorem Bprime (weak convergence, KR norm, counterexample calculus)	Verified
Measure-derived sensitivity metric κ^u	Definition κ_{mu} (event-mass metric along parameter trajectories)	Verified
Tie-break robustness across selection rules	Section v8: five rules (declared + four variants), r in $[+0.386, +0.396]$, Table v8	Verified
Protein-layer null	Sections e26-e27: abundance-change $r = +0.008..+0.032$; protein-level $r = -0.083$	Verified
Empirical validation with genome-scale data	iJO1366 + iML1515; M3D, PRECISE, PaxDb, Schmidt proteomics; 98 automated numeric checks	Verified

Table 1. The companion's cross-claims about zai2026measure, verified against journal_manuscript_v3.tex (1,956 lines, read in full).

2. Eight Concrete Fixes for the Companion Paper

Each fix below names its target lines in companion_categorical_v2.tex, then gives the replacement text (one line or one paragraph, ready to adapt), the reason it works, and the interaction with the application paper. The fixes are ordered as in the repair set; together they close every item the severity triage marked fatal-if-refereed. Nothing here touches the verified-sound list of the ledger deliverable - each fix is scoped to the defective lines only.

Fix 1 - One optic formalism (Riley pairs); retype the unification corollary

Targets: Def *def:optic* (L372-382); Thm *thm:composition proof* (L1563-1589); Cor *cor:unification* (L1591-1599); Def *def:hott-optic* (L3618-3627); Table 1 residual column.

Replace Def def:optic by Riley's pair form: an object of $\text{Optic}(C)$ is a pair (M, C) of morphisms of the monoidal category C , with the source/target conventions of Riley's Categories of Optics; a morphism (M, C) to (M', C') is an equivalence class $[R, f, g]$ of triples with $f: M \otimes R \rightarrow M'$ and $g: C' \rightarrow C \otimes R$, quotiented by the coend relation (Riley's hom-formula); composition is Riley's Proposition 2.3, which the composition theorem's proof already cites. Delete the clause "satisfying the residual-compatibility condition" (or define that condition as the coend equivalence). Declare Def def:hott-optic the DIAGONAL specialization (forward and backward carriers identified per interface), not "the" generalization. Retype Cor cor:unification: replace " $T(O^*) = O^*$ " by " $R(O^*)$ is the unique fixed point of the realization map $R(O): S \rightarrow S$ (Thm unconditional-banach)" - T is an optic morphism and has no object-level fixed points. Replace the English residual phrases of Table 1 by the residual objects of the instantiation.

Why it works. The definition, the filtered-colimits theorem, and the composition proof currently use three incompatible shapes (F1); Riley's Prop. 2.3 is already the imported composition law, so restating the definition in his shape removes the translation gap rather than adding machinery. The corollary currently re-asserts exactly what its own theorem disclaims (a morphism does not act on objects); the realization $R(O)$ is the object that has fixed points.

Application-paper interaction. The application paper mentions optics only in prose ("the optic-composition semantics with its contraction analysis") with no formalism of its own, so the realignment creates no cross-paper conflict. No application-paper edit needed.

Fix 2 - Rewrite (or demote) the optic-colimit Step 3

Targets: Thm thm:filtered-colimits-optic, Step 3 (L3262-3275); Step 2 justification.

Corrected Step 3 (Set-level): fix i and a common upper bound $k \geq i$; define the cocone leg at i as the optic morphism $[R_i, f_i, g_i]$ with $f_i: M_i \otimes R_i \rightarrow M_{\max}$ obtained from the covariant colimit cocone of the M -components, and $g_i: C_{\max} \rightarrow C_i \otimes R_i$ the backward structure map declared by the construction (in the RAF instantiation, built from the same inclusions). Commutativity for $j \geq i$ reduces, by associativity of the optic composition law, to the componentwise identities $f_j \cdot (f_{ij} \otimes \text{id}) = f_i$ after the cocone identifications (same on the backward side), each of which is compatibility already required of the directed system; well-definedness of the classes follows from directedness (any two bounds have a common refinement). Remove the undefined map and the element-level pullback; replace Step 2's justification with "representatives exist by the construction of Set-colimits." If the general $\text{Optic}(C)$ statement is retained, add the missing hypothesis - a covariant transport for the backward carriers along the diagram arrows, which generally does not exist - or restrict the theorem to diagrams whose backward components are carried by declared structure, as in the instantiation.

Why it works. The backward carriers are contravariant along a directed diagram's arrows, so the displayed cocone map (which needs a $C_j \rightarrow C_i$ structure map that does not exist) cannot be written down; the corrected construction builds the backward legs from declared structure instead, which is exactly what the RAF instantiation provides, and proves what the scale-up verification actually checks.

Application-paper interaction. The application paper's delegation sentence lists "the filtered-colimit construction of autocatalytic sets" among results developed "with its own proofs and machine verifications." After this fix that sentence is safe for the RAF-level statement; if the demotion option is chosen, qualify the delegation wording per finding C1.

Fix 3 - One catalysis condition (Hordijk-Steel closure) across definition, operator, and appendix

Targets: Def def:raf (L384-393, conditions (ii) and (iii)); Def def:deflationary (L3479-3484); Appendix claim (L4342-4349).

Define once, for a reaction set R' and food set F , the closure $cl_{R'}(F) := \bigcup_{n \geq 0} S_n$, where $S_0 = F$ and $S_{n+1} = S_n \cup \{\text{react}(r), \text{prod}(r) : r \in R', \text{react}(r) \subseteq S_n\}$ - the least set containing F and closed under the reactants and products of applicable reactions (the Hordijk-Steel closure). Then (i) in Def def:raf replace the catalyst condition by "for each $r \in R'$, $\text{cat}(r) \cap cl_{R'}(F) \neq \emptyset$ " and replace the dangling "closed under the induced closure operator" by " $cl_{R'}(F) \subseteq S$ "; (ii) in Def def:deflationary replace the catalyst pool $\text{prod}(S) \cup F$ by $cl_S(F)$; (iii) qualify the appendix claim to "equivalent to the RAF conditions of Hordijk-Steel under well-founded (food-generated) dependency; without well-foundedness the operator admits futile autocatalytic cycles ($X \rightarrow Y$, $Y \rightarrow X$ mutually catalyzing, nothing food-reachable) that food-generation excludes."

Why it works. This is the alignment that makes all three stated conditions the same condition. The closure operator is monotone in S , so the deflationary operator remains monotone and the Tarski fixed-point argument of the appendix - verified sound in the ledger - is preserved verbatim. The qualification is needed because $cl(F \cup \text{prod}(S)) = F \cup \text{prod}(S)$ identically, so the current product-pool version is strictly weaker than Hordijk-Steel in exactly the futile-cycle direction.

Application-paper interaction. No catalysis content appears in the application paper; the fix is companion-internal. The delegation sentence (C1) is unaffected because the posetal terminal-coalgebra proof is complete and correct.

Fix 4 - Align the gluing hypothesis; relabel the remainder

Targets: Thm thm:2cat-gluing hypothesis (L612-614); eq:piecewise-F (L688-693).

In the gluing theorem, replace the hypothesis " $g^*ij A_j = A_i + d(\log g_{ij}) - [A_i, \log g_{ij}]$ " by the matching condition of Definition (O3), verbatim: $A_j|_{U_i \cap U_j} = g_{ij}^{-1} A_i g_{ij} + g_{ij}^{-1} dg_{ij}$. In eq:piecewise-F replace " $O(\epsilon^3)$ " by " $O(\epsilon^2)$ " with the parenthetical: "(for affine wall transitions the terms neglected beyond the displayed $O(\epsilon)$ and $O(\epsilon^2)$ contributions are $O(\epsilon^3)$; for general transitions, second-order Taylor terms of $\log g$ along the wall and commutators of the two resets enter at $O(\epsilon^2)$)" - matching Open Problem 1.

Why it works. The proof already glues "via the matching condition (O3)" (L631-635); only the stated hypothesis lags behind the repaired definition. The remainder relabel is forced by the Taylor expansion of $\log g$ for non-affine transitions; the affine regime is the one the numerics verify, so the parenthetical preserves the verified claim while making the general statement honest.

Application-paper interaction. STRONG CORROBORATION: the application paper's categorical subsection already describes the companion's gluing in the repaired form - "connection 1-forms on each stratum, glued along the switching facets by a boundary transition satisfying a gauge matching condition" (L466-473). The application paper is describing the definition; after this fix the theorem matches both. No application-paper edit needed.

Fix 5 - Re-normalize the Lévy-area term; restate the fatigue law in fluctuation language

Targets: *Lemma ito-expand* (L2281-2344); *Thm levy-3half* (L2346-2382, scale at L2370-2371); *figure caption* (L2419); *Claims C/D* (L1393-1405); *Prop stress proof* (L2221-2223).

(a) In the lemma's statement define $X_2 :=$ half the loop integral of $(W_1 dW_2 - W_2 dW_1)$; replace " $\text{Var}(X_2) = 1/12$ " by " $\text{Var}(X_2) = 1/4$ " and the proof's "variance $\sigma^4/12$ " by " $\sigma^4/4$ "; insert the derivation line: "Var of the loop integral = 1, by Ito integration by parts - the integral equals $2 \int W_1 dW_2 - W_1(1)W_2(1)$, with $\text{Var}(\int W_1 dW_2) = 1/2$ and the covariance $1/2$, hence $4(1/2) + 1 - 4(1/2) = 1$; the $1/12$ is the Brownian-bridge time-integral variance, not the free-Brownian-motion Lévy area." (b) Propagation: replace the sub-leading scale " $\sigma^2/(2\sqrt{12})$ " by " $\sigma^2/4$ " (figure caption term " $(\nu/2\sqrt{12})a$ " by " $(\nu/4)a$ "); report the Monte Carlo two-term fit's second coefficient so the normalization is testable; and state the exposure assumption explicitly - the $3/2$ exponent requires the effective diffusion to satisfy $\sigma_{\text{eff}}^2 = \nu \cdot a$ (noise exposure proportional to loop size); with $\sigma^2 = \nu$ fixed, the fluctuation scales linearly in a . (c) Restate Claims C and D and the stress proposition: the per-loop increment has mean $\pi a^2 + O(\sigma^3)$ and leading standard deviation $\sqrt{(5\nu)/2} \cdot a$ (the $a^{3/2}$ form under the amplitude-scaled exposure); the cumulative fatigue sum has mean $\Sigma \pi a_k^2$ and standard deviation the square root of $\Sigma (5\nu/4) a_k^2$; the failure condition reads: $\Sigma \pi a_k^2 + z \cdot (\text{fluctuation scale}) < \text{threshold}$, at confidence z . If the intended object is the mean absolute increment, the constant is $\sqrt{(2/\pi)} \cdot \sqrt{(5\nu)/2}$, not $\sqrt{(5\nu)/2}$. Delete every use of "deterministic mean" for the C_{fat} term.

Why it works. Three-way inconsistency ($1/12$ versus $1/3$ versus the true $1/4$ -or- 1) plus a fluctuation term used as drift: both were re-derived by hand in the re-read. The $5/4$ variance computation for the Gaussian term - the source of C_{fat} - is correct and untouched; only the sub-leading scale and the interpretation change. The exponent statement matters because $\text{Var}(G) = (5/4)\sigma^2 a^2$ makes the standard deviation linear in a unless the diffusion itself scales with a - which the paper must state to claim $3/2$.

Application-paper interaction. NONE: the application paper contains no Lévy, fatigue, or stochastic-calculus content (grep-verified) - the entire fatigue thread is companion-internal. This fix requires no application-paper edit and creates no cross-paper inconsistency.

Fix 6 - Enlarge the state box; instantiate K; drop the dangling hypotheses

Targets: *The instantiation* (L1670-1690); *Thm unconditional-banach* / *Prop treg-lip* (L1749-1755); *Prop qbound* (L2625-2630); *the KM gloss* (L1752-1753).

Replace $X := [0,1]^d$ by $X := [-1.5, 1.5]^d$, with the verification paragraph: each contracting $f_i(x) = (1 - \alpha)x + \alpha \cdot \tanh(Wx + b)$ maps any box of half-width B into $(1 - \alpha)$ -times the box plus $\alpha \cdot (-1, 1)^d$ and has global Lipschitz constant $(1 - \alpha) + \alpha \rho \leq 0.92$; the chain's extreme output is attained immediately after $f_2 = 1.15 \cdot \text{Id}$, giving $1.15(0.6B + 0.4s)$; invariance requires $1.15(0.6B + 0.4) \leq B$, i.e. $0.69B + 0.46 \leq B$, i.e. $B \geq 0.46/0.31 \approx 1.484$; $B = 1.5$ satisfies it with margin ($1.495 \leq 1.5$), and every subsequent map contracts further. Hence $T: X \rightarrow X$ and $\text{Lip}(T) \leq 0.92^6 \cdot 1.15 = 0.697 < 1$ - Banach applies on the enlarged box. Disclose W_i, b_i (or state the bounds as uniform over all pairs with the declared norms). Instantiate K : either $K := X$ (then the projection is the identity, $T_{\text{reg}} = T$, and the "reorganizes toward the interior" narrative is dropped) or $K := [0,1]^d$ (then $T_{\text{reg}}(X) \subseteq (1 - \lambda)[-1.5, 1.5] + \lambda[0,1] \subseteq X$, so $T_{\text{reg}}: X \rightarrow X$). Delete the never-used monotonicity and cocoercivity hypotheses from the qbound proposition, or delete the proposition as a near-duplicate. Replace the gloss "Krasnoselskii-Mann-averaged (Bregman-regularized)" by "Krasnoselskii-Mann-averaged (Euclidean projection)."

Why it works. The current hypothesis $T: X \rightarrow X$ is false on $[0,1]^d$ (the chain escapes to roughly $[-1.0, 1.01]^d$, and the KM-averaged map also leaves $[0,1]^d$); the box threshold $B \geq 0.46/0.31$ is sharp for the chain's worst point and 1.5 clears it with margin. No Bregman divergence appears anywhere in the construction, so the gloss must go; the cocoercivity hypotheses are unused in the proof they decorate.

Application-paper interaction. The application paper's delegation sentence explicitly lists "the optic-composition semantics with its contraction analysis" as companion-carried; that delegation is only safe once this fix makes the contraction theorem's hypotheses hold for the stated instantiation. Recommended application-paper action: none required, but the fix is a prerequisite for the delegation wording (C1).

Fix 7 - Make the AcCoA remark match the table; implement or amend step (v)

Targets: Remark netG-accolla-cycle (L3954-3975); Table network-battery (L4095-4096); L3760-3761; L4068-4074; Def def:autopoiesis step (v) (L1209-1213); Table netA rows c and f.

(a) Replace the remark's conversion sentence by: "Network H converts 43/44 to 44/44 and Network I likewise (Table network-battery); Network G remains 41/42 in both phases - its AcCoA limit cycle is NOT converted by the Phase III criterion; the two Phase-I failures that Phase III converts are those of H and I." Fix the metabolite-to-network attachment (AcCoA to H, ALA to I, per the table) and the backwards cross-reference range. (b) Either implement the restoration control - add to the operational protocol: "after recovery-above-threshold, require the return to the reference state (distance to x^* at most ϵ_0 for all steps in the declared window); rows c (3.84 against baseline 0.45) and f (443.8 against 2.65) then fail, and the verdicts must be recomputed" - or amend Definition autopoiesis step (v) to the implemented criterion and add the remark: "the operational test implements steps (i)-(iv) plus collapse-reversal, not the restoration control of step (v); the definition is thereby weakened to the protocol actually run."

Why it works. "42/42" occurs exactly once in the paper (grep-verified) and contradicts the battery table, which shows the conversion at H; and the passing rows of Table netA overshoot their baselines by factors of 8.5 and 167, violating the step-(v) control that the definition still requires. Either arm of (b) restores definition-protocol identity; only the first arm keeps the current wording of the definition.

Application-paper interaction. The network battery is companion-only - the application paper contains no autopoiesis or Phase-III content (grep-verified: no autopoiesis, no RAF beyond the delegation sentence). Companion-internal fix.

Fix 8 - Restrict the envelope domination to active surrogates; pose the ratio form as open

Targets: Thm smooth-envelope item (e) (L2969-2983); its proof (L3019-3026); eq:kappa-alg (the stray a-supremum).

Replace item (e) by two statements. (e1) "For every finite family $\{r_q\}$ and every x , the directional derivative of the envelope satisfies $E'(x; v) = \max \text{ over } q \text{ in } A(x) \text{ of } r_q'(x; v)$, where $A(x)$ is the argmax (the active set); consequently the Clarke identification of item (c) is recovered, and $E'(x; v) \geq r_{q^*}'(x; v)$ for every MAXIMIZER q^* in $A(x)$." (e2) "Domination over non-active surrogates is false in general ($E = \max(0, g)$ with $g(x) = x$ at $x = -1$: $E'(\cdot; +1) = 0 < g' = 1$); the quantity defined in eq:kappa-alg therefore bounds a raw directional derivative of E , and its comparison with the ratio-form κ_α of Definition def:kv is recorded as an open problem, not claimed as a theorem." Define or delete the stray supremum over a in $(0, a^*]$ in eq:kappa-alg (the symbol a^* occurs nowhere else in the paper). If the discretization bridge asserted in the application paper is to be retained, state it as a definition-level correspondence at ACTIVE atoms - the per-event measure mass corresponds to the positive part of the margin change at the crossed facet, normalized by the local margin - and mark it explicitly as a correspondence rather than a proved domination.

Why it works. The pointwise domination $E \geq r_q$ implies derivative domination only for active q (the counterexample above is two-line-verifiable); the finite-max calculus gives the correct active-set statement, which is also exactly what the Clarke subdifferential item needs. The ratio-form comparison was never proved and is what the application paper's bridge sentence actually needs - so the honest move is to restrict the claim and pose the bridge as open.

Application-paper interaction. THE STRONGEST INTERACTION IN THE PAIR: the application paper's categorical subsection asserts that the per-pair form of κ^H "is the measure-theoretic discretization of that object" (the ratio-form viability-weighted curvature) and that the value-flux coupling "supplies the survival covectors." After this fix the companion no longer supports the strong reading of that sentence. Recommended application-paper edit (one line): soften to "is modeled on / is the measure-theoretic analogue of that object (the precise discretization correspondence is open)" - see finding C2.

3. Re-Checking the Repaired Items Against the Application Paper

Every repaired item verified in the ledger (R1-R19) was re-examined against the application paper read at line level. The table records, for each repair, what the application paper contains on the same subject and whether the two papers are consistent. The verdict column uses: CONSISTENT (the app paper corroborates or is silent in a compatible way), CORROBORATED (the app paper positively supports the repair's content or direction), and CONSISTENT-BY-SEPARATION (the app paper has no content on the subject, and the division of labor makes that expected). No contradiction was found between any repaired item and the application paper. The two systemic exposures are recorded separately as findings C1 (delegation wording) and C2 (the discretization bridge).

Repair	Application-paper content on the subject	Verdict
R1 - Vereshchagin-Vitanyi citation	No rate-distortion, Kraft, or enumerability content anywhere (grep-verified); the app paper's tools are parametric LP, Alexandrov, Radon measures	Consistent by separation

Repair	Application-paper content on the subject	Verdict
R2 - pairs-crossing holonomy	Regime dichotomy (Prop dichotomy): $O(\epsilon)$ curvature mass for piecewise-affine maps, measured slope 1.00; smooth maps slope 2.000 - the measure-level analogue of the repaired boundary term	Corroborated
R3 - matching condition (O3)	Categorical subsection (L466-473) describes stratified connections glued by "a boundary transition satisfying a gauge matching condition" - the repaired form	Corroborated
R4 - C2 retraction (softplus)	No C2 claims anywhere; regularity language is measure-theoretic (Radon, Alexandrov, Clarke-free); "No new measure-theoretic theorem is claimed"	Consistent by separation
R5 - uniform-lip repair	No Lipschitz-envelope content	Consistent by separation
R6 - countable dense subfamily	No supremum-enumerability content	Consistent by separation
R7 - envelope honesty (a)-(e)	Interacts via the discretization bridge - see Fix 8 and finding C2	Consistent (after Fix 8)
R8 - geometric-phase corollary	No geometric-phase computation referenced; holonomy claims are measure-level (rerouting mass, genotype-loop memory)	Consistent by separation
R9 - exact C_{fat} derivation	No Levy/fatigue/Ito content (grep-verified)	Consistent by separation
R10 - terminal coalgebra	Delegation sentence lists "the filtered-colimit construction of autocatalytic sets"; the posetal Tarski proof is complete - the delegation is safe for this item at the RAF level (see Fix 2 for the optic-level caveat)	Consistent
R11 - Theorem B demotion	No functorial-realization content; the app paper never relies on the functorial claim	Consistent by separation
R12 - set-level realization	No realization-functor content	Consistent by separation
R13 - colimit/cocone fix	Delegation prose consistent with the corrected cocone language	Consistent
R14 - $0=0$ retraction	No falsifiability-hierarchy content	Consistent by separation
R15 - battery reframing	No network battery, autopoiesis, or Phase I/II/III content (grep-verified)	Consistent by separation
R16 - Noether repair	No Bregman or Noether content	Consistent by separation
R17 - provenance clean	AI-use declaration (GLM + DeepSeek), sole author, SHA-256 manifests, funding/ethics/competing-interests blocks - the same clean standard	Corroborated
R18 - open problems honest	Own open items (the general complex case of the bridge conjecture; the deterministic refinement-sequence form) - complementary scope	Consistent
R19 - smaller repairs	No corresponding content	Consistent by separation

Table 2. The nineteen verified repairs of the ledger, re-examined against the application paper. No contradictions were found.

4. Cross-Paper Consistency Findings (C1-C10)

Ten findings from the cross-check - the systemic exposures (C1, C2), the verified cross-claims (C4, C7), the consistencies that need one sentence of reconciliation (C3, C5), and the practices each paper can borrow from the other (C6, C8). Line numbers prefixed APP- refer to journal_manuscript_v3.tex; unprefixed line numbers refer to companion_categorical_v2.tex.

C1 - The mutual delegation claim is overstated for four of the five delegated items.

Both papers state that the categorical framework is carried by the companion "with its constructions, proofs, and machine verifications" (APP-L457-459) / "developed in full, as a standalone theory paper with its own proofs and machine verifications" (APP-L1119-1128); the companion reciprocates ("the application paper cites it for constructions and proofs", L180-181). Checked item by item against the ledger: the Fisher-minimal transport law IS fully proved (Prop kkt, with a complete KKT proof, L960-1001); the gluing theorem is descent-by-citation for stratified bases with a hypothesis/definition mismatch (V3, V1); the optic-composition contraction analysis has the backward-component, self-map, and K-instantiation gaps (F3, F4, F21); the filtered-colimit construction has the ill-formed Step 3 (N5); and the HoTT extension is proof-sketch-marked (F9). So one of the five delegated subjects is fully proved as delegated; four carry documented defects or partiality. Recommended joint action: qualify the delegation wording in both papers ("with constructions, machine verifications, and proofs at the marked status per result") - or complete the eight repairs first, after which the wording is accurate for every item except the sketch-marked HoTT extension (which the companion itself flags).

C2 - The kappa-to-kappa-mu discretization bridge is asserted on both sides and proved on neither.

The companion's Definition def:kv remark (L475-477): "the measure-theoretic discretization used in the application paper replaces the 2-form F by an atomic crease measure." The application paper (APP-L482-491): "The per-pair form of κ^μ is the measure-theoretic discretization of that object - the atomic crease measure replaces the curvature 2-form, and the value-flux coupling supplies the survival covectors." Both sentences describe the same bridge; neither paper contains a theorem establishing it. In the companion, the envelope domination item that should carry it has the form mismatch and the active-surrogates-only gap (F6, fixed by Fix 8); in the application paper, the phrase "per-pair form of κ^μ " is itself undefined (the defined metric is per-gene; the per-pair object it alludes to is the companion's per-pair $\kappa_\alpha(\theta, x; u, v)$, which IS defined - but the correspondence between the two is exactly the unproved bridge). The bridge is currently a stated analogy. After Fix 8, the companion no longer supports the strong reading; the application paper should soften its categorical-subsection sentence accordingly (one-line edit), or the pair should state the correspondence at active atoms as a definition and mark its status.

C3 - The tie-break counting is consistent; only the adjective drifts.

The companion says "tie-break robustness across five lexicographic selection rules" (L1274); the application paper's claim structure says " $r \in [+0.386, +0.396]$ across five selection rules" (APP-L154-155), its tie-break section defines exactly five (TB0 declared plus TB1-TB4 variants, APP-L886-899), its figure caption says "the five stage-3 selection rules (declared plus four variants)" (APP-L1032-1033), and its abstract says "four alternative selection rules" (APP-L104). All five readings reconcile (five total = declared plus four). The companion's adjective "lexicographic" is loose - the variants are stage-3 rule swaps within the lexicographic protocol class (TB3 changes the stage-3 objective; TB4 maximizes it) - matching the application paper's own phrase "a property of the lexicographic protocol class." Wording only; no numerical discrepancy.

C4 - The companion's cross-claims about the application paper are accurate.

Each element of the companion's description was located and verified in the application paper: the atomic curvature measure on the parametric-FBA active-set complex (Definition mu; Theorem Bprime); the measure-theoretic sensitivity metric κ^u (Definition kmu); the selection-rule dependence exposed and stress-tested rather than masked (Remark lock + the tie-break experiment); the explicit protein-layer null (sections e26-e27, with the honest abundance-level association reported alongside); and the genome-scale empirical validation (the M3D association with its layer decision, path robustness, and event-measure stabilization). The companion describes its counterpart correctly - a nontrivial check, since the description was written from a paper the reading passes could not see.

C5 - The two holonomy-scaling headlines are consistent, but only after regime delineation - which neither paper states.

The application paper's abstract: "loop holonomy scales linearly in loop size (slope 1.00, versus the quadratic scaling of any smooth map)" - the atomic/wall-crossing regime (curvature mass captured by a loop of radius ϵ is $O(\epsilon)$ for piecewise-affine maps, $O(\epsilon^2)$ for smooth maps; APP Prop dichotomy proves both arms and the regime-dial measures crossover $\epsilon^*/\sigma \in [2.45, 4.11]$). The companion's headline: $\kappa(a) = a^2$ and the small-loop theorem's $\epsilon^2 F_{12}$ - the smooth-connection regime (loops contained in a single stratum; the Fisher-minimal transport law is the smooth avatar, per the application paper's own summary sentence). The two claims are complementary, not contradictory: the app paper's dichotomy proves BOTH regimes, and the companion's quadratic law lives in the smooth arm. But a referee reading the pair back-to-back meets "linear (slope 1.00)" in one abstract and "quadratic $\kappa(a) = a^2$ " in the other without the reconciliation sentence. Recommended: one bridging remark in each paper - the companion noting that wall-crossing loops accumulate $O(\epsilon)$ per crossing (already in its pairs-crossing theorem), the application paper noting that its smooth arm is the regime where the companion's quadratic law applies.

C6 - The application paper already contains the fix pattern the companion needs for its kappa-notation conflation.

The companion's F23 (three objects under one $\kappa_v(a)$ notation) has a ready-made resolution in the application paper's Discussion (APP-L1099-1109): "The three sensitivity objects of the earlier literature - the geometric κ_v , the flux-statistical κ_{flux} , and the time-course κ_{time} - are one measure at different resolutions: the atomic measure of Definition mu, its coarse-grainings, and its integrals along parameter trajectories." The application paper thereby names and unifies exactly the three objects the SAVGS lineage conflated. Recommended companion action (beyond the ledger's triage): import this resolution framing in the hierarchy/verdict sections, so that the pointwise $\kappa_v(\theta, x)$, the loop-amplitude $\kappa(a)$, and the loop-averaged $D_v(a)$ are declared as the smooth-avatar, the trajectory functional, and the depth statistic of the same object.

C7 - The circularity retirement is real on both sides.

The companion concedes the indicator circularity in words (F7) and removes the definition. The application paper's predictor is κ^u , whose per-event mask $w_g(g)$ selects events whose crossed stratum changes the activity of the gene's reactions - a structural property of the event (active-set change), not a knockout-experiment label; no essentiality call enters the predictor. The precursor geometric κ_v is retained only as a declared, weak baseline (panel-level r from -0.063 unmasked to +0.084 masked, APP-L774-775), reported honestly as the discrepancy that motivated the recalibration. The circularity exposure documented in the v3 read is closed in the pair as it now stands.

C8 - Count hygiene: the application paper models the practice the companion still needs.

The application paper's near-colliding counts (424/433/435/438/440/525/537, plus 1,516/2,779 and 25,107/4,000) are all disambiguated in a dedicated appendix table, each traced to a deposited artifact file (APP-L1260-1286), and the sweep census reconciles exactly (13 = glucose + oxygen + ten gene knockdowns + the iJO1366 replication; two affine; eleven crossing; the eno sweep at mass 0.934). This is the cure for the count-inconsistency disease documented in the SAVGS v3 read (six different network counts) - and the companion's own counting frictions (the G/H row, the 2/5 and 0/10 anchors, the 41/42 family) would benefit from the same one-page disambiguation table.

C9 - The application paper carries one visible provenance residue.

Figure paths embed the generation pipeline's name: the primary-association, path-robustness, tie-break, and stabilization figures resolve from the subdirectory deepseek_bridge/ (APP-L1006, L1020, L1031, L1041). Given the declared, supervised AI use this is harmless - but it is the same class of signal the v3 read flagged, and a one-command rename before submission removes it.

C10 - The declared division of labor is structurally sound and mostly holds.

"The two manuscripts share no verbatim passages of length, and neither depends on the other's claims" (APP-L499-502; the companion's status line concurs). Verified: the application paper's empirical claims are self-contained (statistical; every number traces to a deposited artifact), its mathematical layer is classical and fully proved internally, and its categorical subsection states "nothing below is a premise of the empirical results" - so the companion's defects do not touch the application paper's evidence chain. The residual exposure is confined to the prose of the categorical subsection (C2) and the mutual delegation wording (C1). The pair is decoupled in substance; it is coupled only in framing.

5. New Findings in the Application Paper

The full read of journal_manuscript_v3.tex also produced findings internal to the application paper itself. None is severe; the paper is in substantially better internal shape than either SAVGS manuscript or the companion. They are logged here for completeness because the user's instruction series has been line-level review throughout, and because two of them (APP-1, APP-2) sit inside load-bearing statements (the engine's regularity lemma and a proposition's stated law).

APP-1 (moderate) - The engine's regularity lemma proves a different selection rule than the engine uses.

Lemma lex (APP-L333-341) states regularity for "the three-stage lexicographic optimal flux" defined in the methods section; the methods engine's stage 3 is a weighted linear objective, $\min w\text{-transpose-}v$ with seeded weights $w \sim U(0.5, 1.5)$ (APP-L1171-1175), and the metric's definition declares the tie-break part of the metric (Remark lock). But the lemma's technical proof (APP-L1735-1748) establishes uniqueness by coordinate pinning: "the final stage minimizes the coordinates $v_1, \dots, v_n$ one at a time, in a fixed order, over the current face" - a DIFFERENT selection rule. The polyhedral chamber machinery that follows does cover linear-objective stages (each next stage is "again a multi-parametric LP value function," APP-L1785-1787), so only the well-definedness/uniqueness step is specific to coordinate pinning. For the seeded-weight rule, uniqueness holds only generically in w (outside a measure-zero set of weight vectors, for each chamber structure). One-paragraph fix: add to the proof "for stage-3 objectives of the declared form, the argmin over the compact stage-2 face is a unique vertex for every θ outside a Lebesgue-null set of weight vectors; the chamber and projection argument below then applies verbatim, and the locked seed (20240901) realizes a generic draw" - or restate the lemma for the coordinate-pinned completion and derive the engine's selection as its generic case.

APP-2 (minor) - The semiconvexity proposition states two contradictory forms of the same law.

Proposition semiconvex (APP-L384-388): "the required constant blows up like $1/(2h)$; $\lambda \cdot h_{\max} = 0.500$ numerically across $\lambda = 1 \dots 10^6$, the analytic law $\lambda \cdot h_{\max} = 1/(2\lambda)$ holding at every scale." The two clauses contradict: the first says the product $\lambda h_{\max}$ is the constant 0.500 across three orders of magnitude; the second says it equals $1/(2\lambda)$, which varies by six orders. The proof (APP-L1512-1514) repeats the second form. With "blows up like $1/(2h)$," the correct pair is $h_{\max} = 1/(2\lambda)$ and $\lambda \cdot h_{\max} = 1/2$ - so either the numeric clause should read $\lambda^2 \cdot h_{\max} = 0.500$, or the analytic law's denominator should be 2, not 2λ . One-character-plus fix; a referee will catch the contradiction.

APP-3 (cosmetic) - Label artifacts expose the revision history.

The decoupling corollary is labeled conj:valueflux (APP-L642-643) - the label prefix of a conjecture on a corollary environment, the residue of a conjecture-to-corollary promotion; and the refinement-bridge theorem is titled "corrected form" (APP-L514). Both are harmless but visible; rename the label (cor:valueflux) and drop "corrected form" from the title before submission.

APP-4 (cosmetic) - The deepseek_bridge figure directory (see C9).

Four figures load from deepseek_bridge/; renaming the directory (and updating the four includegraphics paths) removes the residue.

APP-5 (trivial) - Abstract versus claim-structure phrasing.

The abstract says "four alternative selection rules" where the claim structure and the tie-break section say "five selection rules" - consistent readings (declared + four variants), flagged only for copyediting.

Arithmetic and consistency spot-checks performed on the application paper during this read - all pass:

- The counterexample-calculus ratios: $3 - 1/128 = 2.9922$ and $1.4 - 1/128 = 1.3922$ - both match the claimed measured values exactly.
- The OR-gate concavity violation: with the max-gate value function, the midpoint violation is exactly $1/2$ (the two corners give 1, the midpoint gives $1/2$).
- $438/2,583 = 17.0\%$ active reactions; $537/2,583 = 20.8\%$ union - both match.

- The sweep census: 13 = glucose + oxygen + ten gene knockdowns + the iJO1366 replication; two affine; eleven crossing; the eno sweep's mass 0.934 with the sub-threshold kink explanation.
- 57 grid points = 7 segments at 8x refinement + 1; the four design corners carrying the event mass 288.77; value-kink census = 4.
- The tie-break table: r within $[+0.386, +0.396]$; maximum deviation 0.0093; the TB1 deviation traced to two boundary genes (n : 424 to 426).
- The p -values are plausible throughout ($r = 0.395$ at $n = 424$ gives t of about 8.8; p near 10 to the minus 17).
- F1-style consistency of the reported precision/recall pairs; the epistasis, path-robustness, and stabilization numbers are internally consistent with their sections and the count table.

6. Joint Submission Readiness

The pair is closer to submission than either SAVGS manuscript ever was: the application paper is internally sound (two one-line repairs and cosmetics), and the companion needs its eight drafted fixes plus the framing edits below. The actions are listed in the order they should be executed; the first three close every fatal-if-refereed item of the ledger's triage.

Companion - apply Fixes 1-8 (chapter 2).

Fixes 1, 2, 3, 4, and 6 close the fatal items (optic formalism and the unification corollary; the optic-colimit Step 3; the catalysis triple mismatch; the gluing hypothesis; the self-map and K gaps). Fix 5 repairs the sub-leading scale and the mean/fluctuation language without touching the verified C_{fat} constant. Fixes 7 and 8 close the two serious table/protocol items and the domination claim.

Companion - align the delegation and bridge language.

After the fixes, add to the status line: per-result proof status is marked in the paper (sketch-level for the HoTT extension; citation-level for the stratified descent); and restate the discretization remark of Definition def:kv as the correspondence-at-active-atoms version of Fix 8.

Application paper - four one-line edits.

(i) Soften the categorical-subsection bridge sentence per C2/Fix 8 ("is modeled on ... the precise discretization correspondence is open"). (ii) Add the generic-weights uniqueness step (or restate Lemma lex for the coordinate-pinned selection) per APP-1. (iii) Fix the semiconvexity law contradiction per APP-2. (iv) Rename the conj:valuelux label and drop "corrected form" per APP-3.

Both - the bridging remark for the scaling regimes (C5).

One sentence each: the companion notes that wall-crossing loops accumulate $O(\epsilon)$ per crossing (already in its pairs-crossing theorem, so this is a pointer, not new math); the application paper notes that its smooth arm is where the companion's quadratic law applies.

Both - ship the bibliographies.

Neither .bib file was provided with either upload; reference resolution cannot be checked for either paper (the companion cites *bhattacharya2007* and the *zai2026measure/zai2026categorical* pair; the application paper inputs *journal_manuscript_v3_bmb_refs.tex*, absent from the upload).

Companion - borrow the count-hygiene practice (C8) and the kappa-unification framing (C6).

A one-page disambiguation table for the battery counts and anchors; and the three-objects-one-measure resolution sentence imported from the application paper's Discussion to close F23's notation conflation.

Optional polish: rename the deepseek_bridge figure directory (C9/APP-4); correct the "fifteen orders of magnitude" to fourteen (companion N16); fix the stray-parenthesis and hyphenation artifacts (V6).